

Recall from last time ...

The Problem of Overfitting

A photograph of a wooden bed frame with a mattress that has been cut into the shape of the number 4. The mattress is white with a quilted pattern. The bed frame is made of dark wood and is placed on a light-colored wooden floor. The background is a plain white wall.

**THE BEST WAY TO
EXPLAIN OVERFITTING**




Underfitting (High Bias)



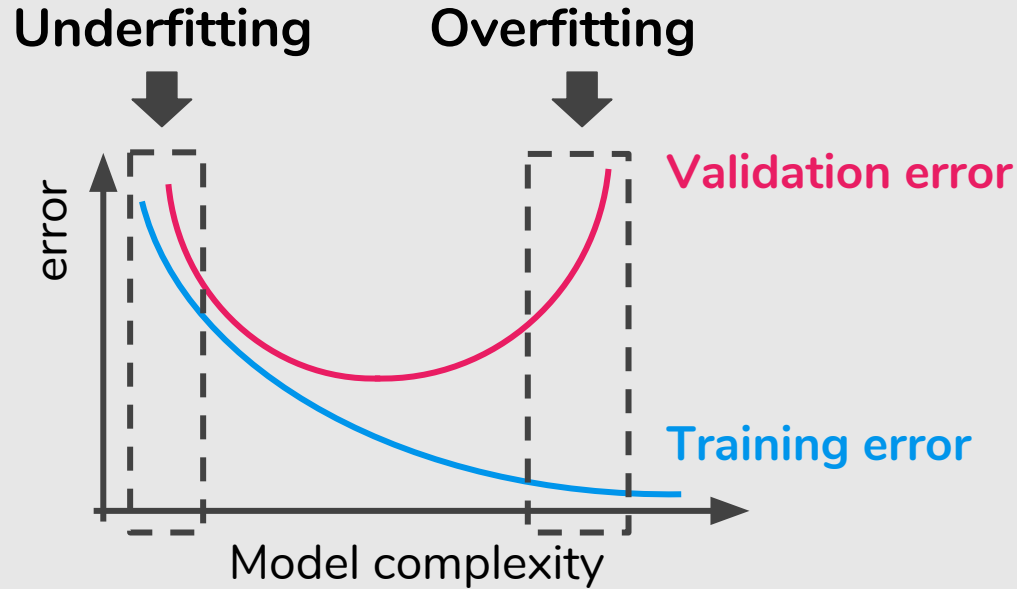
Overfitting (High Variance)

Diagnosing Bias vs. Variance

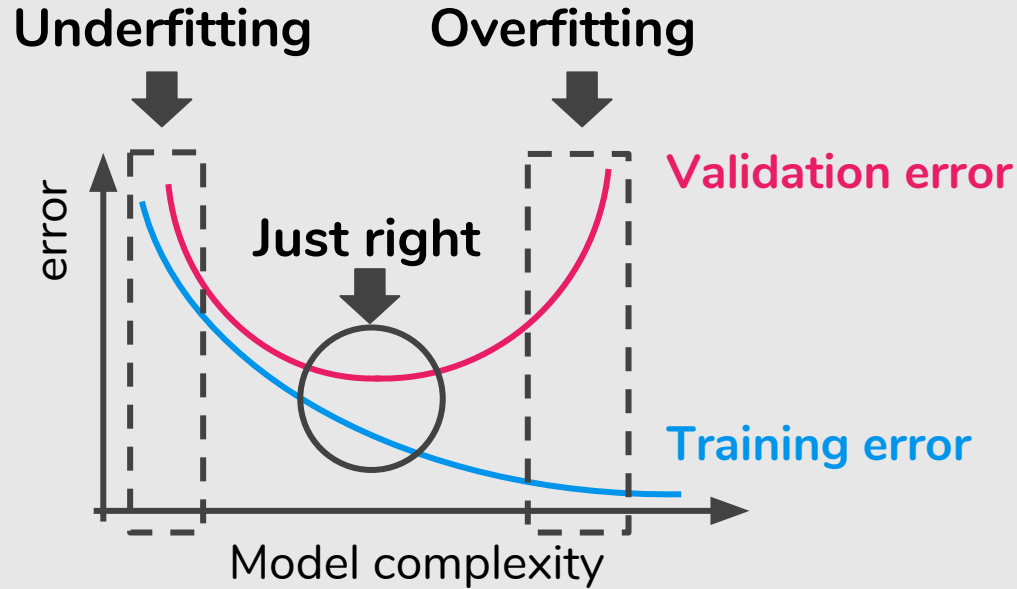
Diagnosing Bias vs. Variance



Diagnosing Bias vs. Variance



Diagnosing Bias vs. Variance




Cost Function

Regularization

$$J(\theta) = \frac{1}{2m} \left[\underbrace{\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2}_{\text{to fit the training data well}} + \underbrace{\lambda \sum_{j=1}^n \theta_j^2}_{\text{to keep the parameters small}} \right]$$

Regularization parameter



Regularized Linear Function

Gradient Descent

repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

} (simultaneously update θ_j for $j = \text{✗ } 1, \dots, n$)

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Gradient Descent

repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

} (simultaneously update θ_j for $j = \text{✗} 1, \dots, n$)

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Gradient Descent

$$h_{\theta}(x) = \theta^T x \rightarrow h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

} (simultaneously update θ_j for $j = \text{✗} 1, \dots, n$)

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

References

— — —

Machine Learning Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chap. 4
- Pattern Recognition and Machine Learning, Chap. 3
- Understanding the Bias-Variance Tradeoff:

<http://scott.fortmann-roe.com/docs/BiasVariance.html>

Machine Learning Courses

- <https://www.coursera.org/learn/machine-learning>, Week 3 & 6

Testing and Error Metrics

Machine Learning and Pattern Recognition

(Largely based on slides from Luis Serrano)

Prof. Sandra Avila
Institute of Computing (IC/Unicamp)

MC886/MO444, August 30, 2018



How well is my model doing?

Today's Agenda



— — —

- Testing and Error Metrics
 - Training, Testing
 - Accuracy
 - Precision
 - Recall
 - F-Score



Which model is better?




Which model is better?




Which model is better?





Why testing?





Why testing?





Why testing?





Why testing?





Why testing?



Why testing?



Why testing?



Friends don't let friends
use testing data
for training

Data



```
graph TD; Data[Data] --> Training1[Training]; Data --> Test1[Test]; Training1 --> Training2[Training]; Training1 --> Validation[Validation]; Training2 --> Training3[Training]; Training2 --> Validation3[Validation]; Training3 --> Test2[Test]; Validation --> Test2;
```

Training


Test

Training

Validation

Test


k-fold Cross Validation



k-fold Cross Validation




$k = 5$



k-fold Cross Validation




$k = 5$



k-fold Cross Validation




$k = 5$



k-fold Cross Validation




$k = 5$




k-fold Cross Validation




$k = 5$




k-fold Cross Validation




$k = 5$



$k \times 2$ -fold Cross Validation

 Training
 Validation


$k = 5$



$k \times 2$ -fold Cross Validation



$k = 5$



← randomized

A black arrow pointing left towards the blue bar, with the word 'randomized' next to it.


$k \times 2$ -fold Cross Validation




$k = 5$



$k \times 2$ -fold Cross Validation




$k = 5$



$k \times 2$ -fold Cross Validation



$k = 5$



← randomized

$k \times 2$ -fold Cross Validation



$k = 5$



$k \times 2$ -fold Cross Validation



$k = 5$



$k \times 2$ -fold Cross Validation



$k = 5$




...

k times = $k \times 2$ folds


Randomizing in Cross Validation

- Training
- Validation




Randomizing in Cross Validation

- Training
- Validation




Randomizing in Cross Validation

- Training
- Validation




Randomizing in Cross Validation

● Training
● Validation




Randomizing in Cross Validation

● Training
● Validation



Randomizing in Cross Validation

● Training
● Validation



MO850A: Tópicos Avançados em Ciência da Computação I — **Scientific Methodology**

Prof. Jacques Wainer (IC/Unicamp)


Evaluation Metrics

How well is my model doing?


Credit Card Fraud



Credit Card Fraud




Credit Card Fraud



Model: All transactions are good.


Credit Card Fraud



Model: All transactions are good.

$$\text{Correct} = \frac{284,335}{284,807} = 99.83\%$$


Credit Card Fraud




Model: All transactions are good.

Problem: I'm not catching any of the bad ones!

Credit Card Fraud




Credit Card Fraud



Model: All transactions are fraudulent.

Credit Card Fraud



Model: All transactions are fraudulent.
Problem: I'm accidentally catching all the good ones!

Medical Model



Health



Sick

Spam Classifier Model




Not Spam





Spam




Confusion Matrix Table

	Diagnosed Sick	Diagnosed Healthy
Sick		
Healthy		





Confusion Matrix Table

	Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	
Healthy		






Confusion Matrix Table

	Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	
Healthy		True Negative 

Confusion Matrix Table

	Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	False Negative 
Healthy		True Negative 

Confusion Matrix Table

	Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	False Negative 
Healthy	False Positive 	True Negative 






Confusion Matrix Table



10,000
patients

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
Sick	1000	200
Healthy	800	8000


Confusion Matrix Table

	Sent to Spam Folder	Sent to Inbox
Spam	True Positive 	False Negative 
Not Spam	False Positive 	True Negative 

Confusion Matrix Table


		Folder	
		Spam Folder	Inbox
1,000 emails	Spam	100	170
	Not Spam	30	700

Confusion Matrix Table




		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive		
	Negative		

Confusion Matrix Table




		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	
	Negative		

Confusion Matrix Table




		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	
	Negative		5 True negatives

Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	
	Negative	2 False positives	5 True negatives

Confusion Matrix Table




		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	1 False negative
	Negative	2 False positives	5 True negatives

Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●




Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●



Predicted Class


True Class		Predicted Class		
		Guessed Class 1	Guessed Class 2	Guessed Class 3
	Class 1			
	Class 2			
	Class 3			

Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●




True Class

Predicted Class

	Predicted Class		
	Guessed Class 1	Guessed Class 2	Guessed Class 3
Class 1	5	2	1
Class 2	3	6	0
Class 3	0	1	7

Confusion Matrix Table (n classes)



Accuracy



Diagnosis

Patients

	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Accuracy



Diagnosis


Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
Patients	1,000	200
	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Diagnosis



	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Accuracy =

$$1,000 + 8,000$$

Patients

Accuracy



Diagnosis

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
Sick	1,000	200
Healthy	800	8,000


Accuracy:

Out of all the **patients**, how many did we classify correctly?

Accuracy =

$$\frac{1,000 + 8,000}{10,000} = 90\%$$

Accuracy




	Folder		
	Spam Folder	Inbox	
Email	Spam	100	170
	Not Spam	30	700

Accuracy:

Out of all the **emails**, how many did we classify correctly?

Accuracy

 Email	Folder	
	Spam Folder	Inbox
	Spam	Not Spam
Spam	100	170
Not Spam	30	700


Accuracy:

Out of all the **emails**, how many did we classify correctly?

Accuracy =

$$\frac{100 + 700}{1,000} = 80\%$$


Accuracy



Accuracy:

Out of all the **data**, how many points did we classify correctly?

Accuracy




Accuracy:

Out of all the **data**, how many points did we classify correctly?

Accuracy =

$$\frac{\text{Correctly Classified Points}}{\text{All points}}$$

Accuracy



Accuracy:


Out of all the **data**, how many points did we classify correctly?

Accuracy =

$$\frac{\text{Correctly Classified Points}}{\text{All points}}$$

$$\frac{11}{11 + 3} = 78.57\%$$

Accuracy



Transactions	Prediction	
	Fraudulent	Not Fraudulent
Fraudulent	0	472
Not Fraudulent	0	284,335


Accuracy:

Out of all the **transactions**, how many did we classify correctly?

Accuracy =


$$\frac{0 + 284,335}{284,807} = 99.83\%$$

Normalized Accuracy



Transactions	Prediction	
	Fraudulent	Not Fraudulent
Fraudulent	0	472
Not Fraudulent	0	284,335

Normalized Accuracy




Transactions	Prediction	
	Fraudulent	Not Fraudulent
Fraudulent	0	472
Not Fraudulent	0	284,335

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$

Normalized Accuracy




Transactions	Prediction	
	Fraudulent	Not Fraudulent
Fraudulent	0	472
Not Fraudulent	0	284,335

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$
$$\frac{\frac{0}{0 + 472} + \frac{284,335}{284,335 + 0}}{2} =$$

Normalized Accuracy




Transactions	Prediction	
	Fraudulent	Not Fraudulent
Fraudulent	0	472
Not Fraudulent	0	284,335

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$
$$\frac{\frac{0}{0 + 472} + \frac{284,335}{284,335 + 0}}{2} =$$
$$\frac{0 + 100}{2} = 50\%$$

Normalized Accuracy

Accuracy = 80%

	Folder	
	Spam Folder	Inbox
	Spam	Inbox
Email	100	170
	30	700

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$
$$\frac{\frac{100}{100 + 170} + \frac{700}{700 + 30}}{2} =$$
$$\frac{37.0 + 95.9}{2} = 66.5\%$$

Normalized Accuracy

Accuracy = 90%











Diagnosis

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
Sick	1,000	200
Healthy	800	8,000

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$
$$\frac{\frac{1000}{1000 + 200} + \frac{8000}{8000 + 800}}{2} =$$
$$\frac{83.3 + 90.9}{2} = 87.1\%$$

	Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	False Negative 
Healthy	False Positive 	True Negative 

	<p>Diagnosed Sick</p>	<p>Diagnosed Healthy</p>
<p>Sick</p>		<p>False Negative</p> 
<p>Healthy</p>	<p>False Positive</p> 	



Sent to Spam
Folder

Sent to Inbox

Spam

True
Positive



False
Negative






Not Spam

False
Positive



True
Negative



	Sent to Spam Folder	Sent to Inbox
Spam		False Negative 
Not Spam	False Positive 	

Evaluation Metrics



Medical Model

False positives ok
False negatives **NOT** ok



Spam Detector

False positives **NOT** ok
False negatives ok

Evaluation Metrics



Medical Model

False positives ok
False negatives **NOT** ok
High Recall



Spam Detector

False positives **NOT** ok
False negatives ok
High Precision

Precision



Diagnosis

Patients

	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Precision



Diagnosis

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
	1,000	200
	800	8,000

Precision:

Out of all the patients we diagnosed with illness, how many were actually sick?

Precision



Diagnosis

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Precision:

Out of all the patients we diagnosed with illness, how many were actually sick?

Precision



		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000


Precision:

Out of all the patients we diagnosed with illness, how many were actually sick?

Precision =

$$\frac{1,000}{1,000 + 800} = 55.7\%$$

Precision




Email	Folder	
	Spam Folder	Inbox
Spam	100	170
Not Spam	30	700

Precision:

Out of all the emails sent to the spam inbox, how many did were actually spam?

Precision

Email	Folder	
		
	Spam Folder	Inbox
Spam	100	170
Not Spam	30	700


Precision:

Out of all the emails sent to the spam inbox, how many did were actually spam?

Precision =

$$\frac{100}{100 + 30} = 76.9\%$$


Precision



Precision:

Out of all the points we've predicted to be positive, how many are correct?


Precision



Precision:

Out of all the points we've predicted to be positive, how many are correct?

Precision



Precision:

Out of all the points we've predicted to be positive, how many are correct?

Precision =

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Recall



Diagnosis

Patients

	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Recall



Diagnosis

Patients	Diagnosis	
	Diagnosed Sick	Diagnosed Healthy
	Sick	Healthy
Sick	1,000	200
Healthy	800	8,000

Recall:

Out of all the sick patients, how many did we correctly diagnose as sick?

Recall



Diagnosis

Patients

	Diagnosed Sick	Diagnosed Healthy
Sick	1,000	200
Healthy	800	8,000

Recall:

Out of all the sick patients, how many did we correctly diagnose as sick?

Recall



Patients

		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Recall:

Out of all the sick patients, how many did we correctly diagnose as sick?

Recall =

$$\frac{1,000}{1,000 + 200} = 83.3\%$$


Recall

	Folder	
	Spam Folder	Inbox
Email	100	170
	30	700

Recall:

Out of all the spam emails, how many were correctly sent to the spam folder?

Recall



Email	Folder	
	Spam Folder	Inbox
Spam	100	170
Not Spam	30	700

Recall:

Out of all the spam emails, how many were correctly sent to the spam folder?


Recall =

$$\frac{100}{100 + 170} = 37\%$$

Recall

Recall:


Out of all the points labelled positive, how many did we correctly predict?




Recall

Recall:

Out of all the points labelled positive, how many did we correctly predict?



Recall




Recall:

Out of all the points labelled positive, how many did we correctly predict?

Recall =

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Recall



Recall:

Out of all the points labelled positive, how many did we correctly predict?

Recall =

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$\frac{6}{6 + 1} = 85.7\%$$

Precision and Recall



Medical Model

Precision: 55.7%

Recall: 83.3%



Spam Detector

Precision: 76.9%

Recall: 37%

One Score?



Medical Model

Precision: 55.7%

Recall: 83.3%

Average = 69.5%




Spam Detector

Precision: 76.9%

Recall: 37%


Average = 56.9%

Credit Card Fraud



Model: All transactions are fraudulent.


Credit Card Fraud



Model: All transactions are fraudulent.

$$\text{Precision} = \frac{472}{284,807} = 0.016\%$$

Credit Card Fraud




Model: All transactions are fraudulent.

$$\text{Precision} = \frac{472}{284,807} = 0.016\%$$

$$\text{Recall} = \frac{472}{472} = 100\%$$

Harmonic Mean

Arithmetic Mean = $\frac{x + y}{2}$




Harmonic Mean

Arithmetic Mean = $\frac{x + y}{2}$

Harmonic Mean = $\frac{2xy}{x + y}$

Harmonic Mean



Arithmetic Mean = $\frac{x + y}{2}$

Harmonic Mean = $\frac{2xy}{x + y}$


Precision: 1

Recall: 0

Average = 0.5

Harmonic Mean = 0

Harmonic Mean



Precision: 1

Recall: 0

Average = 0.5

Harmonic Mean = 0


Precision: 0.2

Recall: 0.8

Average = 0.5

Harmonic Mean = 0.32

Harmonic Mean



Precision: 1

Recall: 0

Average = 0.5

Harmonic Mean = 0

Precision: 0.2

Recall: 0.8

Average = 0.5

Harmonic Mean = 0.32

F1 Score = Harmonic Mean (Precision, Recall)

F1 Score



Medical Model

Precision: 55.7%

Recall: 83.3%

Average = 69.5%

F1 Score = 66.8%

F1 Score



Spam Detector


Precision: 76.9%

Recall: 37%

Average = 56.9%

F1 Score = 50.0%

F1 Score



Precision: 75%

Recall: 85.7%

Average = 80.3%

F1 Score = 80%

F_β Score

F_β Score



Precision



Recall

F_β Score



Precision

F0.5 Score

F1 Score

F2 Score



Recall

F_β Score



Precision

F0.5 Score



F1 Score

F2 Score



Recall

F_β Score



Precision

F0.5 Score

F1 Score

F2 Score

F10 Score



Recall

F_β Score



Precision

F0.5 Score

F1 Score

F2 Score

F10 Score



Recall

F_β Score

F1 Score = Harmonic Mean (Precision, Recall)

F_β Score

F1 Score = Harmonic Mean (Precision, Recall)

$$H = \frac{n}{\frac{1}{x_1} + \frac{1}{x_2} + \dots + \frac{1}{x_n}}$$

F_β Score

F1 Score = Harmonic Mean (Precision, Recall)

$$H = \frac{n}{\frac{1}{x_1} + \frac{1}{x_2} + \dots + \frac{1}{x_n}}$$


$$F_1 = 2 \frac{1}{\frac{1}{\text{recall}} + \frac{1}{\text{precision}}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

F_β Score


$$F_1 = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

$$F_\beta = (1 + \beta^2) \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

Type I Error (false positive)



Type II Error (false negative)



References

— — —

- https://en.wikipedia.org/wiki/Precision_and_recall
- https://en.wikipedia.org/wiki/Binary_classification
- https://en.wikipedia.org/wiki/F1_score
- <https://www.quora.com/What-is-an-intuitive-explanation-of-F-score>
- “Approximate Statistical Tests for Comparing Supervised Classification Learning Algorithms”, Neural Computation, p. 1895-1923, 1998
<https://www.mitpressjournals.org/doi/10.1162/089976698300017197>

Machine Learning Courses

- Luis Serrano: <https://www.youtube.com/watch?v=aDW44NPhNw0>