

# Deep Neural Networks

## Machine Learning and Pattern Recognition

(Largely based on slides from Fei-Fei Li & Justin Johnson & Serena Yeung)

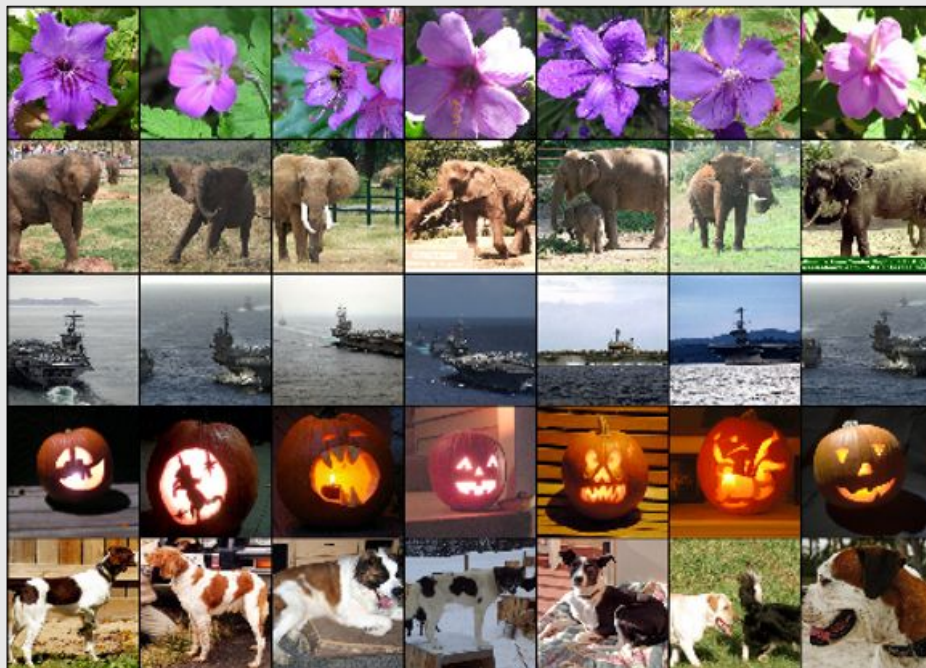
**Prof. Sandra Avila**  
Institute of Computing (IC/Unicamp)

MC886/MO444, October 2, 2018

# DNNs are everywhere ...

## Classification

## Retrieval



Credit: Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012

# DNNs are everywhere ...

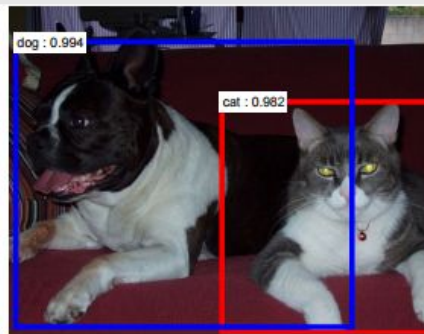

Classification: ElsaGate vs. Safe





# DNNs are everywhere ...

## Detection




## Segmentation



Credit: Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, 2015. Clement Farabet, 2012.

# DNNs are everywhere ...



Detection

Andreza Santos, Thiago Teixeira Santos, Sandra Avila. 2018.


<https://youtu.be/YgZbTca1hl8>

# DNNs are everywhere ...

## Pose Estimation



## Playing Games



Credit: Toshev & Szegedy 2014. Xiaoxiao Guo, Satinder Singh, Honglak Lee, Richard Lewis, and Xiaoshi Wang, 2014.



# DNNs are everywhere ...

No errors



A white teddy bear sitting in the grass

Minor errors



A man in baseball uniform throwing a ball

Somewhat related



A woman is holding a cat in her hand

## Image Captioning



A man riding a wave on top of a surfboard



A cat sitting on a suitcase on the floor



A woman standing on a beach holding a surfboard

Captions generated by Justin Johnson using Neuraltalk.

# DNNs are everywhere ...




## Image Style Transfer

Gatys et al, "Image Style Transfer using Convolutional Neural Networks", CVPR 2016

Gatys et al, "Controlling Perceptual Factors in Neural Style Transfer", CVPR 2017



# DNNs are everywhere ...



## Image Colorization

Zhang et al., "Colorful Image Colorization", ECCV 2016  
<https://demos.algorithmia.com/colorize-photos/>

# DNNs are everywhere ...

*Proof.* Omitted. □

**Lemma 0.1.** *Let  $\mathcal{C}$  be a set of the construction.*

*Let  $\mathcal{C}$  be a gerber covering. Let  $\mathcal{F}$  be a quasi-coherent sheaves of  $\mathcal{O}$ -modules. We have to show that*

$$\mathcal{O}_{\mathcal{O}_X} = \mathcal{O}_X(\mathcal{L})$$

*Proof.* This is an algebraic space with the composition of sheaves  $\mathcal{F}$  on  $X_{\acute{e}tale}$  we have

$$\mathcal{O}_X(\mathcal{F}) = \{morph_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F})\}$$

where  $\mathcal{G}$  defines an isomorphism  $\mathcal{F} \rightarrow \mathcal{F}$  of  $\mathcal{O}$ -modules. □

**Lemma 0.2.** *This is an integer  $\mathbb{Z}$  is injective.*

*Proof.* See Spaces, Lemma ?? □

**Lemma 0.3.** *Let  $S$  be a scheme. Let  $X$  be a scheme and  $X$  is an affine open covering. Let  $\mathcal{U} \subset \mathcal{X}$  be a canonical and locally of finite type. Let  $X$  be a scheme. Let  $X$  be a scheme which is equal to the formal complex.*

*The following to the construction of the lemma follows.*

*Let  $X$  be a scheme. Let  $X$  be a scheme covering. Let*

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$


*be a morphism of algebraic spaces over  $S$  and  $Y$ .*

*Proof.* Let  $X$  be a nonzero scheme of  $X$ . Let  $X$  be an algebraic space. Let  $\mathcal{F}$  be a quasi-coherent sheaf of  $\mathcal{O}_X$ -modules. The following are equivalent

- (1)  $\mathcal{F}$  is an algebraic space over  $S$ .
- (2) If  $X$  is an affine open covering.

Consider a common structure on  $X$  and  $X$  the functor  $\mathcal{O}_X(U)$  which is locally of finite type. □

This since  $\mathcal{F} \in \mathcal{F}$  and  $x \in \mathcal{G}$  the diagram



is a limit. Then  $\mathcal{G}$  is a finite type and assume  $S$  is a flat and  $\mathcal{F}$  and  $\mathcal{G}$  is a finite type  $f_*$ . This is of finite type diagrams, and

- the composition of  $\mathcal{G}$  is a regular sequence,
- $\mathcal{O}_{X'}$  is a sheaf of rings.

□

*Proof.* We have see that  $X = \text{Spec}(R)$  and  $\mathcal{F}$  is a finite type representable by algebraic space. The property  $\mathcal{F}$  is a finite morphism of algebraic stacks. Then the cohomology of  $X$  is an open neighbourhood of  $U$ . □

*Proof.* This is clear that  $\mathcal{G}$  is a finite presentation, see Lemmas ??.

A reduced above we conclude that  $U$  is an open covering of  $\mathcal{C}$ . The functor  $\mathcal{F}$  is a "field

$$\mathcal{O}_{X,x} \longrightarrow \mathcal{F}_x \xrightarrow{-1(\mathcal{O}_{X_{\acute{e}tale})}} \mathcal{O}_{X_d}^{-1} \mathcal{O}_{X_\lambda}(\mathcal{O}_{X_d}^{\mathcal{P}})$$

is an isomorphism of covering of  $\mathcal{O}_{X_d}$ . If  $\mathcal{F}$  is the unique element of  $\mathcal{F}$  such that  $X$  is an isomorphism.

The property  $\mathcal{F}$  is a disjoint union of Proposition ?? and we can filtered set of presentations of a scheme  $\mathcal{O}_X$ -algebra with  $\mathcal{F}$  are opens of finite type over  $S$ .

If  $\mathcal{F}$  is a scheme theoretic image points. □

If  $\mathcal{F}$  is a finite direct sum  $\mathcal{O}_{X_\lambda}$  is a closed immersion, see Lemma ?? . This is a sequence of  $\mathcal{F}$  is a similar morphism.

## Text Generation

# DNNs are everywhere ...

Proof. Omitted.

**Lemma 0.1.** Let  $\mathcal{C}$  be a set of  
Let  $\mathcal{C}$  be a gerber covering.  $\mathcal{C}$   
have to show that

*Proof.* This is an algebraic space  
have

$$\mathcal{O}_X(\mathcal{F}) = \{ \text{morph}_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F}) \}$$

where  $\mathcal{G}$  defines an isomorphism  $\mathcal{F} \rightarrow \mathcal{F}$  of  $\mathcal{O}_X$ -modules.

**Lemma 0.2.** This is an integer  $\mathbb{Z}$  is injective.

*Proof.* See Spaces, Lemma ??.

**Lemma 0.3.** Let  $S$  be a scheme. Let  $X$  be a scheme and  $X$  is an affine open  
covering. Let  $\mathcal{U} \subset \mathcal{X}$  be a canonical and locally of finite type. Let  $X$  be a scheme.  
Let  $X$  be a scheme which is equal to the formal complex.

The following to the construction of the lemma follows.

Let  $X$  be a scheme. Let  $X$  be a scheme covering. Let

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over  $S$  and  $Y$ .

*Proof.* Let  $X$  be a nonzero scheme of  $X$ . Let  $X$  be an algebraic space. Let  $\mathcal{F}$  be a  
quasi-coherent sheaf of  $\mathcal{O}_X$ -modules. The following are equivalent

- (1)  $\mathcal{F}$  is an algebraic space over  $S$ .
- (2) If  $X$  is an affine open covering.

Consider a common structure on  $X$  and  $X$  the functor  $\mathcal{O}_X(U)$  which is locally of  
finite type.

This since  $\mathcal{F} \in \mathcal{F}$  and  $x \in \mathcal{G}$  the diagram

$$\begin{array}{ccc} & & \mathcal{O}_{X'} \\ & \nearrow & \downarrow \\ & \mathcal{O}_{X'} & \\ \alpha' \longrightarrow & & \alpha \\ \downarrow & & \downarrow \\ \alpha' \longrightarrow & & \alpha \end{array} \quad \begin{array}{c} X \\ \downarrow \\ \text{MorSets} \\ \downarrow \\ d(\mathcal{O}_{X/k}, \mathcal{G}) \end{array}$$

is a limit. Then  $\mathcal{G}$  is a finite type and assume  $S$  is a flat and  $\mathcal{F}$  and  $\mathcal{G}$  is a finite  
type  $f_*$ . This is of finite type diagrams, and

- the composition of  $\mathcal{G}$  is a regular sequence,
- $\mathcal{O}_{X'}$  is a sheaf of rings.

*Proof.* We have see that  $X = \text{Spec}(R)$  and  $\mathcal{F}$  is a finite type representable by  
algebraic space. The property  $\mathcal{F}$  is a finite morphism of algebraic stacks. Then the  
cohomology of  $X$  is an open neighbourhood of  $U$ .

*Proof.* This is clear that  $\mathcal{G}$  is a finite presentation, see Lemmas ??.

A reduced above we conclude that  $U$  is an open covering of  $\mathcal{C}$ . The functor  $\mathcal{F}$  is a  
“field

$$\mathcal{O}_{X,x} \longrightarrow \mathcal{F}_x^{-1}(\mathcal{O}_{X_{\text{étale}}}) \longrightarrow \mathcal{O}_{X_d}^{-1}(\mathcal{O}_{X_\lambda}(\mathcal{O}_{X_d}^{\mathcal{F}}))$$

is an isomorphism of covering of  $\mathcal{O}_{X_d}$ . If  $\mathcal{F}$  is the unique element of  $\mathcal{F}$  such that  $X$   
is an isomorphism.

The property  $\mathcal{F}$  is a disjoint union of Proposition ?? and we can filtered set of  
presentations of a scheme  $\mathcal{O}_X$ -algebra with  $\mathcal{F}$  are opens of finite type over  $S$ .

If  $\mathcal{F}$  is a scheme theoretic image points.

If  $\mathcal{F}$  is a finite direct sum  $\mathcal{O}_{X_\lambda}$  is a closed immersion, see Lemma ??.

Text  
Generation



# Training: “Maior dúvida da aula” 27/october/2017

[https://github.com/II Sourcell/recurrent\\_neural\\_network](https://github.com/II Sourcell/recurrent_neural_network)

## ##### GoogLeNet, Inception Module

Não entendi muito bem sobre as inception layers na GoogLeNet. Entendi a ideia de fazer a mesma coisa de um filtro grande com vários filtros menores. Com vários filtros menores temos menos parâmetros que um filtro grande?

Quando fazemos inception e concatenamos os resultados, podemos comparar isso à criação de vetor de características? Porque estamos retirando tipos diferentes de informações de uma mesma camada de input e juntando elas pra formar um output.

Acho que não consegui entender muito bem o inception module da arquitetura GoogLeNet. Para que ele serve exatamente? Obrigada.

no modelo de inception v4, usa a paralelizacao para obter menos parametros, entao isso quer dizer que enquanto menos parametros e mais profundo da melhores resultados?

Não entendi exatamente que fator possibilitou a remoção das camadas fully connected na GoogLeNet. Pelo que eu entendi, as redes mais modernas voltaram com a camada fully connected. Então quando usá-la ou não usá-la?

## ## Números de parâmetros

Em relação a arquitetura proposta na rede GoogLeNet, não ficou muito claro para mim as camadas internas, principalmente na parte em que aplicar vários filtros menores, equilibra a aplicar um filtro maior (embora o resultado não seja o mesmo).

Não ficou claro para mim qual a vantagem de se utilizar, por exemplo, 3 pequenos filtros 3x3 ao invés de um 7x7. Na aula você comentou que é para evitar diminuir drasticamente a imagem, mas qual a desvantagem disso?

Eu não entendi aquelas contas dos filtros que reduziam o número de parâmetros

## ##### ResNet Filtro 1x1

Achei um pouco confuso as dimensões do filtro 1x1. Achei confuso a parte da convolução de tal filtro.

# Training: “Maior dúvida da aula” 27/october/2017

```
iter 0, loss: 107.601633
```

```
----
```

```
'ōqIE:ō:3(é
```

```
0 Q.L"cÉhíL'uàfM0)êoâz.àãâéláč-)D(iéêdàF(1LF LRcFA0nC(Pô(á#HM5éI?#ázHrtGTRF)5wlGaúa2éj?pd7,u  
xp5LQ"r24F7é1efL"CabvêúhyLdã 7àã2à0bm xv?qnAodí'P)mTg4(u4F7ú13ómrQnmeFNbãoúvâ3i?xsuRăjáécó.-  
záy-
```

```
----
```

# Training: “Maior dúvida da aula” 27/october/2017

```
iter 0, loss: 107.601633
```

```
----
```

```
'ōqIE:ō:3(é  
0 Q.L"cÉhíL'uàfM0)êoâz.àãâéláč-)D(iéêdàF(1LF LRcFA0nC(Pô(á#HM5éI?#ázHrtGTRF)5wlGaúa2éj?pd7,u  
xp5LQ"r24F7é1efL"CabvêúhyLdã 7ãã2à0bm xv?qnAodí'P)mTg4(u4F7ú13ómrQnmeFNbãoúvâ3i?sx suRãjáécó.-  
záy
```

```
--- iter 46000, loss: 23.238596
```

```
-----  
és GoogLeNet. E a rede aprende?
```

0 Daras dúvrvilg. ( ende no pré-tro "rar outlara destidas? Com uttres dessar algo us filtros parte novados aplicar au mula.

e narepteno Retênne camada entros lemos e m



# Training: “Maior dúvida da aula” 27/october/2017

```
iter 0, loss: 107.601633
```

```
-----  
'õqIE:õ:3(é  
0 Q.L"çÉhíL'uàfM0)êoâz.ããâélác-)D(iéêdàF(1LFLrRcFA0nC(Pô(á#HM5éI?#ázHrtGTRF)5wlGaúa2éj?pd7,u  
xp5LQ"r24F7élefl"CabvêúhyLdã 7ãã2à0bm xv?qnAodí'P)mTg4(u4F7ú13ómrQnmeFNbãoúvâ3i?sx suRãjáécó.-  
Záy-----
```

```
--- iter 46000, loss: 23.238596
```

```
-----  
és GoogLeNet. E a rede aprende?
```

```
0 Daras dúvrvilg. ( ende no pré-tro "rar outlara destinadas? Com uttres dessar algo us filtros  
parte novados aplicar au mula.
```

```
e nar iter 204000, loss: 10.733449
```

```
-----  
to, ina utir alpal asvelum motrio tarada mexexexterna mai reviso de enter meiss grandas
```

```
##### ResNet Filtro 1x1? Alheing?
```

```
Não entendi exatamente que fia, confenhalo deset desecta..
```

```
##### Como as
```

# DNNs are everywhere ...



Google AI

Google Duplex



Hi, how can I help?

Google Assistant



AI Assistant Calls:  
Google Duplex

<https://www.youtube.com/watch?v=WPzu6W2rWNs>

# DNNs are everywhere ...

## Synthesizing Audio



Without Re-timing



With Re-timing  
(Our Result)

Suwajanakorn et al., "Synthesizing Obama  
Learning Lip Sync from Audio", SIGGRAPH, 2017


<https://youtu.be/mKxgAnuvaZk>

<https://youtu.be/cQ54GDm1eL0>



# DNNs are everywhere ...


## Synthesizing Skin Lesion




Alceu Bissoto, Fábio Perez, Eduardo Valle, Sandra Avila. "Skin lesion synthesis with generative adversarial networks", ISIC Skin Image Analysis Workshop at MICCAI, 2018.

# Convolutional Neural Networks (CNNs)

# Fully Connected Layer




CIFAR-10




$32 \times 32 \times 3$  image  $\Rightarrow$  stretch to  $3072 \times 1$

# Fully Connected Layer



CIFAR-10




$32 \times 32 \times 3$  image  $\Rightarrow$  stretch to  $3072 \times 1$






# What is a Convolution?


# What is a Convolution?



0	1	0
1	-4	1
0	1	0





# What is a Convolution?




Edge  
Detection

0	1	0
1	-4	1
0	1	0




# What is a Convolution?



Emboss

-2	-1	0
-1	1	1
0	1	2





# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter


# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter




# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4		

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &0*1 + 0*0 + 1*1 = 4 \end{aligned}$$


# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	

$$\begin{aligned} &1*1 + 1*0 + 0*1 + \\ &1*0 + 1*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 3 \end{aligned}$$




# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4

$$\begin{aligned} &1*1 + 0*0 + 0*1 + \\ &1*0 + 1*1 + 0*0 + \\ &1*1 + 1*0 + 1*1 = 4 \end{aligned}$$


# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4
2		

$$\begin{aligned} &0*1 + 1*0 + 1*1 + \\ &0*0 + 0*1 + 1*0 + \\ &0*1 + 0*0 + 1*1 = 2 \end{aligned}$$

# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4
2	4	

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 4 \end{aligned}$$


# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4
2	4	3

$$\begin{aligned} &1*1 + 1*0 + 0*1 + \\ &1*0 + 1*1 + 1*0 + \\ &1*1 + 1*0 + 0*1 = 3 \end{aligned}$$




# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4
2	4	3
2		

$$\begin{aligned} &0*1 + 0*0 + 1*1 + \\ &0*0 + 0*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 2 \end{aligned}$$


# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter



4	3	4
2	4	3
2	3	

$$\begin{aligned} &0*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &1*1 + 1*0 + 0*1 = 3 \end{aligned}$$

# What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5 × 5 matrix

1	0	1
0	1	0
1	0	1

3 × 3 filter




4	3	4
2	4	3
2	3	4

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &1*0 + 1*1 + 0*0 + \\ &1*1 + 0*0 + 0*1 = 4 \end{aligned}$$


# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




**Convolve** the filter with the image i.e.  
“slide over the image spatially,  
computing dot products”

$5 \times 5 \times 3$  filter




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure



**Convolve** the filter with the image i.e.  
“slide over the image spatially,  
computing dot products”

$5 \times 5 \times 3$  filter




**Filters always extend the full  
depth of the input volume**




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




$5 \times 5 \times 3$  filter  $\rightarrow$




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure



# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




1 number:

$5 \times 5 \times 3 = 75$ -dimensional  
dot product + bias)

$$w^T x + b$$

# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




1 number:

$5 \times 5 \times 3 = 75$ -dimensional  
dot product + bias)

$$w^T x + b$$

# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure



1 number:


$5 \times 5 \times 3 = 75$ -dimensional  
dot product + bias)

$$w^T x + b$$


That region in the input image is called the *local receptive field* for the hidden neuron.

# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




Convolve (slide) over  
all spatial locations





# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




Convolve (slide) over  
all spatial locations




$32 \times 32 \times 3$  image  
 **$5 \times 5 \times 3$  filter**

activation map




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




Convolve (slide) over  
all spatial locations



$32 \times 32 \times 3$  image  
 **$5 \times 5 \times 3$  filter**


(considering a second filter)

**activation maps**




# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure



Convolve (slide) over  
all spatial locations




$32 \times 32 \times 3$  image  
 **$5 \times 5 \times 3$  filter**


If we had 6  $5 \times 5 \times 3$  filters ...

# Convolution Layer

$32 \times 32 \times 3$  image  $\Rightarrow$  preserve spatial structure




Convolve (slide) over  
all spatial locations



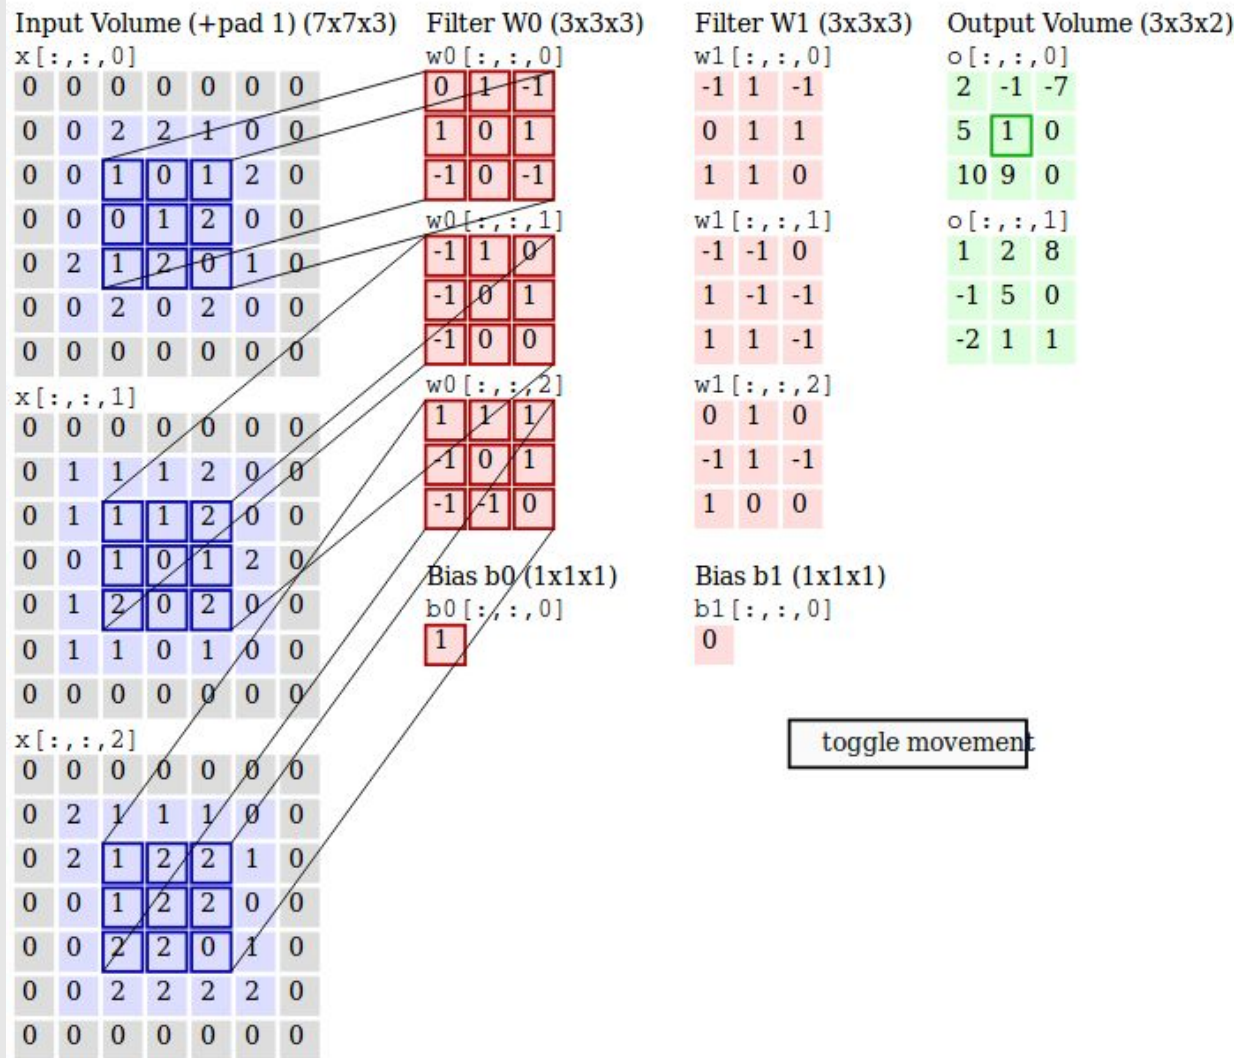
$32 \times 32 \times 3$  image  
 **$5 \times 5 \times 3$  filter**

If we had 6  $5 \times 5 \times 3$  filters ...

**6 activation maps**




<http://cs231n.github.io/convolutional-networks>




# Convolutional Networks

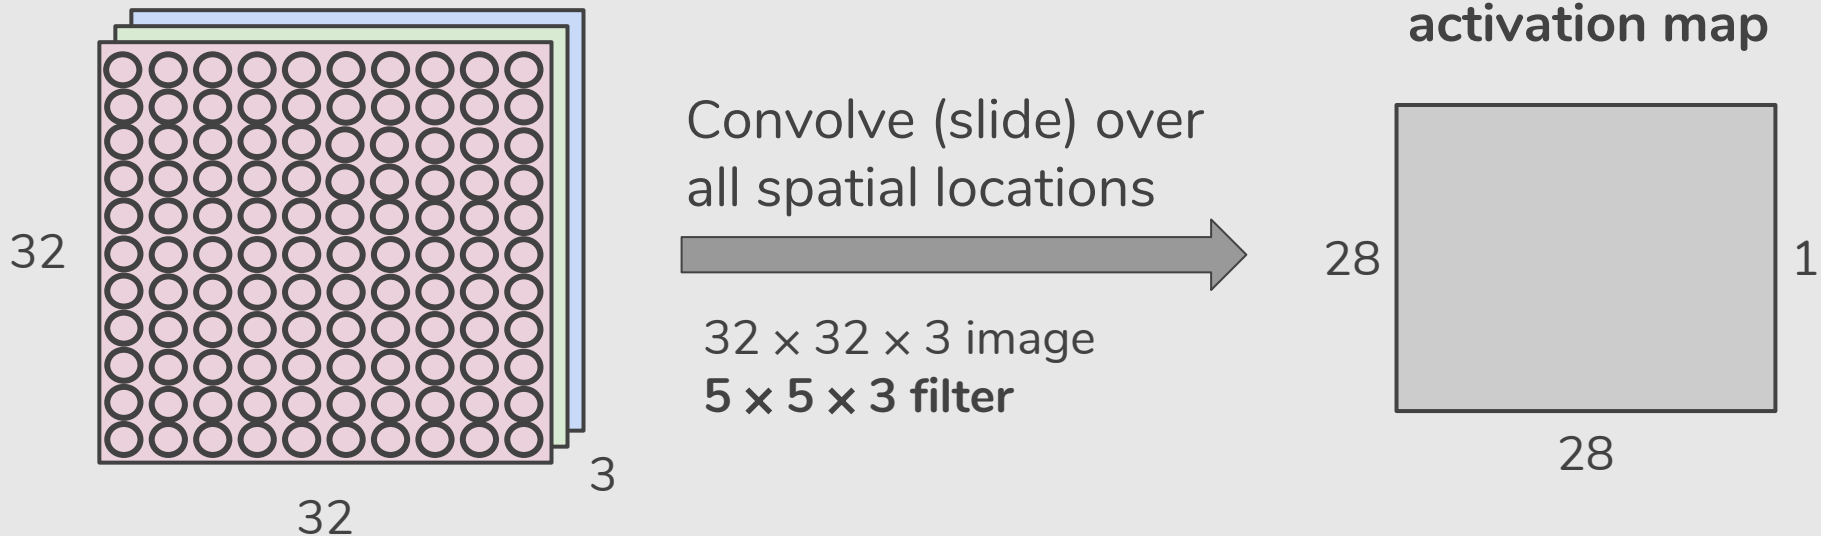
Sequence of Convolutional Layers, interspersed with activation functions.








# A Closer Look at Spatial Dimensions




# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter

# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter

# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter


# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter




# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter

# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter


# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter


$\Rightarrow 5 \times 5$  output

# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 2**

# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 2**

# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 2**

# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 2**

# A Closer Look at Spatial Dimensions




$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 2**

$\Rightarrow 3 \times 3$  output




# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 3**?


# A Closer Look at Spatial Dimensions



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter  
applied with **stride 3**?

**Doesn't fit!**  
**cannot apply  $3 \times 3$  filter on**  
 **$7 \times 7$  input with stride 3.**


# A Closer Look at Spatial Dimensions



Output size:

$$(N - F) / \text{stride} + 1$$

# A Closer Look at Spatial Dimensions




Output size:

$$(N - F) / \text{stride} + 1$$

e.g.  $N = 7, F = 3$ :

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

# A Closer Look at Spatial Dimensions



Output size:


$$(N - F) / \text{stride} + 1$$

e.g.  $N = 7, F = 3$ :

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

# A Closer Look at Spatial Dimensions



Output size:

$$(N - F) / \text{stride} + 1$$

e.g.  $N = 7, F = 3$ :

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3) / 3 + 1 = 2.33$$



## In Practice: Common to zero pad the border

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

$7 \times 7$  input,  
 $3 \times 3$  filter applied  
with **stride 1** with **pad 1**

What is the output?



## In Practice: Common to zero pad the border

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

$7 \times 7$  input,  
 $3 \times 3$  filter applied  
with **stride 1** with **pad 1**

What is the output?  
 **$7 \times 7$  output**

## In Practice: Common to zero pad the border

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

In general, common to see CONV layers with stride 1, filters of size  $F \times F$ , and zero-padding with  $(F-1)/2$  (will preserve size spatially).


e.g.  $F = 3 \Rightarrow$  zero pad with 1

$F = 5 \Rightarrow$  zero pad with 2

$F = 7 \Rightarrow$  zero pad with 3

Shrinking too fast is not good, doesn't work well.

$32 \rightarrow 28 \rightarrow 24 \rightarrow \dots$



# Number of Parameters

Input volume:  $32 \times 32 \times 3$

10  $5 \times 5$  filters with stride 1, pad 2

Number of parameters in this layer?

# Number of Parameters


Input volume:  $32 \times 32 \times 3$

$10 \times 5 \times 5$  filters with stride 1, pad 2

Number of parameters in this layer?

Each filter has  $5 \times 5 \times 3 + 1 = 76$  parameters (+1 for bias)

$\Rightarrow 76 \times 10 = 760$



# Pooling Layer


- Makes the representations smaller and more manageable
- Operates over each activation map independently

# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with  $2 \times 2$   
filters and stride 2






# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with  $2 \times 2$   
filters and stride 2




6	

# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with  $2 \times 2$   
filters and stride 2




6	8

# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4


Max pooling with  $2 \times 2$   
filters and stride 2



6	8
3	


# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently




# Fully Connected Layer

- Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



http://neuralnetworksanddeeplearning.com/chap6.html#final\_conv



## Convolutional neural networks in practice

---

We've now seen the core ideas behind convolutional neural networks. Let's look at how they work in practice, by implementing some convolutional networks, and applying them to the MNIST digit classification problem. The program we'll use to do this is called `network3.py`, and it's an improved version of the programs `network.py` and `network2.py` developed in earlier chapters\*. If you wish to follow along, the code is available [on GitHub](#). Note that we'll work through the code for `network3.py` itself in the next section. In this section, we'll use `network3.py` as a library to build convolutional networks.

\*Note also that `network3.py` incorporates ideas from the Theano library's documentation on convolutional neural nets (notably the implementation of LeNet-5), from Misha Denil's [implementation of dropout](#), and from [Chris Olah](#).

# CNNs Architectures

# CNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)



INPUT 32x32

C1: feature maps 6@28x28

C3: f. maps 16@10x10

S2: f. maps 6@14x14

S4: f. maps 16@5x5

C5: layer 120

F6: layer 84

OUTPUT 10

Convolutions

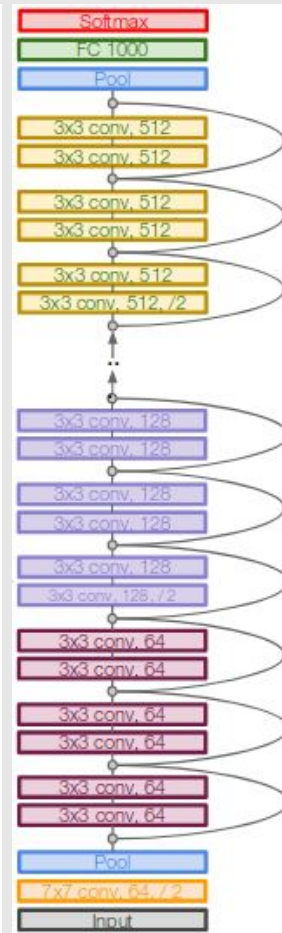



Subsampling

Convolutions

Subsampling

Full connection

Full connection



# References

— — —

## Machine Learning Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chap. 11 & 13

## Machine Learning Courses

- <https://www.coursera.org/learn/neural-networks>
- “The 3 popular courses on Deep Learning”:  
<https://medium.com/towards-data-science/the-3-popular-courses-for-deeplearning-ai-ac37d4433bd>