Machine Learning Tools

Edson Riberto Bollis





Index

- Deep learning hardware
 - CPU, GPU, TPU
- Deep learning software
 - PyTorch and TensorFlow
 - Static vs Dynamic computation graphs





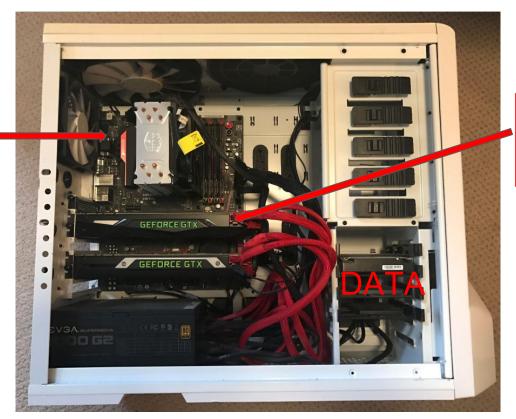








CPU





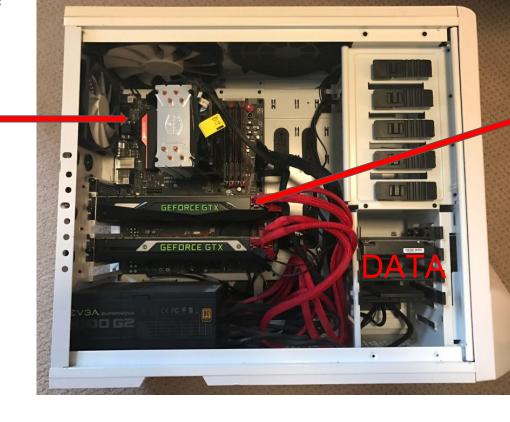
GPU







CPU





GPU

- Read all data into RAM
- Use SSD instead of HDD
- Use multiple CPU threads to prefetch data





CPU x GPU

	Cores	Clock Speed	Memory	Price	Speed
CPU (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.2 GHz	System RAM	\$339	~540 GFLOPs FP32
GPU (NVIDIA GTX 1080 Ti)	3584	1.6 GHz	11 GB GDDR5 X	\$699	~11.4 TFLOPs FP32





NVIDIA vs AMD





HIP: converts CUDA code to run AMD GPUs

NVIDIA vs AMD

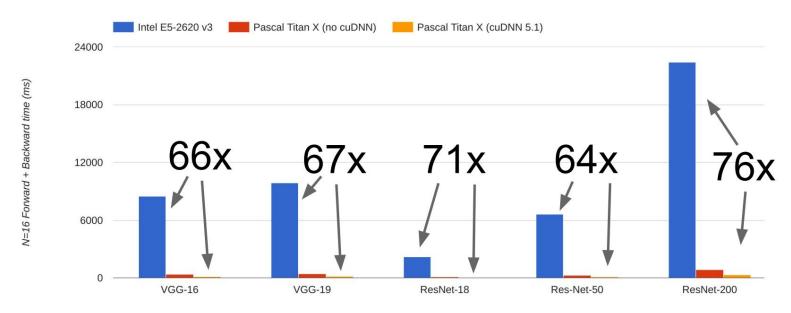
CUDA, cuDNN,cuBLAS, cuFFT, etc OpenCL (slower CUDA)





CPU vs GPU in practice

(CPU performance not well-optimized, a little unfair)



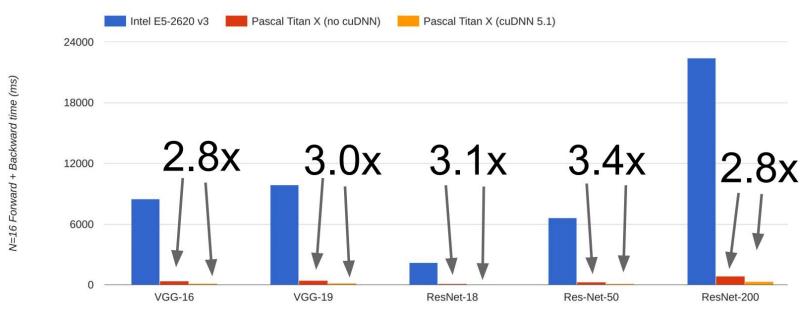
Data from https://github.com/jcjohnson/cnn-benchmarks





CPU vs GPU in practice

cuDNN much faster than "unoptimized" CUDA



Data from https://github.com/jcjohnson/cnn-benchmarks





CPU x GPU x TPU

CPU: Fewer cores, but each core is much faster and much more capable; great at sequential tasks;

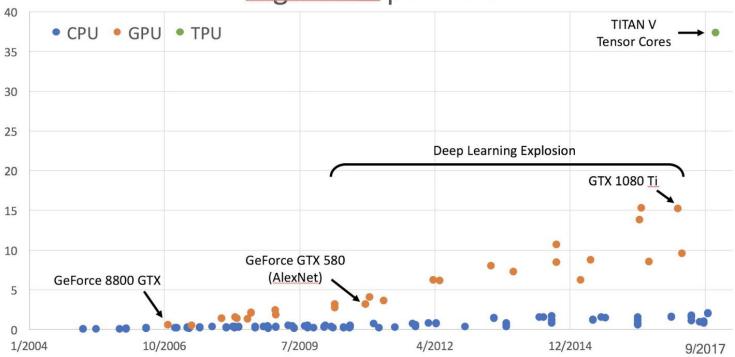
GPU: More cores, but each core is much slower and "dumber"; great for parallel tasks 14;

TPU: Specialized hardware for deep learning;





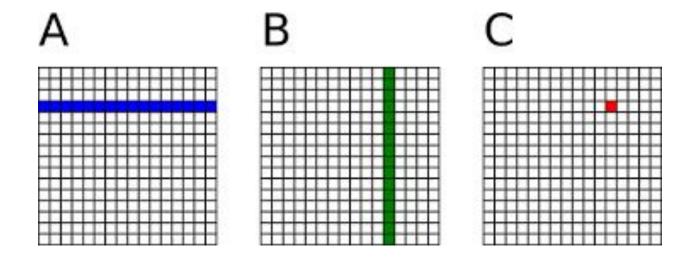
GigaFLOPs per Dollar







Example: Matrix Multiplication







Deep Learning Software





Frameworks

Caffe Caffe2
(UC Berkeley) Facebook)

Torch (NYU / Facebook) — PyTorch (Facebook)

Theano _____ TensorFlow (Google)

PaddlePaddle Chainer (Baidu)

MXNet
(Amazon)
Developed by U Washington, CMU, MIT,
Hong Kong U, etc but main framework of

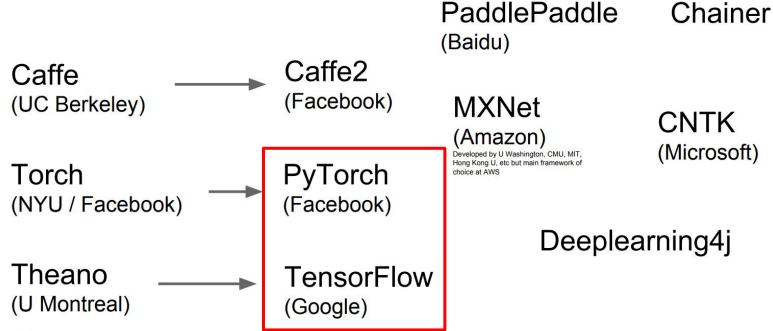
CNTK (Microsoft)

Deeplearning4j





Frameworks







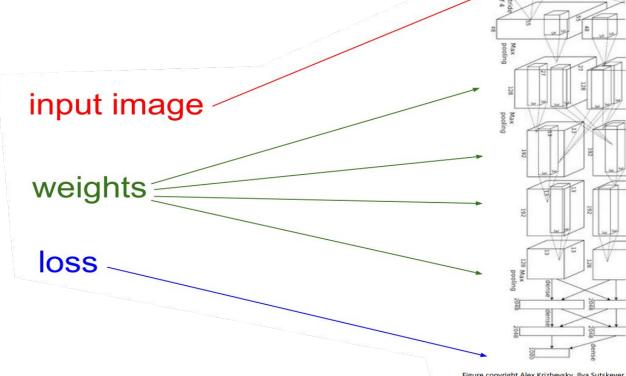
Main idea of deep learning frameworks

- Quick to develop and test new ideas;
- Automatically compute gradients;
- Run it all efficiently on GPU (optimized cuDNN, cuBLAS, etc).





Graphs









Graphs

input image

loss

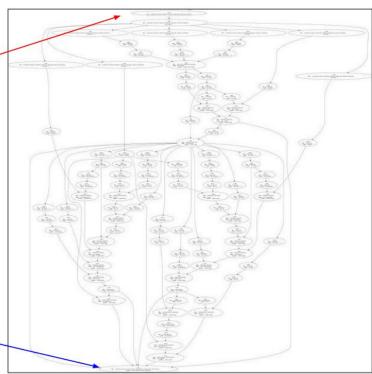


Figure reproduced with permission from a Twitter post by Andrej Karpathy.

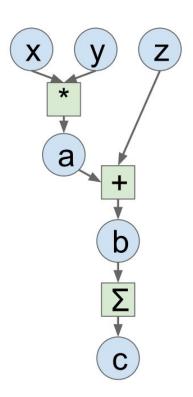




Graphs Numpy

UNICAMP

```
import numpy as np
np.random.seed(0)
N, D = 3, 4
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
a = x * y
c = np.sum(b)
grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```

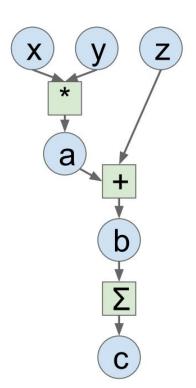




Graphs Numpy

UNICAMP

```
import numpy as np
np.random.seed(0)
N, D = 3, 4
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
a = x * y
b = a + z
c = np.sum(b)
grad c = 1.0
grad b = grad c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad y = grad a * x
```



PyTorch

```
import torch

N, D = 3, 4
x = torch.randn(N, D,
y = torch.randn(N, D)
z = torch.randn(N, D)

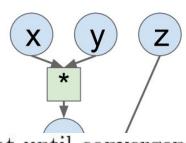
a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```



Graphs Numpy

```
import numpy as np
np.random.seed(0)
N, D = 3, 4
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
a = x * y
b = a + z
c = np.sum(b)
grad c = 1.0
grad b = grad c * np.ones((N, D))
grad_a = grad_b.copy()
grad z = grad b.copy()
grad x = grad a * y
grad y = grad a * x
```



PyTorch

Repeat until convergence {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

0 = 3, 4torch.randn(N, D, requires_grad=True) torch.randn(N, D) torch.randn(N, D)

x * y a + z torch.sum(b)

c.backward() print(x.grad)

import torch







- Tensor: Like a numpy array, but can run on GPU

- Module: A neural network layer; may store state or learnable weights

 Autograd: Package for building computational graphs out of Tensors, and automatically computing gradients





```
import torch
device = torch.device('cpu')
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in, device=device)
y = torch.randn(N, D out, device=device)
w1 = torch.randn(D in, H, device=device)
w2 = torch.randn(H, D out, device=device)
learning rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h relu = h.clamp(min=0)
    y pred = h relu.mm(w2)
    loss = (y pred - y).pow(2).sum()
    grad y pred = 2.0 * (y pred - y)
    grad w2 = h relu.t().mm(grad y pred)
    grad h relu = grad y pred.mm(w2.t())
    grad h = grad h relu.clone()
    grad h[h < 0] = 0
    grad w1 = x.t().mm(grad h)
    w1 -= learning rate * grad w1
    w2 -= learning rate * grad w2
```





```
import torch
device = torch.device('cpu')
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in, device=device)
y = torch.randn(N, D_out, device=device)
w1 = torch.randn(D_in, H, device=device)
w2 = torch.randn(H, D_out, device=device)
```

Tensor

```
learning rate = 1e-6
for t in range(500):
   h = x.mm(w1)
   h relu = h.clamp(min=0)
   y pred = h relu.mm(w2)
   loss = (y pred - y).pow(2).sum()
   grad y pred = 2.0 * (y pred - y)
    grad w2 = h relu.t().mm(grad y pred)
    grad h relu = grad y pred.mm(w2.t())
    grad h = grad h relu.clone()
    grad h[h < 0] = 0
    grad w1 = x.t().mm(grad h)
   w1 -= learning rate * grad w1
   w2 -= learning rate * grad w2
```





```
import torch
device = torch.device('cpu')
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in, device=device)
y = torch.randn(N, D_out, device=device)
w1 = torch.randn(D_in, H, device=device)
w2 = torch.randn(H, D_out, device=device)
```

Forward pass: compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2</pre>
```





Tensor

```
import torch
device = torch.device('cpu')
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in, device=device)
y = torch.randn(N, D_out, device=device)
w1 = torch.randn(D_in, H, device=device)
w2 = torch.randn(H, D_out, device=device)
```

Tensor

Forward pass: compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)</pre>
```

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```

Backward pass: manually compute gradients





```
import torch
device = torch.device('cpu')
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in, device=device)
y = torch.randn(N, D_out, device=device)
w1 = torch.randn(D_in, H, device=device)
w2 = torch.randn(H, D_out, device=device)
```

Tensor

Forward pass: compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)</pre>
```

Backward pass: manually compute gradients



w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2

Gradient descent



```
import torch
```

device = torch.device('cpu') use GPU

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in, device=device)
y = torch.randn(N, D_out, device=device)
w1 = torch.randn(D_in, H, device=device)
w2 = torch.randn(H, D_out, device=device)
```

Tensor

Forward pass: compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)</pre>
```

Backward pass: manually compute gradients



Gradient descent

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```



```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
w1 = torch.randn(D in, H, requires_grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y pred - y).pow(2).sum()
    loss.backward()
    with torch.no grad():
        w1 -= learning rate * w1.grad
        w2 -= learning rate * w2.grad
        wl.grad.zero ()
        w2.grad.zero ()
```





PyTorch PyTorch

```
import torch
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
  = torch.randn(N, D in)
 = torch.randn(N, D out)
w1 = torch.randn(D in, H, requires_grad=True)
w2 = torch.randn(H, D out, requires grad=True)
```

Tensor

```
learning rate = 1e-6
for t in range(500):
```

predictions and loss

```
Forward pass: compute | y pred = x.mm(w1).clamp(min=0).mm(w2)
                    loss = (y pred - y).pow(2).sum()
```

```
loss.backward()
```

Backward pass: automatically compute gradients

```
with torch.no grad():
          wl -= learning rate * wl.grad
          w2 -= learning rate * w2.grad
          wl.grad.zero ()
Gradient
          w2.grad.zero ()
descent
```





```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
```

Forward pass: compute predictions and loss

```
learning_rate = 1e-2
for t in range(500):
    y_pred = model(x)
    loss = torch.nn.functional.mse loss(y pred, y)
```

```
LINICAMP
```

```
loss.backward()
with torch.no_grad():
    for param in model.parameters():
        param -= learning_rate * param.grad
model.zero_grad()
```



Model

```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
model = torch.nn.Sequential(
          torch.nn.Linear(D in, H),
          torch.nn.ReLU(),
          torch.nn.Linear(H, D out))
learning_rate = 1e-2 optimizer = torch.optim.Adam(model.parameters(),
for t in range(500):
                                                lr=learning rate)
    y pred = model(x)
    loss = torch.nn.functional.mse loss(y pred, y)
    loss.backward()
    with torch no grad():
        for param in model parameters():
            param -= learning rate * param.grad
    model.zero grad()
```





```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
model = torch.nn.Sequential(
          torch.nn.Linear(D_in, H),
          torch.nn.ReLU(),
          torch.nn.Linear(H, D out))
learning_rate = 1e-2 optimizer = torch.optim.Adam(model.parameters(),
for t in range(500):
                                                lr=learning rate)
    y pred = model(x)
    loss = torch.nn.functional.mse loss(y pred, y)
    loss.backward()
    optimizer.step()
```



```
model.zero grad()
```



PyTorch: Dynamic Computation Graphs

```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
w1 = torch.randn(D in, H, requires grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y pred - y).pow(2).sum()
    loss.backward()
```





PyTorch: Dynamic Computation Graphs

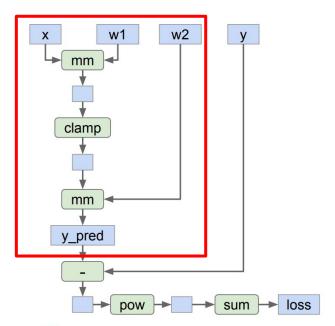


```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D out)
w1 = torch.randn(D in, H, requires grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y pred - y).pow(2).sum()
    loss.backward()
```





PyTorch: Dynamic Computation Graphs

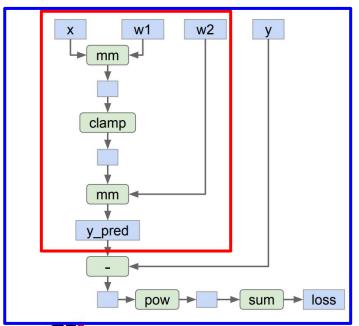


```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y pred - y).pow(2).sum()
    loss.backward()
```





PyTorch: Dynamic Computation Graphs

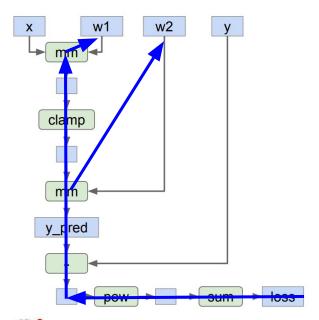


```
import torch
N, D in, H, D out = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()
    loss.backward()
```





PyTorch: Dynamic Computation Graphs

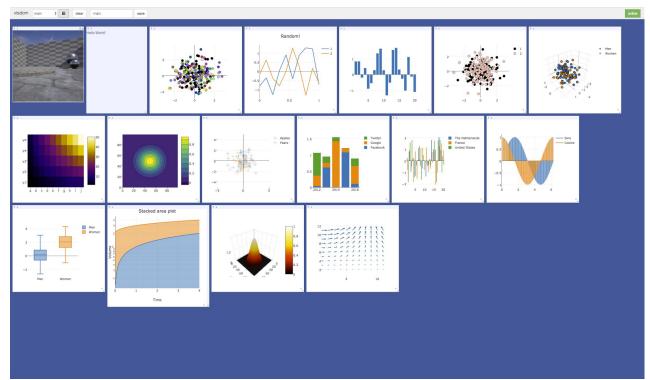


```
import torch
N, D_{in}, H, D_{out} = 64, 1000, 100, 10
x = torch.randn(N, D in)
y = torch.randn(N, D out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D out, requires grad=True)
learning rate = 1e-6
for t in range(500):
    y pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y pred - y).pow(2).sum()
    loss.backward()
```





PyTorch: Visdom



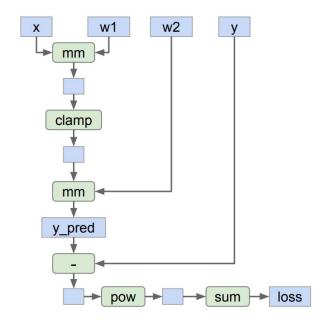




- Static graphs

 Build computational graph describing our computation (including finding paths for backprop)

- Reuse the same graph on every iteration







```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
h = tf.maximum(tf.matmul(x, w1), 0)
y pred = tf.matmul(h, w2)
diff = y pred - y
loss = tf.reduce mean(tf.reduce sum(diff ** 2, axis=1))
grad wl, grad w2 = tf.gradients(loss, [w1, w2])
with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              wl: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad w1, grad w2],
                   feed dict=values)
    loss val, grad wl val, grad w2 val = out
```





```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
h = tf.maximum(tf.matmul(x, wl), 0)
y pred = tf.matmul(h, w2)
diff = y pred - y
loss = tf.reduce mean(tf.reduce sum(diff ** 2, axis=1))
grad wl, grad w2 = tf.gradients(loss, [w1, w2])
with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              wl: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad w1, grad w2],
                   feed dict=values)
    loss val, grad wl val, grad w2 val = out
```

Declaration of Placeholders





```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

Declaration of Placeholders

Forward pass: just build model

```
h = tf.maximum(tf.matmul(x, wl), 0)
y pred = tf.matmul(h, w2)
diff = y pred - y
loss = tf.reduce mean(tf.reduce sum(diff ** 2, axis=1))
grad wl, grad w2 = tf.gradients(loss, [w1, w2])
with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              wl: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad w1, grad w2],
                   feed dict=values)
    loss val, grad wl val, grad w2 val = out
```





```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

Declaration of Placeholders

Forward pass: just build model

```
h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
```

Add info. to calculate Gradient and losses (W1, W2)

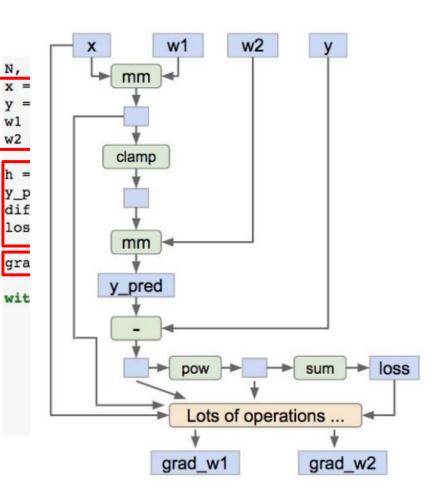
loss_val, grad_wl_val, grad_w2_val = out

grad w1, grad w2 = tf.gradients(loss, [w1, w2])





Forward pass: just build model



Declaration of Placeholders

Add info. to calculate Gradient and loss (W1, W2)





```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

grad w1, grad w2 = tf.gradients(loss, [w1, w2])

Declaration of Placeholders

Forward pass: just build model

```
h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
```

Add info. to calculate Gradient and loss (W1, W2)

Compile model and start the treinment





```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
wl = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

Declaration of Placeholders

Forward pass: just build model

Compile model and start the treinment

```
h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
```

```
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])
```

Add info. to calculate Gradient and loss (W1, W2)

Run





```
N, D, H = 64, 1000, 100

x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
wl = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

Declaration of Placeholders

Forward pass: just build model

Compile model and start the treinment

```
h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
```

grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

```
with tf.Session() as sess:
```

Problem: every step allocate w1, w2 placeholder

Add info. to calculate Gradient and loss (W1, W2)

Run





Forward pass:

just build model

```
N, D, H = 64, 1000, 100
        x = tf.placeholder(tf.float32, shape=(N, D))
                                                             Declaration
        y = tf.placeholder(tf.float32, shape=(N, D))
                                                                 of
        w1 = tf.placeholder(tf.float32, shape=(D, H))
        w2 = tf.placeholder(tf.float32, shape=(H, D))
                                                             Placeholders
        h = tf.maximum(tf.matmul(x, wl), 0)
        y pred = tf.matmul(h, w2)
        diff = y pred - y
w1 = tf.Variable(tf.random_normal((D, H)))
     = tf.Variable(tf.random_normal((H, D))) and loss
                                                                 (W1, W2)
        with tf.Session() as sess:
            values = {x: np.random.randn(N, D),
                     w1: np.random.randn(D, H),
                     w2: np.random.randn(H, D),
                                                             Run
                     y: np.random.randn(N, D),}
         for t in range(50):
             out = sess.run([loss, grad wl, grad w2],
                            feed dict=values)
             loss val, grad wl val, grad w2 val = out
             values[w1] -= learning rate * grad_w1_val
```

values[w2] -= learning rate * grad w2 val

persists in the graph between calls



Compile model

start the treinment

Add assign operations to update w1 and w2 as part of the graph

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random normal((D, H)))
w2 = tf.Variable(tf.random normal((H, D)))
h = tf.maximum(tf.matmul(x, wl), 0)
y pred = tf.matmul(h, w2)
diff = v pred - v
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad wl, grad w2 = tf.gradients(loss, [w1, w2])
learning rate = 1e-5
new w1 = w1.assign(w1 - learning rate * grad w1)
new w2 = w2.assign(w2 - learning rate * grad w2)
with tf.Session() as sess:
                                                    run graph to Init variables
    sess.run(tf.global variables initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    for t in range(50):
train
        loss val, = sess.run([loss], feed dict=values)
```



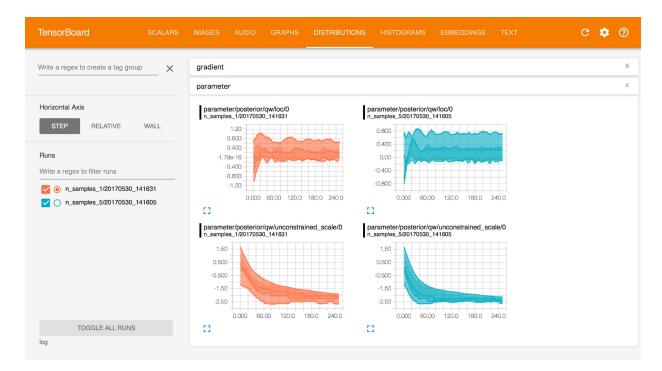
TensorFlow and Keras (TensorFlow, Theano and CNTK)



train



TensorFlow: Tensorboard







TensorFlow: Static x Dinamic :PyTorch

PyTorch
Dynamic Graphs
Static: ONNX, Caffe2

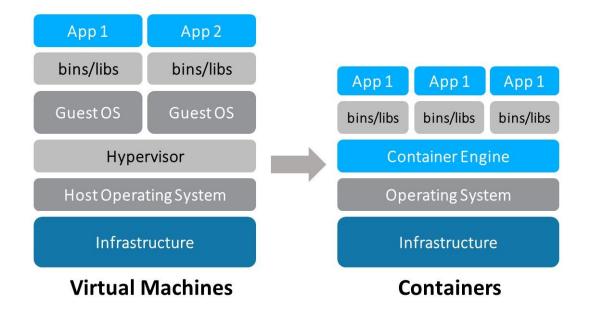
TensorFlow Static Graphs Dynamic: Eager

- Recurrent networks
- Recursive networks
- Modular Networks





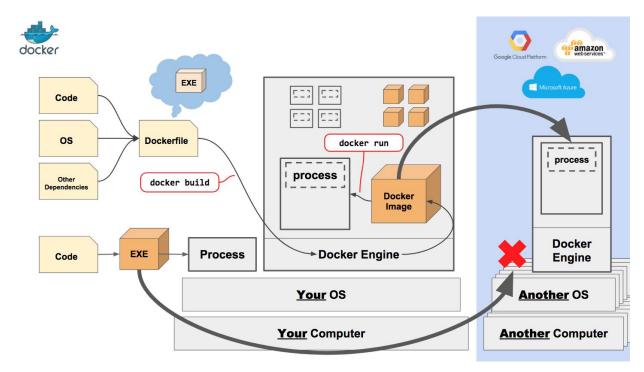
Docker







Docker







Docker

Start CPU only container

\$ docker run -it -p 8888:8888 tensorflow/tensorflow
Go to your browser on http://localhost:8888/

Start GPU (CUDA) container

Install nvidia-docker and run

\$ docker run --runtime=nvidia -it -p 8888:8888 tensorflow/tensorflow:latest-gpu
Go to your browser on http://localhost:8888/

Other versions (like release candidates, nightlies and more)

See the list of tags. Devel docker images include all the necessary dependencies to build from source whereas the other binaries simply have TensorFlow installed.

For more details details see

https://github.com/tensorflow/tensorflow/blob/master/tensorflow/tools/docker/README.md



