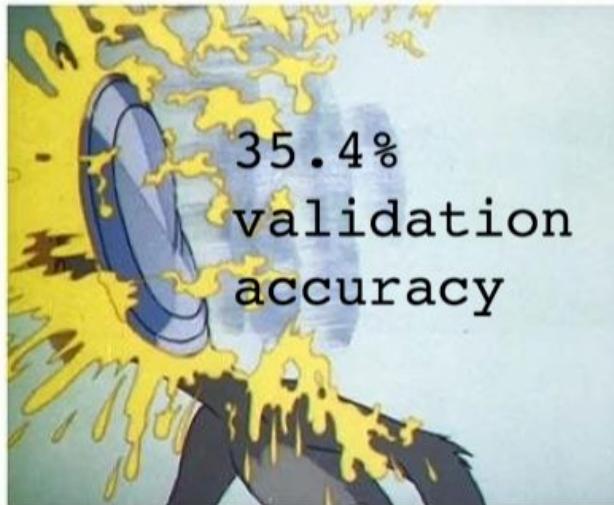


Recall from last time ...

Friends don't let friends
use testing data
for training



@NPCompleteTeens

35.4%
validation
accuracy

Data



Training

Test



Training

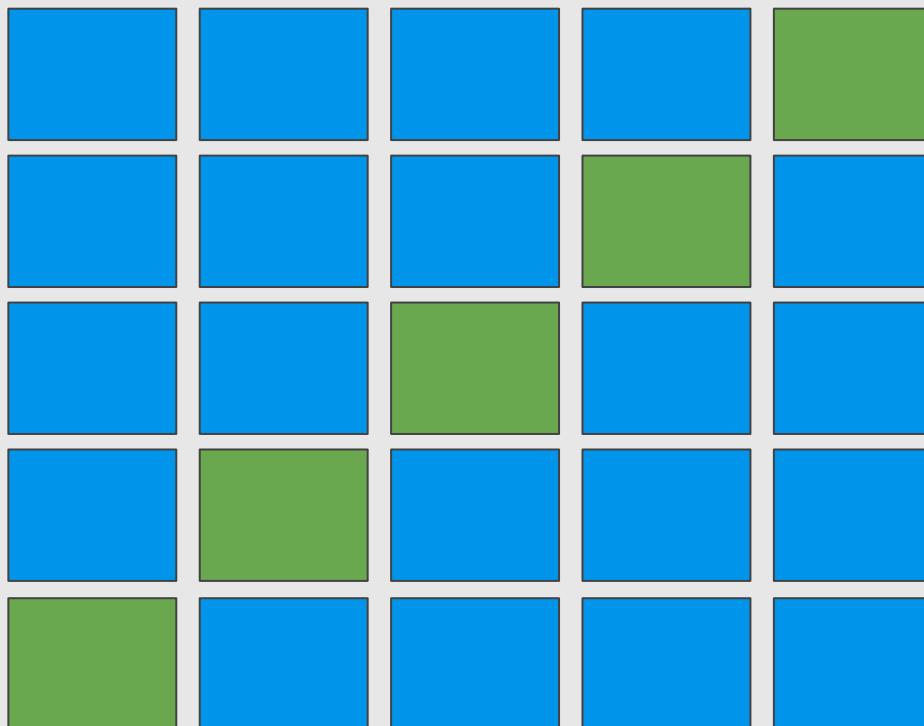
Validation

Test

k-fold Cross Validation



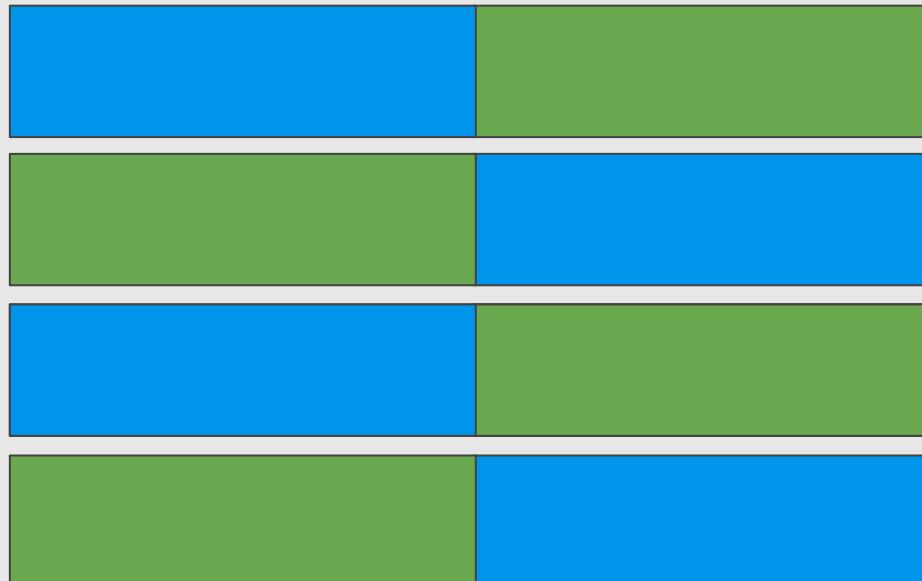
$k = 5$



$k \times 2$ -fold Cross Validation



$k = 5$

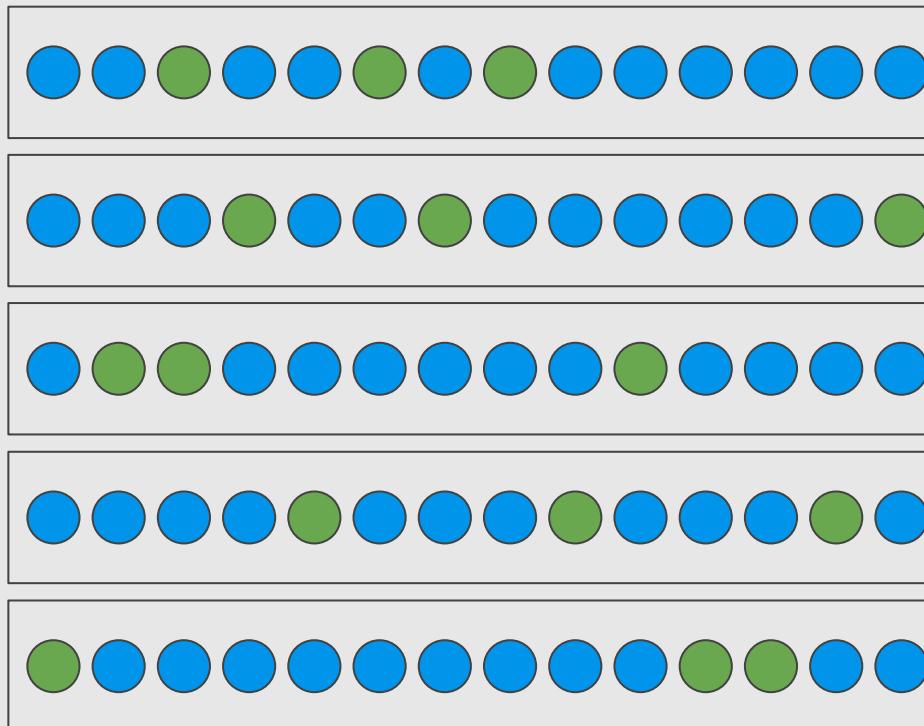


...

k times = $k \times 2$ folds

Randomizing in Cross Validation

● Training
● Validation



Evaluation Metrics

How well is my model doing?

Confusion Matrix Table

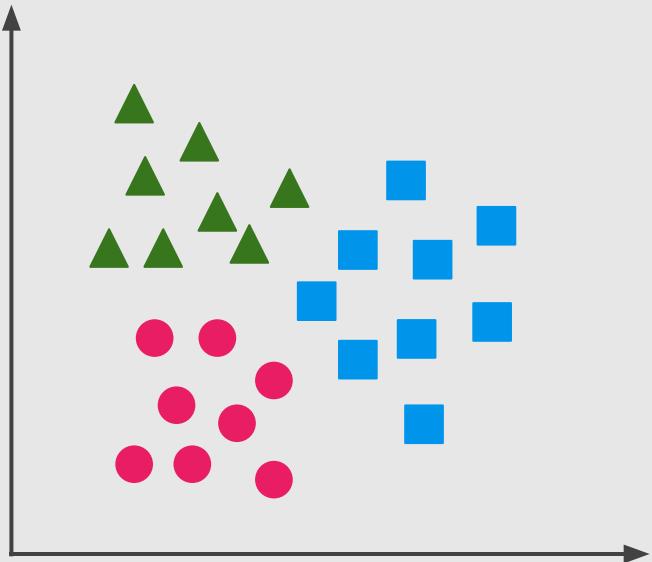
		Diagnosed Sick	Diagnosed Healthy
Sick	True Positive 	False Negative 	
Healthy	False Positive 	True Negative 	

Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●



		Predicted Class		
		Guessed Class 1	Guessed Class 2	Guessed Class 3
True Class	Class 1	5	2	1
	Class 2	3	6	0
	Class 3	0	1	7

Accuracy



		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Accuracy =

$$\frac{TP + TN}{TP + TN + FP + FN} =$$

$$\frac{1,000 + 8,000}{10,000} = 90\%$$

Normalized Accuracy



		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Normalized Accuracy =

$$\frac{\frac{TP}{TP + FN} + \frac{TN}{TN + FP}}{2} =$$

$$\frac{\frac{1000}{1000 + 200} + \frac{8000}{8000 + 800}}{2} =$$

$$\frac{83.3 + 90.9}{2} = 87.1\%$$

Precision



		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Precision:

Out of all the patients we diagnosed with illness, how many were actually sick?

Precision =

$$\frac{1,000}{1,000 + 800} = 55.7\%$$

Recall



		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

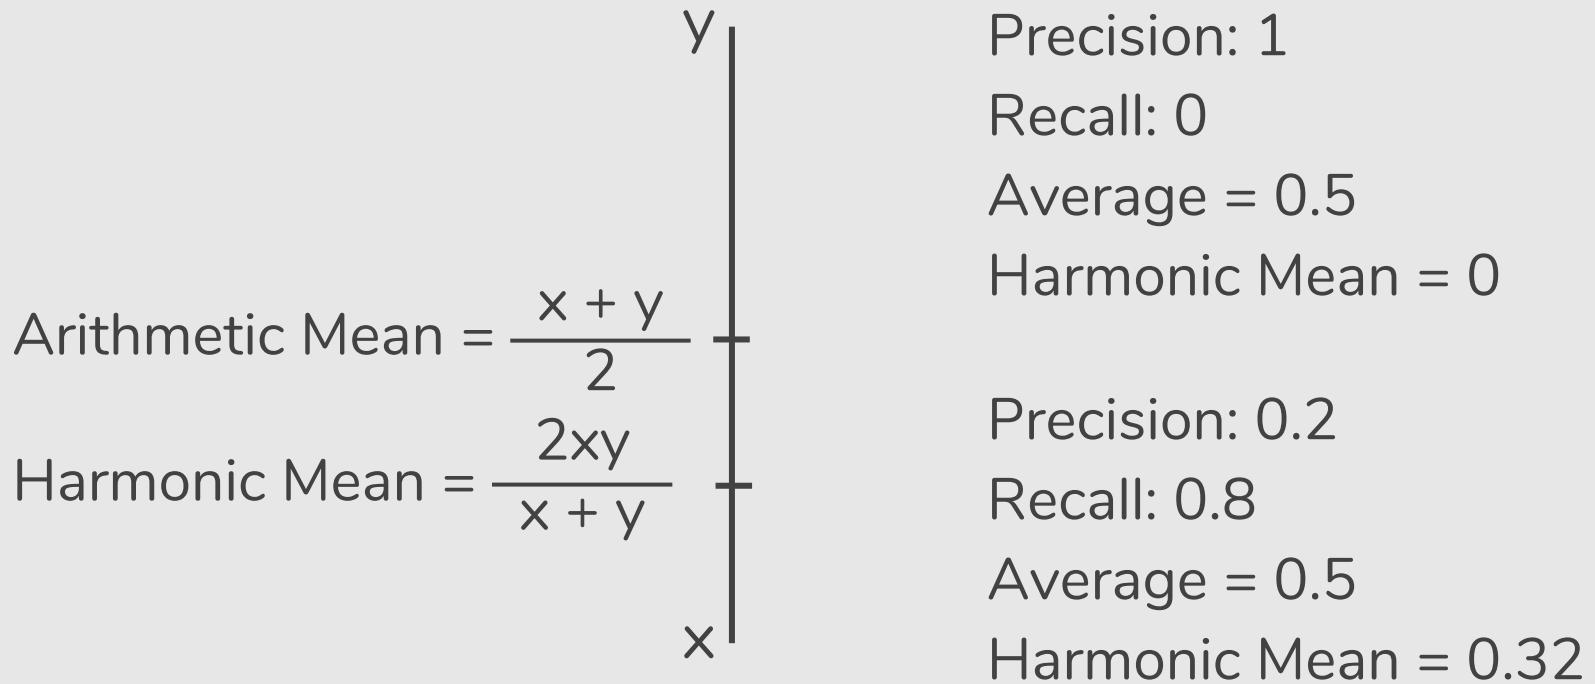
Recall:

Out of all the sick patients, how many did we correctly diagnose as sick?

Recall =

$$\frac{1,000}{1,000 + 200} = 83.3\%$$

Harmonic Mean



F1 Score = Harmonic Mean (Precision, Recall)

F_β Score

$$F_1 = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

$$F_{\beta} = (1 + \beta^2) \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

F_β Score



Precision
(no false positive)

F0.5 Score

F1 Score

F2 Score

Recall
(no false negative)

F10 Score



References

- https://en.wikipedia.org/wiki/Precision_and_recall
- https://en.wikipedia.org/wiki/Binary_classification
- https://en.wikipedia.org/wiki/F1_score
- <https://www.quora.com/What-is-an-intuitive-explanation-of-F-score>
- “Approximate Statistical Tests for Comparing Supervised Classification Learning Algorithms”, Neural Computation, p. 1895-1923, 1998
<https://www.mitpressjournals.org/doi/10.1162/089976698300017197>
- “**A systematic analysis of performance measures for classification tasks**”, 2009
<https://www.sciencedirect.com/journal/information-processing-and-management>
- Luis Serrano: <https://www.youtube.com/watch?v=aDW44NPhNw0>

Artificial Neural Networks

Machine Learning and Pattern Recognition

Prof. Sandra Avila
Institute of Computing (IC/Unicamp)

Many inventions were inspired by Nature ...

Birds inspired us to fly



Many inventions were inspired by Nature ...

Burdock plants inspired velcro



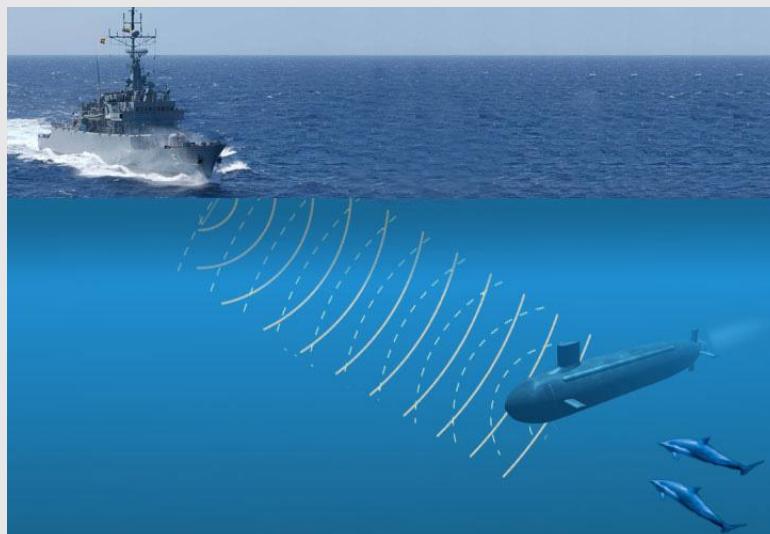
<http://www.youngworldclub.com/wp-content/uploads/2017/04/velcro.jpeg>



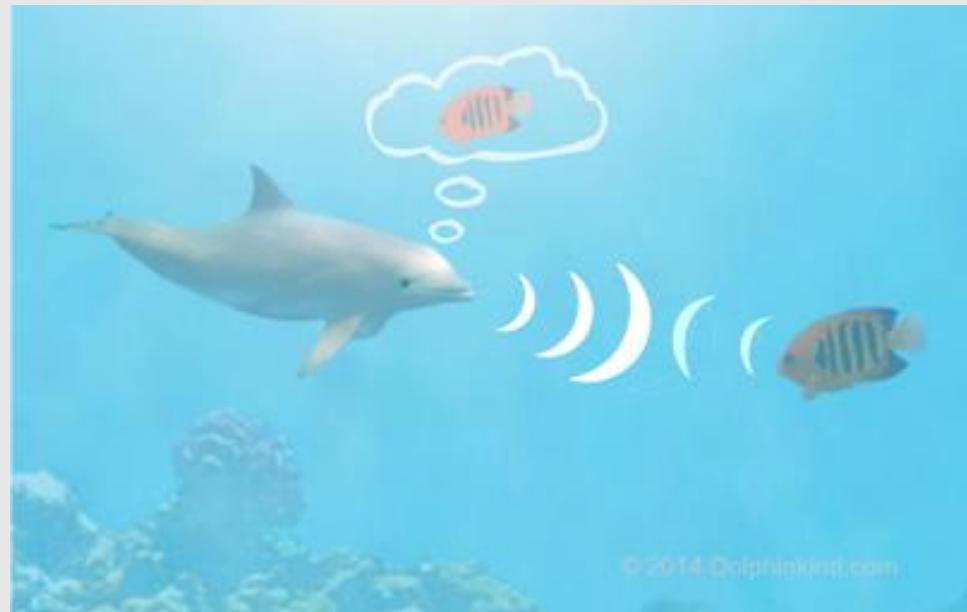
https://c1.staticflickr.com/3/28882961455_f2d6dee40c_b.jpg

Many inventions were inspired by Nature ...

Dolphins inspired sonar development



https://sites.google.com/site/echolocationkawproject/_/rsrc/1459209762464/sonars/image.jpeg



©2014-Dolphinkind.com
http://www.dolphinkind.com/images/dolphin_echolocation.jpg

Many inventions were inspired by Nature ...

Dogs inspired ...

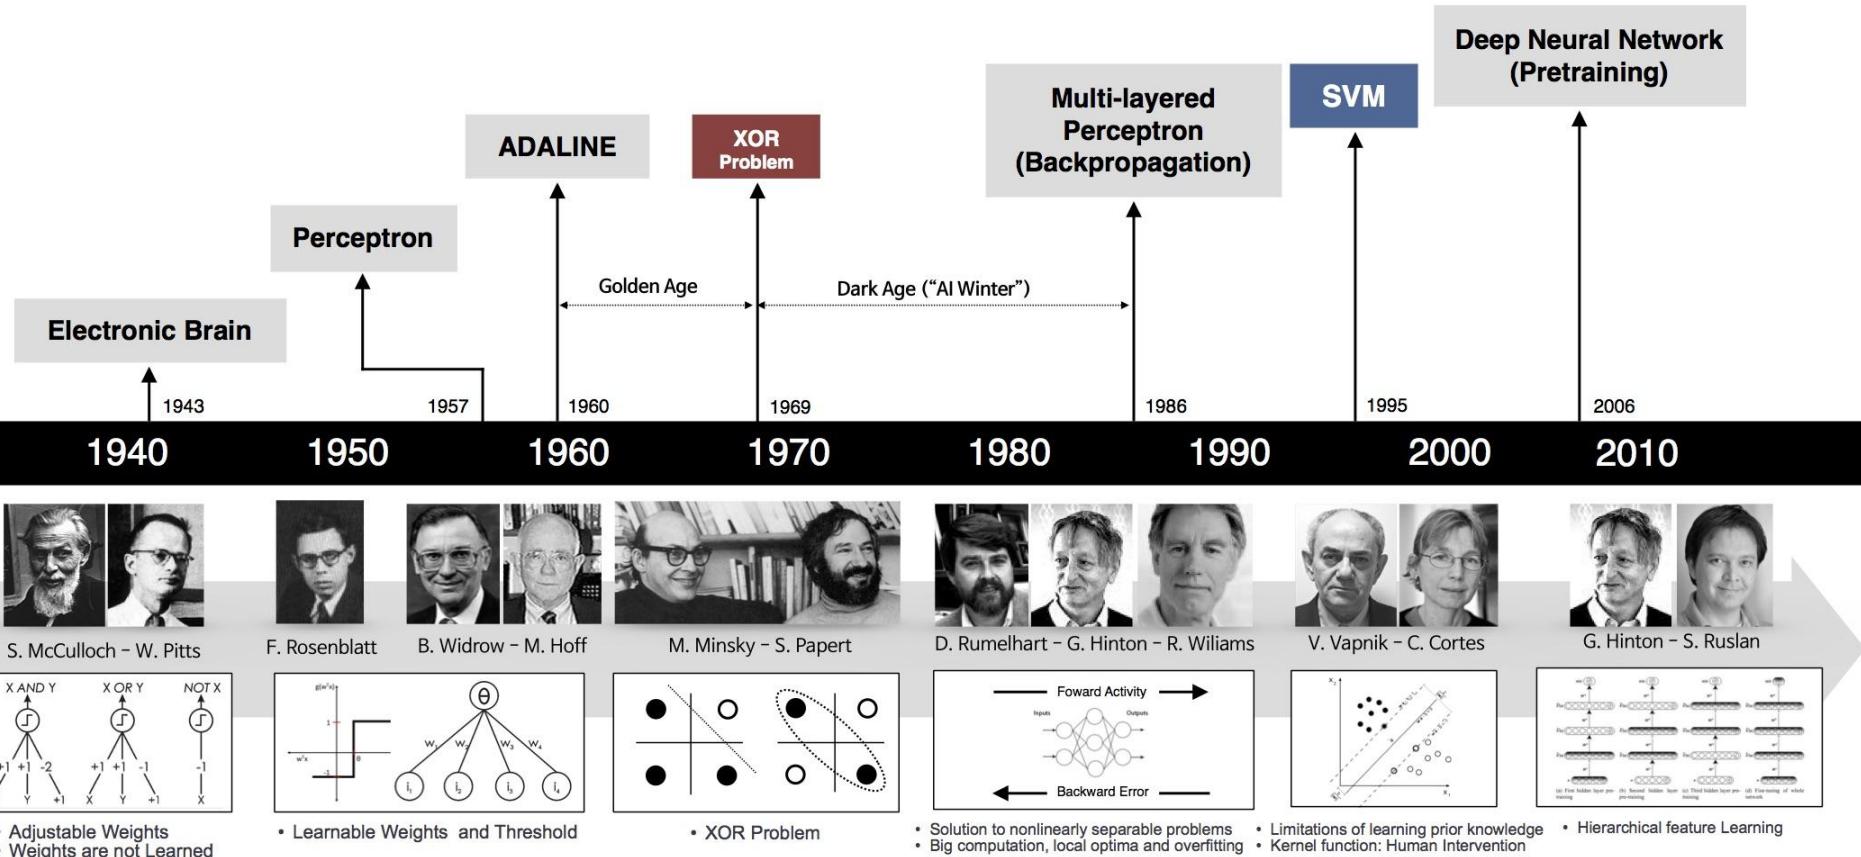


It seems logical to look at the
brain's architecture for inspiration on
how to build an intelligent machine.

Today's Agenda

- Artificial Neural Networks
 - From Biological to Artificial Neurons
 - Biological Neurons
 - Logical Computations with Neurons
 - The Perceptron
 - Examples
 - Neural Network

From Biological to Artificial Neurons



References

- “A logical calculus of the ideas immanent in nervous activity”, 1943:
<https://link.springer.com/content/pdf/10.1007%2FBF02478259.pdf>
- “The Perceptron: A perceiving and recognizing automaton”, 1957
<https://www.import.io/wp-content/uploads/2017/06/rosenblatt-1957.pdf>
- “Perceptrons: An Introduction to Computational Geometry”, 1969
- “Learning representations by back-propagating errors”, 1986
<https://www.nature.com/articles/323533a0>
https://web.stanford.edu/class/psych209a/ReadingsByDate/02_06/PDPVollChapter8.pdf



ILSVRC 2012 – Image Classification task

Rank	Name	Error Rate (%)	Description
1	University of Toronto		Deep Learning
2	University of Tokyo	26.2	Hand-crafted features and learning models
3	University of Oxford	26.9	
4	Xerox/INRIA	27.0	

Object recognition over 1,000,000 images and 1,000 categories (2 GPU)

“ImageNet classification with deep convolutional neural networks”. Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton. In: NIPS, 2012.

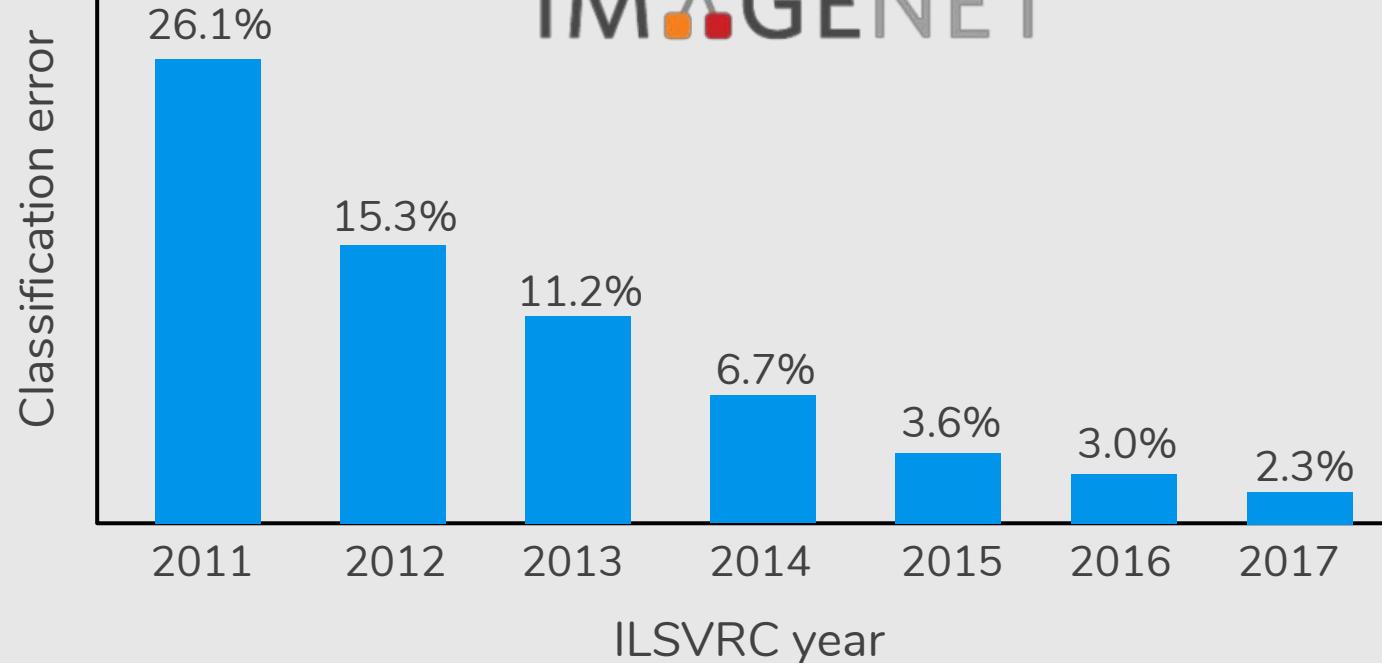


ILSVRC 2012 – Image Classification task

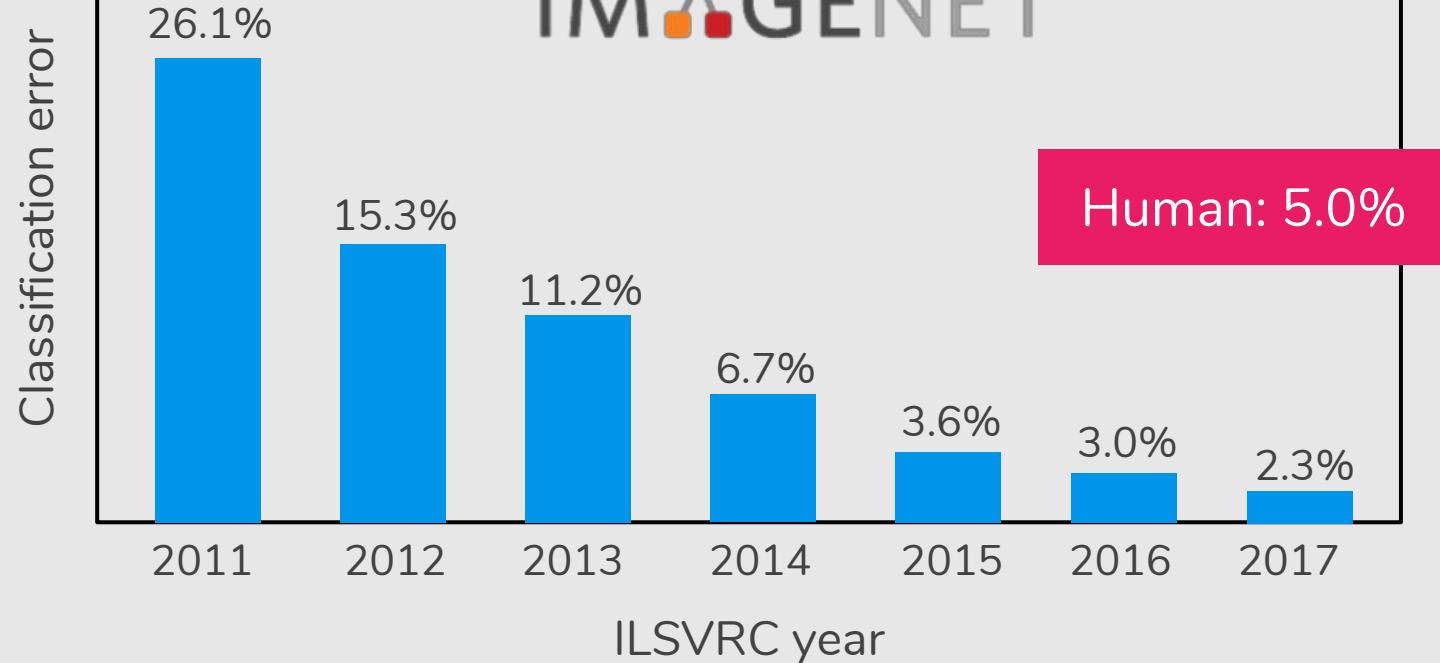
Rank	Name	Error Rate (%)	Description
1	University of Toronto	15.3	Deep Learning
2	University of Tokyo	26.2	Hand-crafted features and learning models
3	University of Oxford	26.9	
4	Xerox/INRIA	27.0	

Object recognition over 1,000,000 images and 1,000 categories (2 GPU)

“ImageNet classification with deep convolutional neural networks”. Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton. In: NIPS, 2012.



“ImageNet classification with deep convolutional neural networks”. Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton. In: NIPS, 2012.



“ImageNet classification with deep convolutional neural networks”. Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton. In: NIPS, 2012.

Will this wave die out like the
previous ones did?

From Biological to Artificial Neurons

1. There is now a **huge quantity of data** available to train neural networks.



www.image-net.org

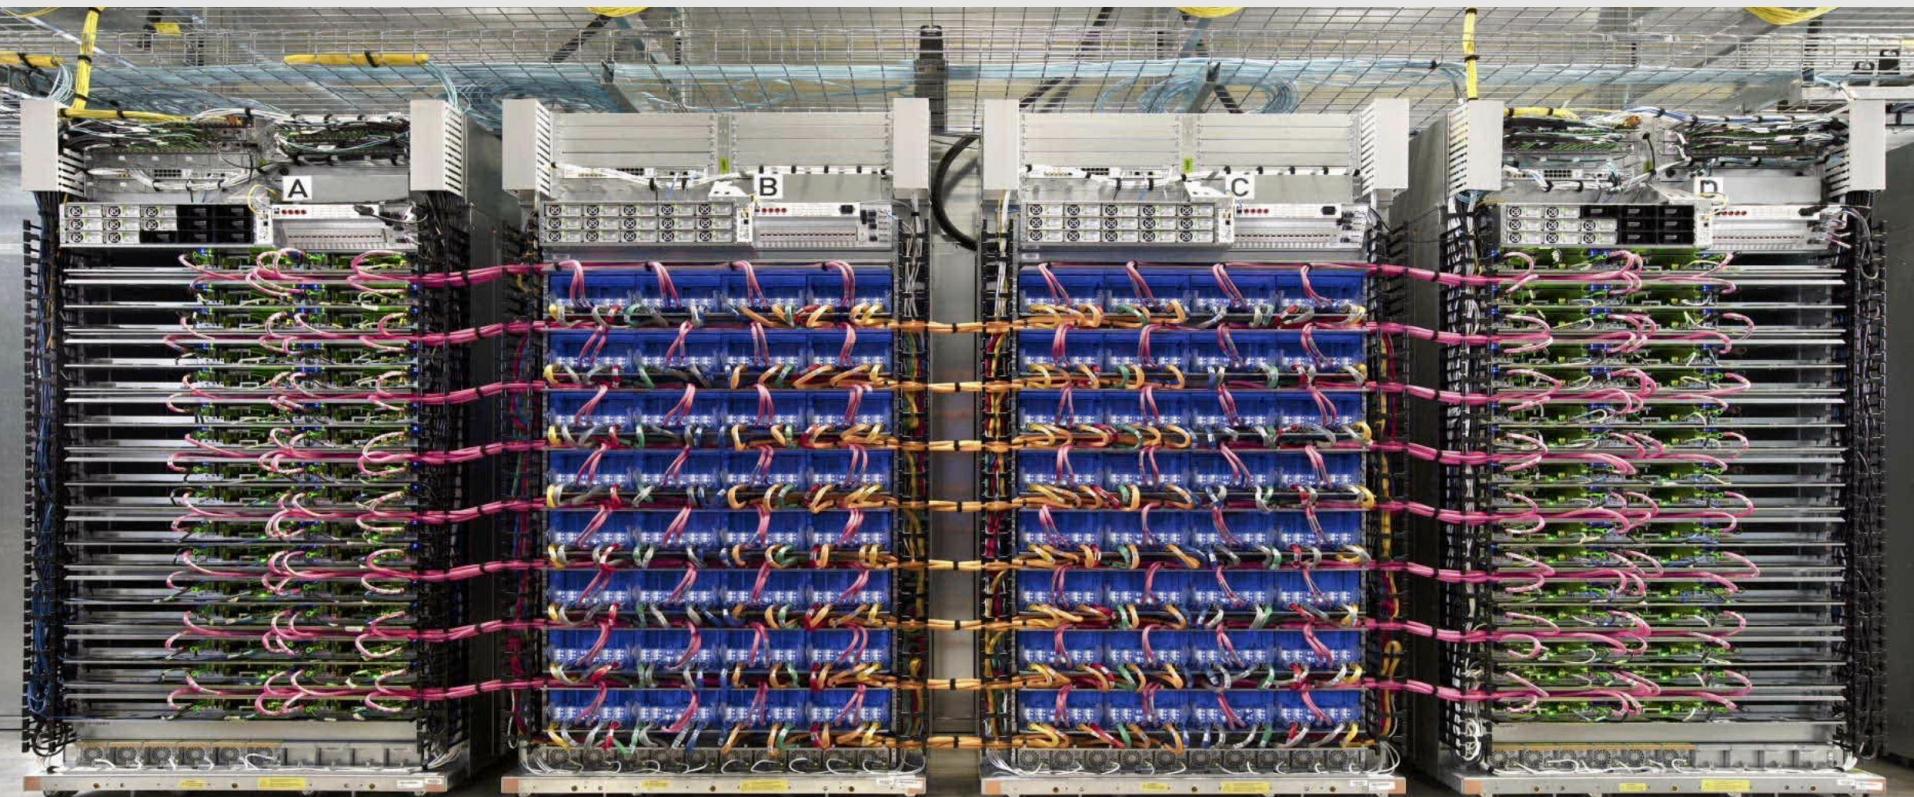
22K categories and **14M** images

- Animals
 - Bird
 - Fish
 - Mammal
 - Invertebrate
- Plants
 - Tree
 - Flower
- Food
- Materials
- Structures
 - Artifact
 - Tools
 - Appliances
 - Structures
- Person
- Scenes
 - Indoor
 - Geological Formations
- Sport Activities



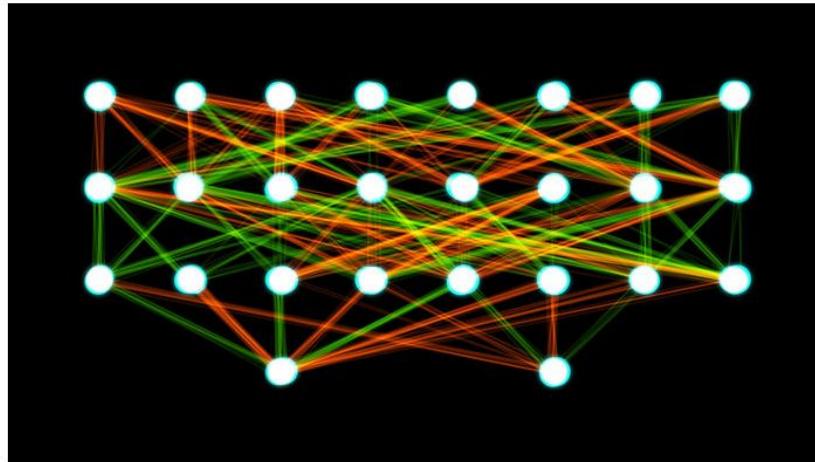
From Biological to Artificial Neurons

1. There is now a **huge quantity of data** available to train neural networks.
2. **Computing power** now makes it possible to train large neural networks in a reasonable amount of time.



From Biological to Artificial Neurons

1. There is now a **huge quantity of data** available to train neural networks.
2. **Computing power** now makes it possible to train large neural networks in a reasonable amount of time.
3. The **training algorithms** have been **improved**.

SHARE

A representation of a neural network.

Akritasa/Wikimedia Commons

Brainlike computers are a black box. Scientists are finally peering inside

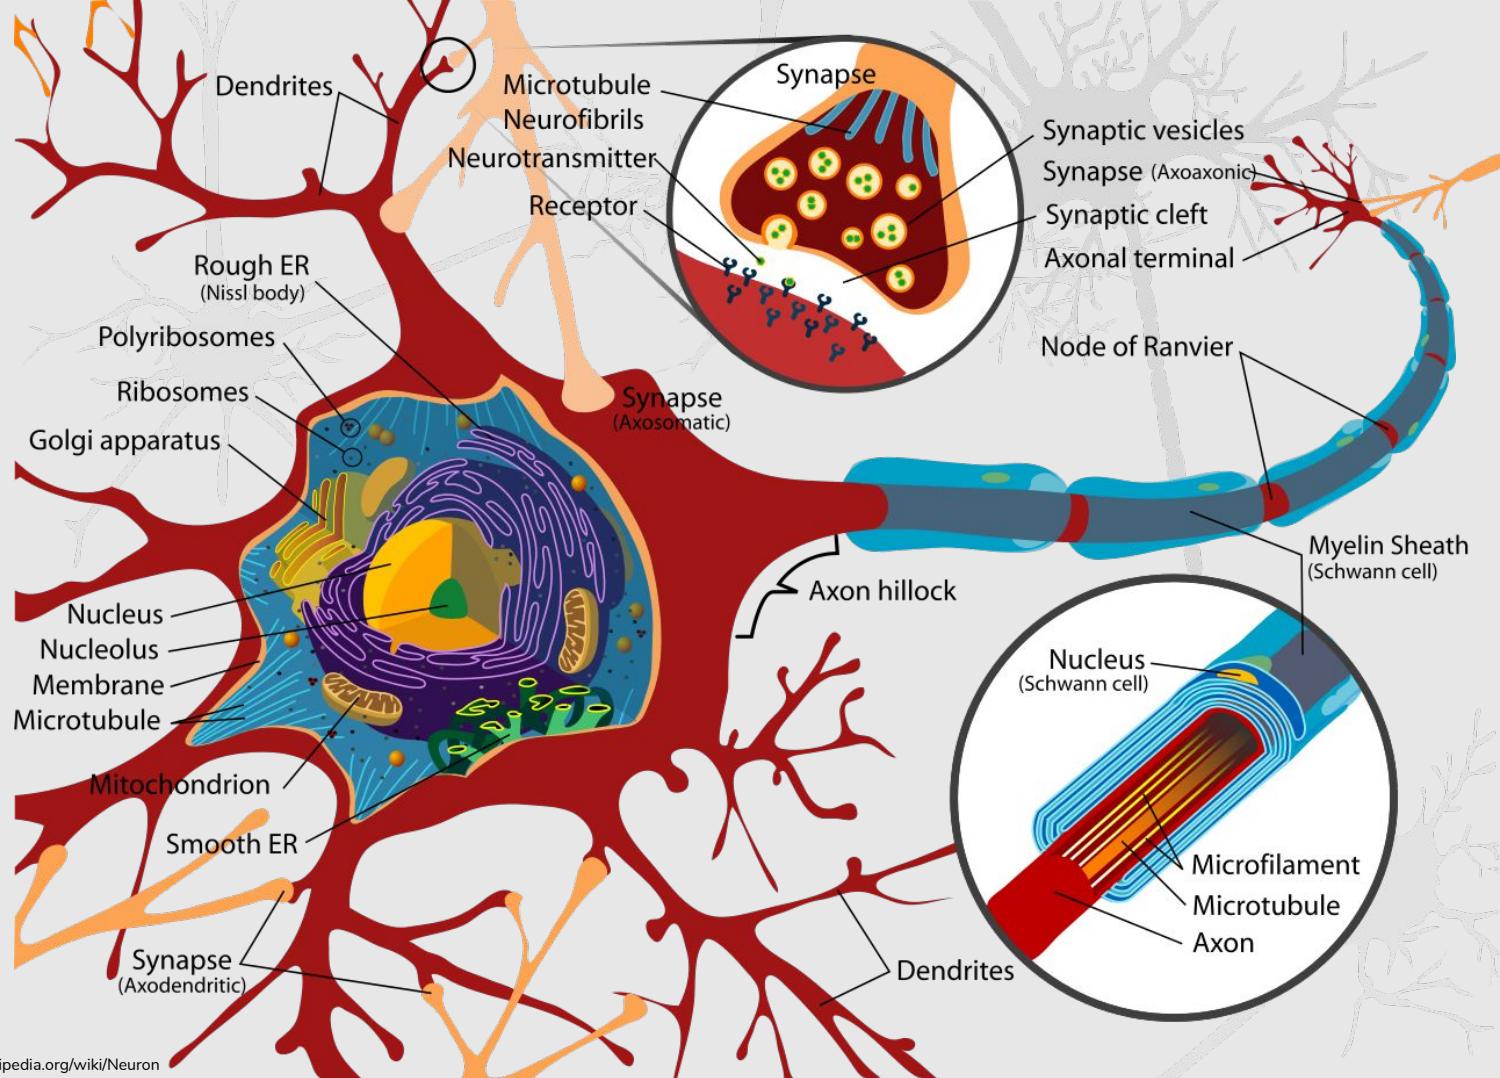
By **Jackie Snow** | Mar. 7, 2017 , 3:15 PM

Last month, Facebook announced software that could simply look at a photo and tell, for example, whether it was a picture of a cat or a dog. A related program identifies cancerous

From Biological to Artificial Neurons

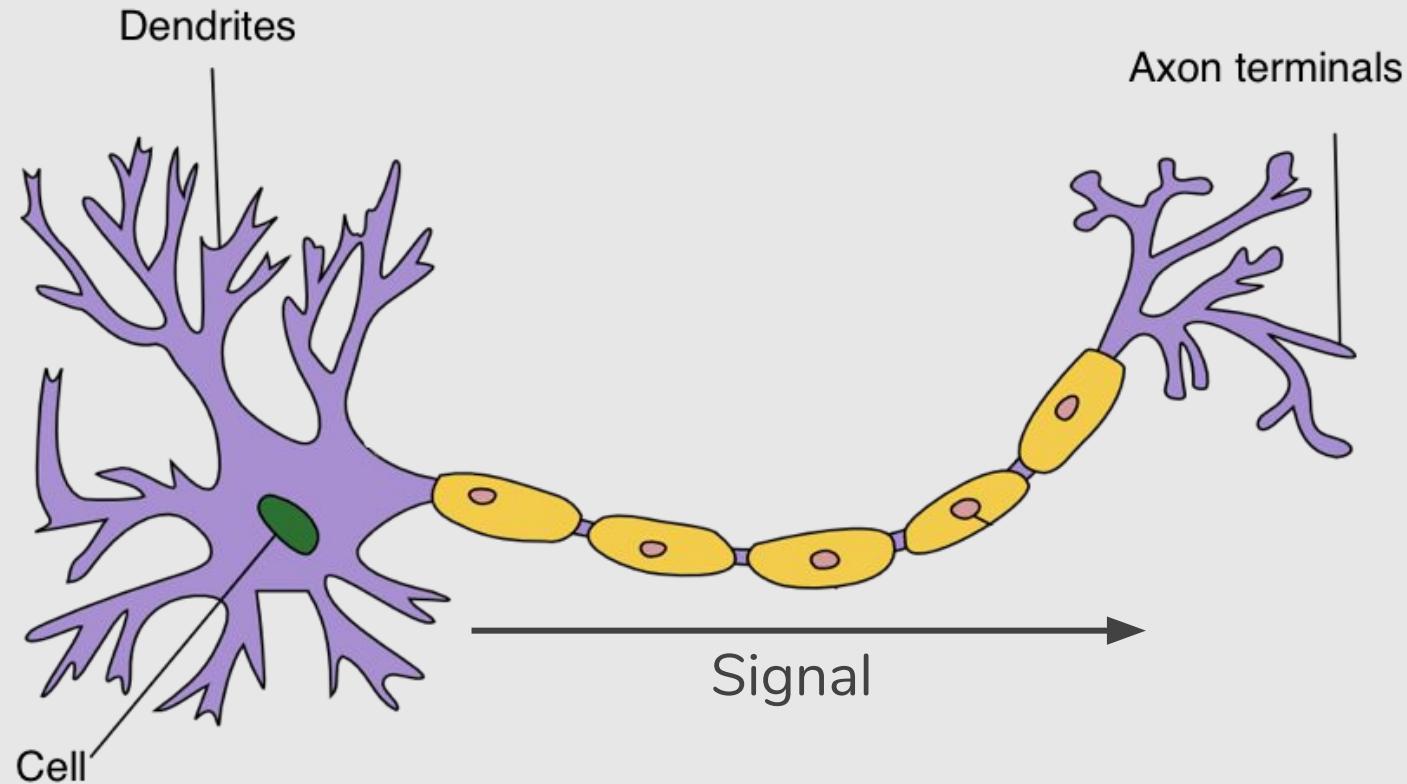
1. There is now a **huge quantity of data** available to train neural networks.
2. **Computing power** now makes it possible to train large neural networks in a reasonable amount of time.
3. The **training algorithms** have been **improved**.
4. ANNs seem to have entered a virtuous circle of funding and progress.

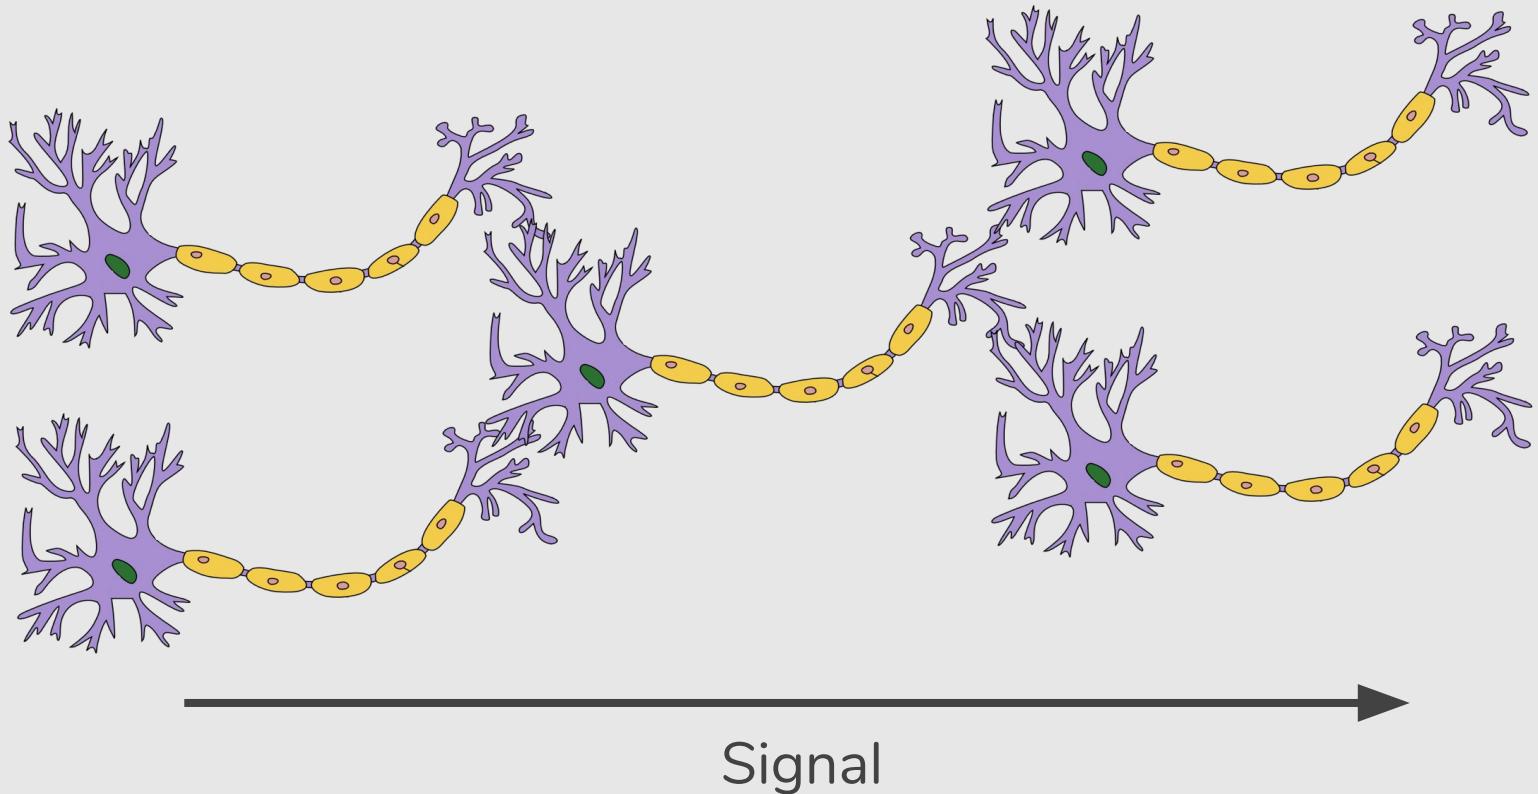
Biological Neurons





<https://www.youtube.com/watch?v=A9Xru1ReRwc>



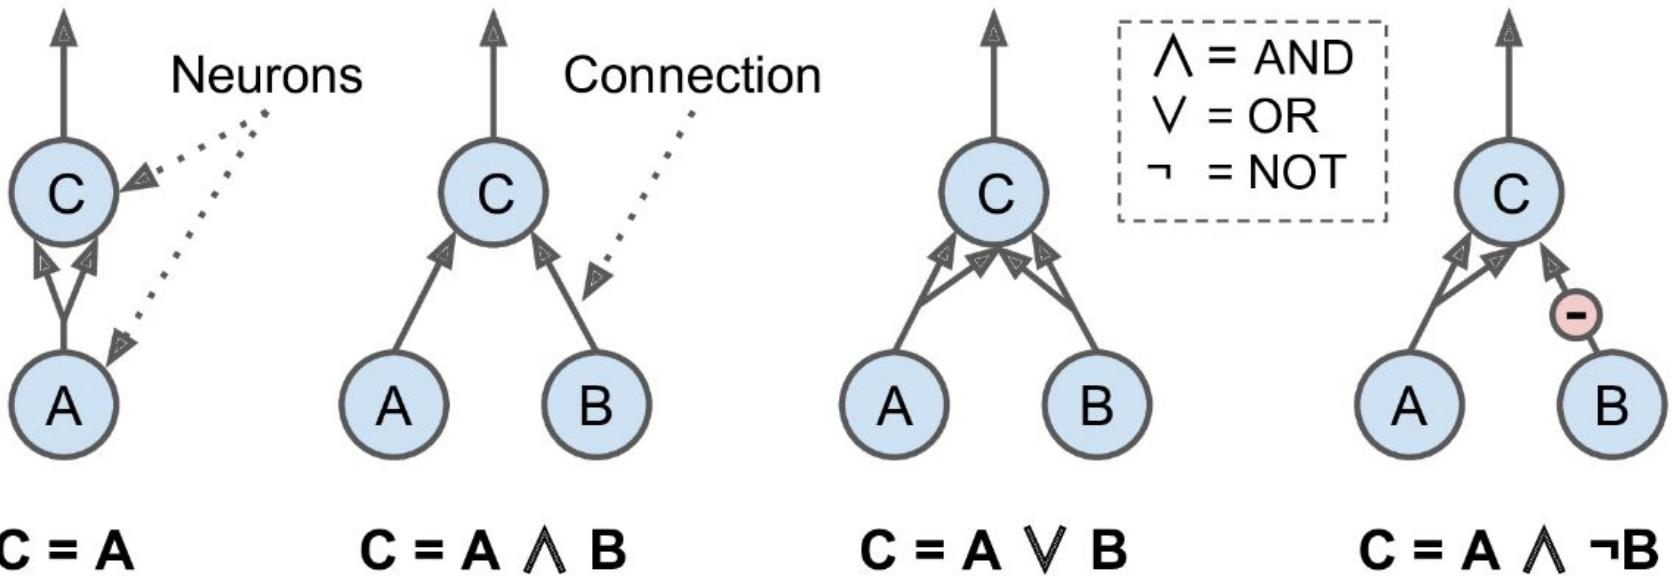


Logical Computations with Neurons

Logical Computations with Neurons

McCulloch and Pitts (1943) proposed a very simple model:

- It has one or more binary (on/off) inputs and one binary output.
- The artificial neuron **simply** activates its output when more than a certain number of its inputs are active.



Artificial Neural Networks performing simple logical computations

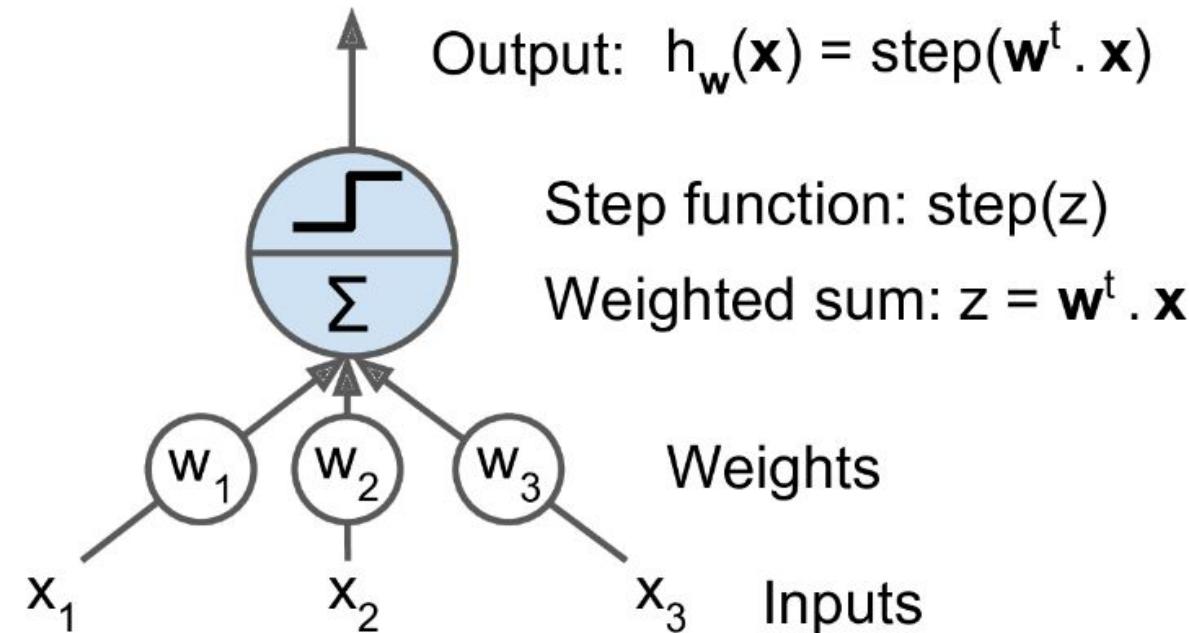
The Perceptron

The Perceptron

Invented in 1957 by Frank Rosenblatt.

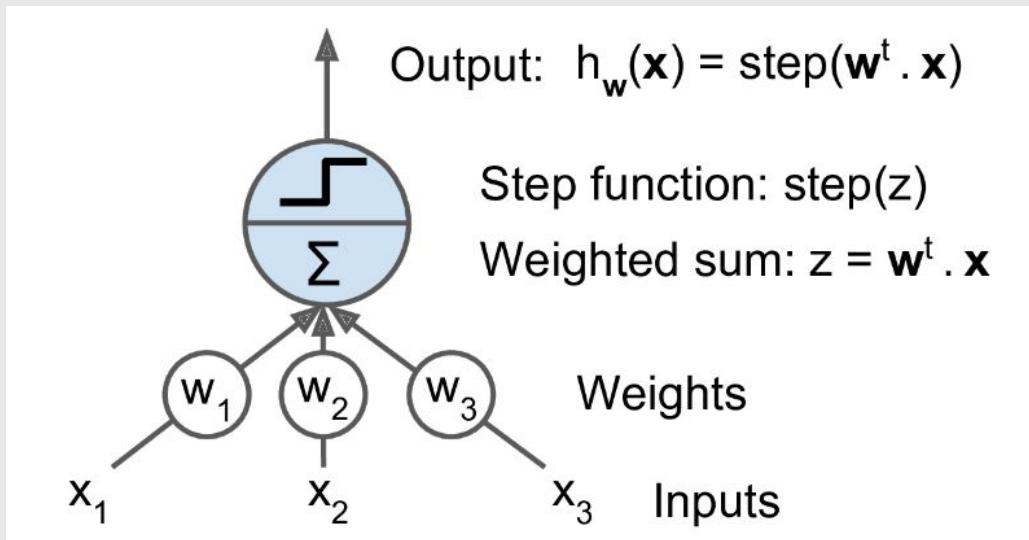
- It is based on a Linear Threshold Unit (LTU):
 - The inputs and output are now **numbers** (instead of binary on/off values) and each input connection is associated with a **weight**.
- The LTU computes a weighted sum of its inputs then it applies a step function to that sum and outputs the result.

The Perceptron



Linear Threshold Unit

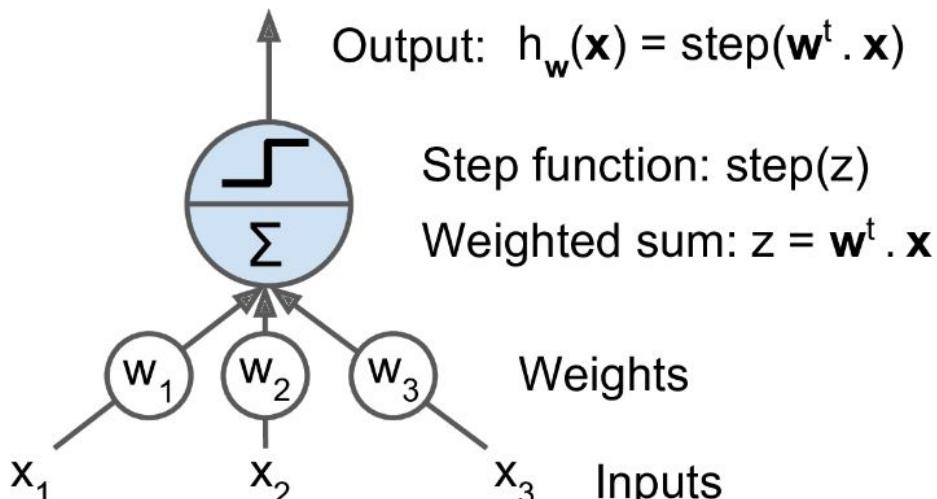
The Perceptron



Linear Threshold Unit

$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

The Perceptron

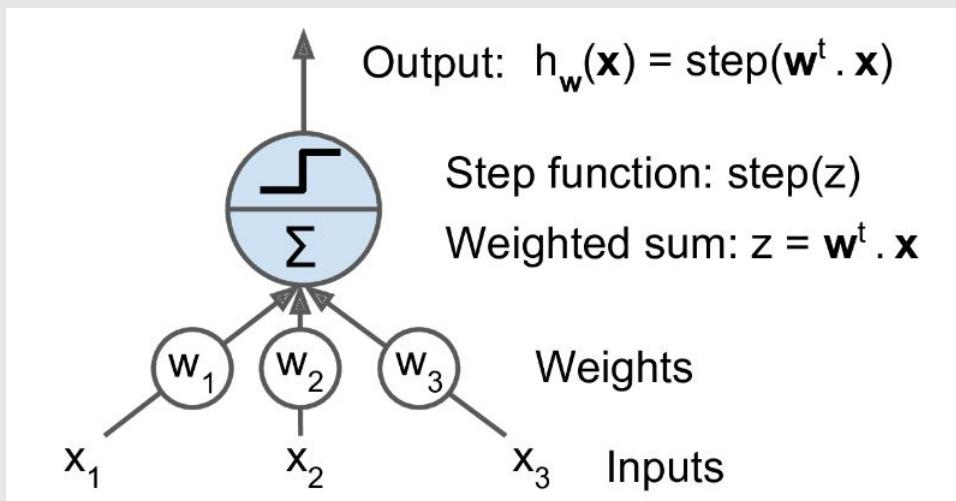


Linear Threshold Unit

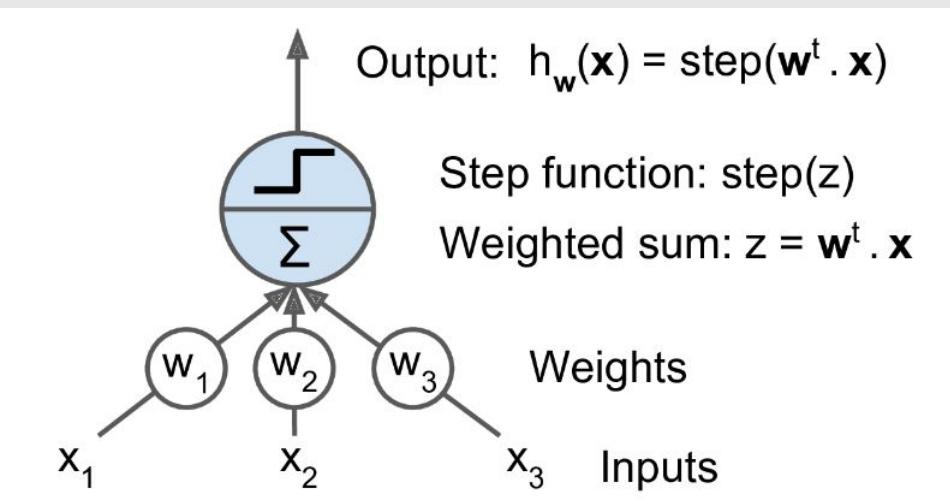
$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

$$\text{sign}(z) = \begin{cases} -1 & \text{if } z < 0 \\ 0 & \text{if } z = 0 \\ +1 & \text{if } z > 0 \end{cases}$$

Neuron Model: Logistic Unit



Neuron Model: Logistic Unit



Inputs

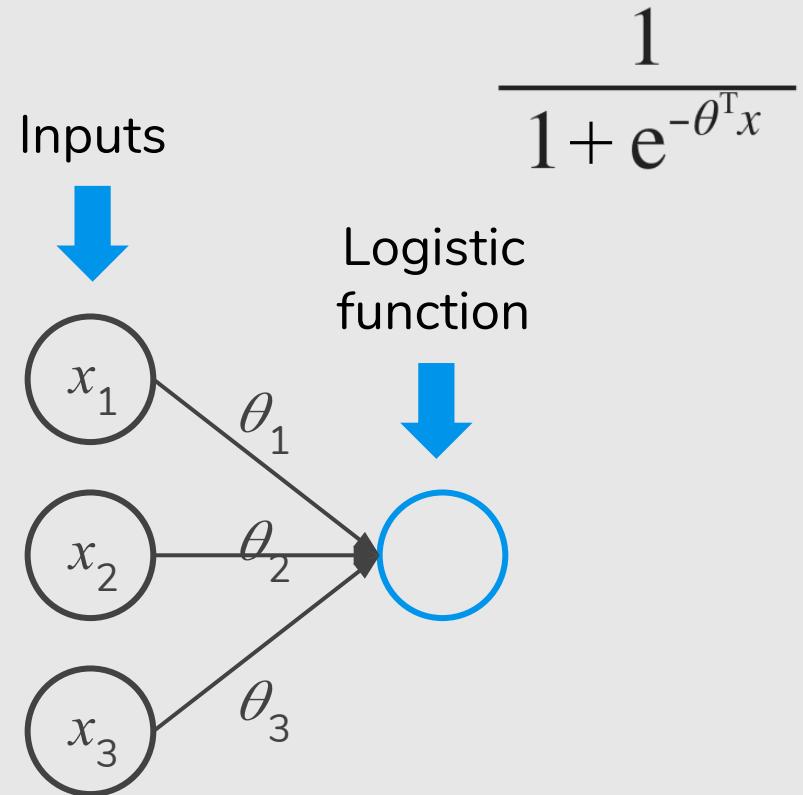
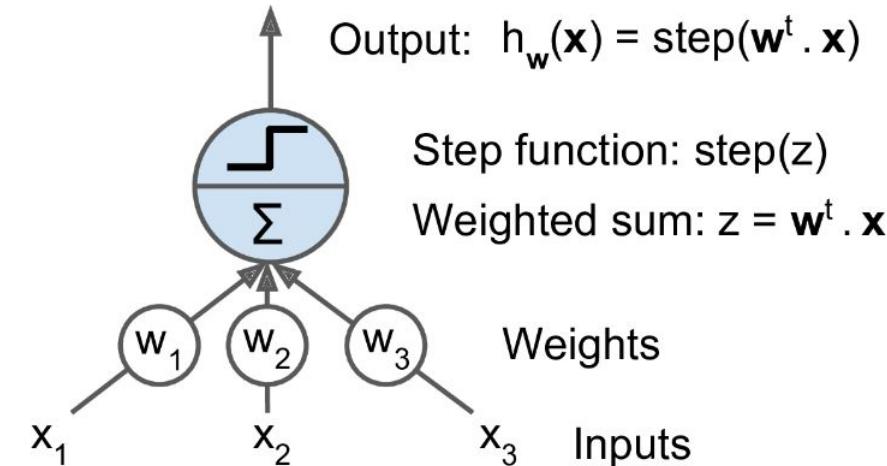


x_1

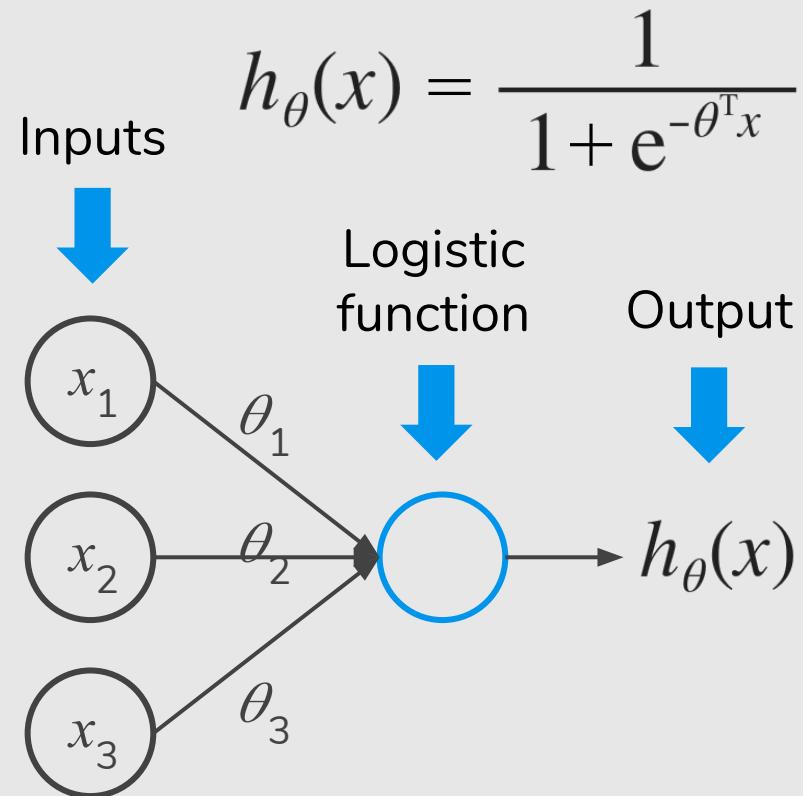
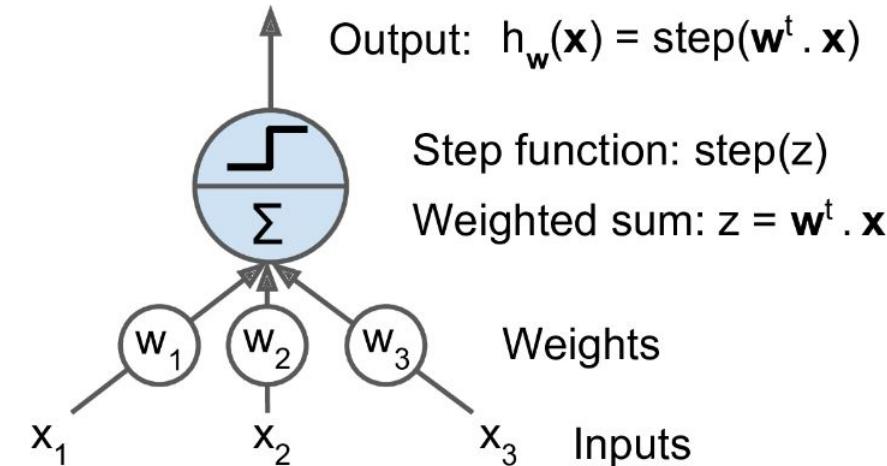
x_2

x_3

Neuron Model: Logistic Unit



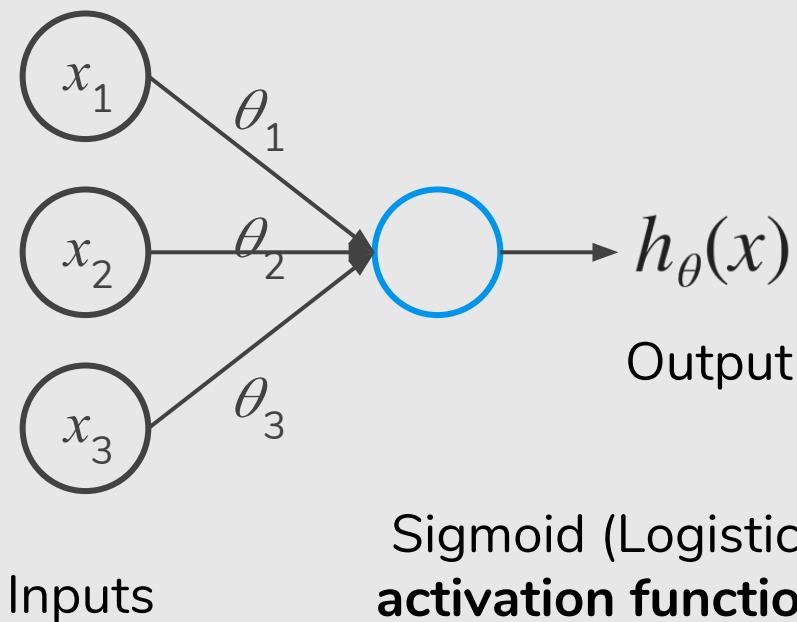
Neuron Model: Logistic Unit



weights



Neuron Model: Logistic Unit



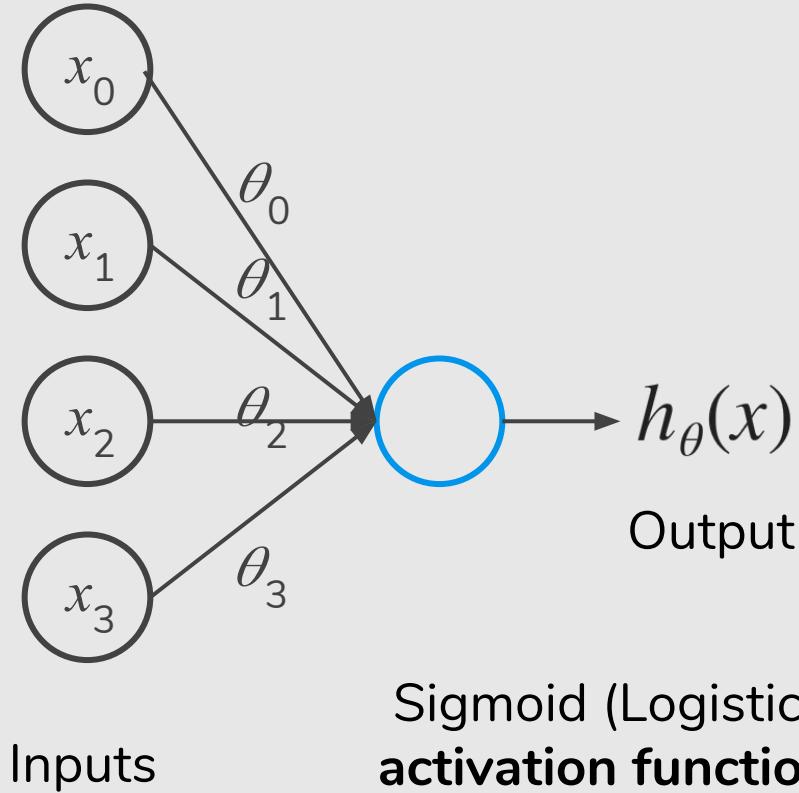
$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

$$h_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$$

weights



Neuron Model: Logistic Unit



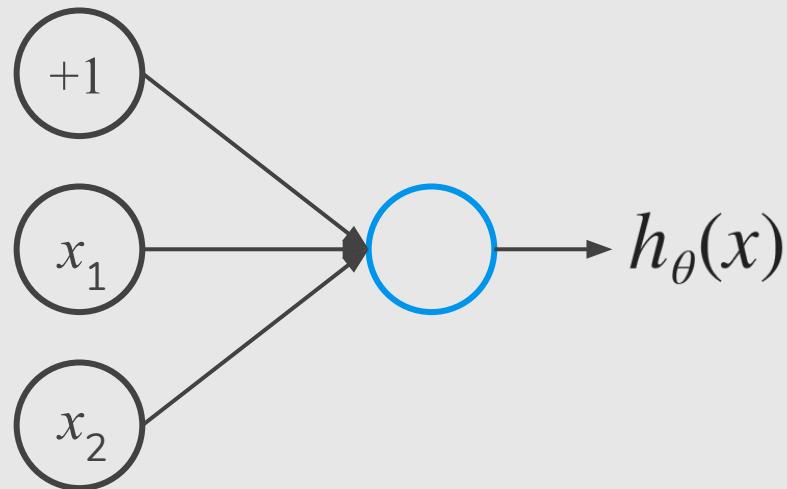
$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Examples

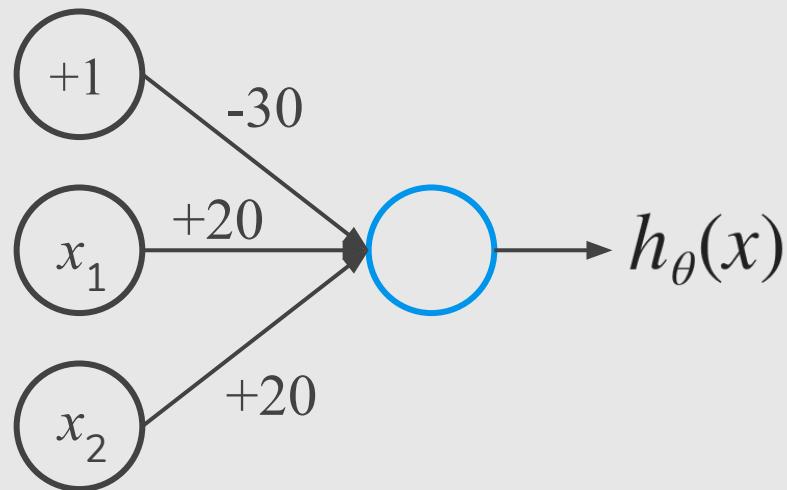
Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



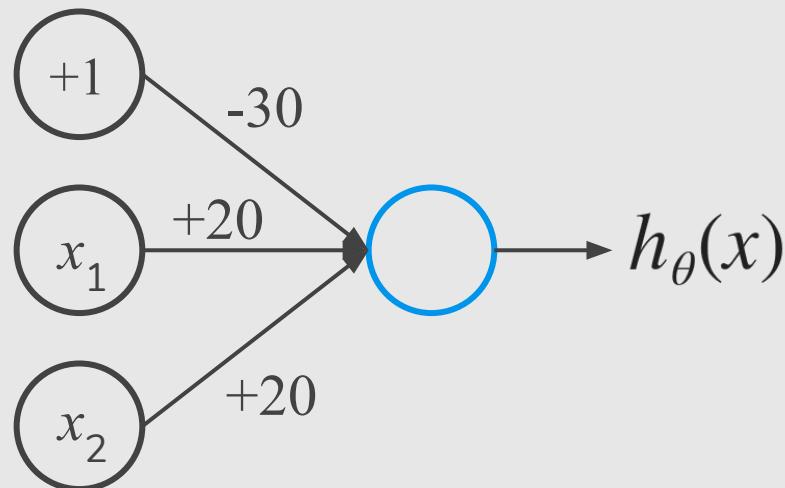
Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



Simple Example: AND

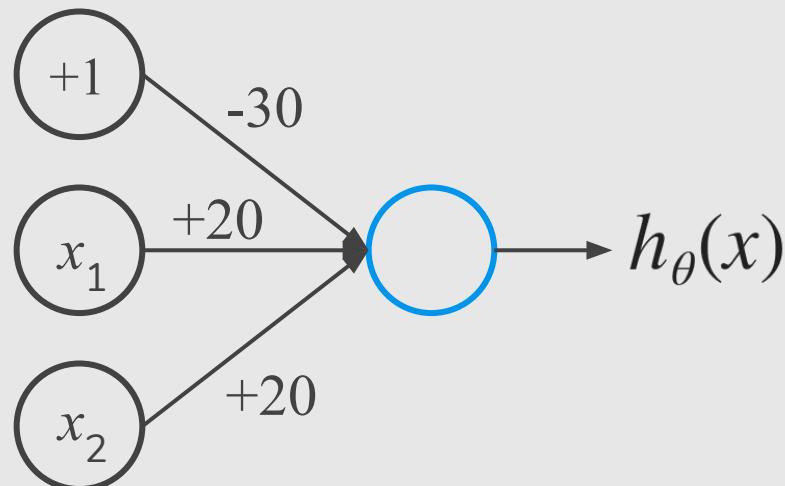
$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



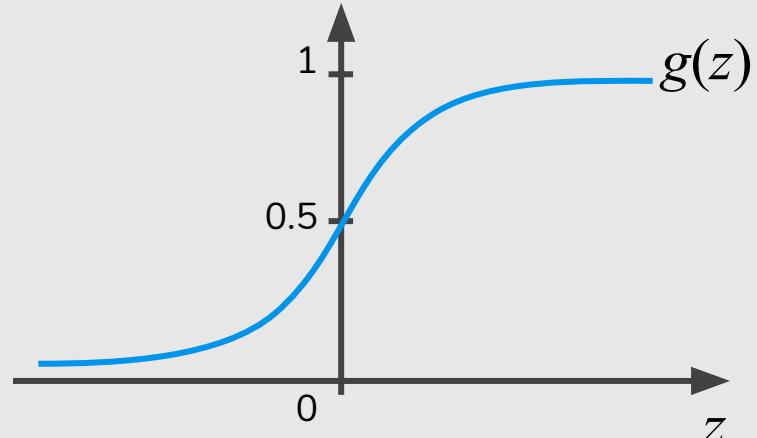
$$h_{\theta}(x) = g(-30 + 20x_1 + 20x_2)$$

Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



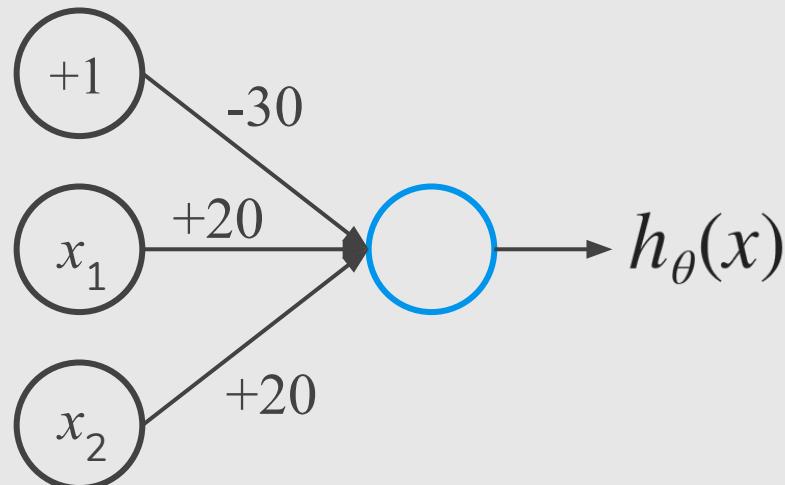
$$h_\theta(x) = g(-30 + 20x_1 + 20x_2)$$



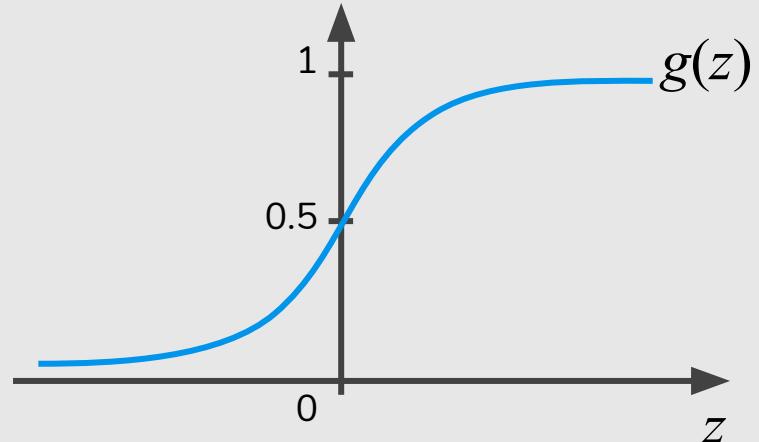
x_1	x_2	$h_\theta(x)$
0	0	0
0	1	0.5
1	0	0.5
1	1	1

Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



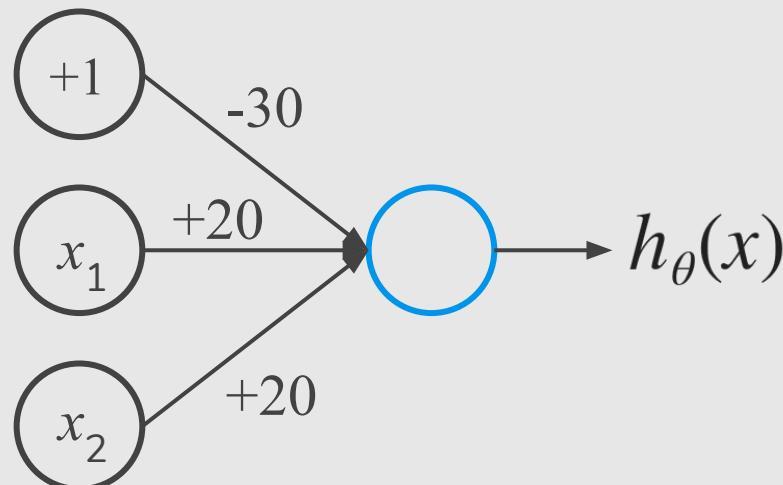
$$h_\theta(x) = g(-30 + 20x_1 + 20x_2)$$



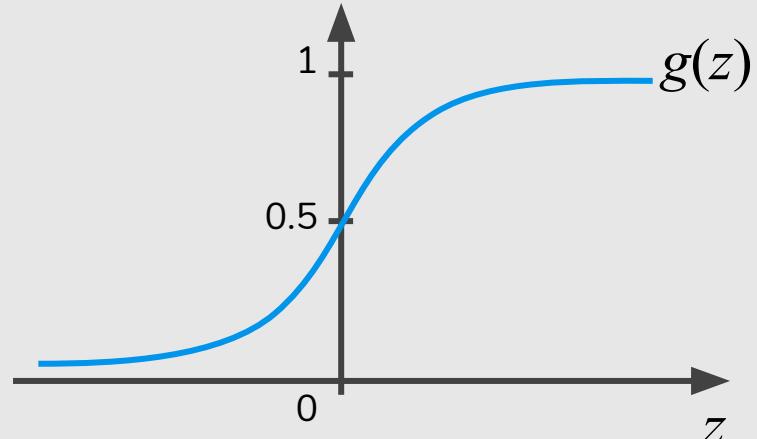
x_1	x_2	$h_\theta(x)$
0	0	$g(-30) \approx 0$
0	1	
1	0	
1	1	

Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



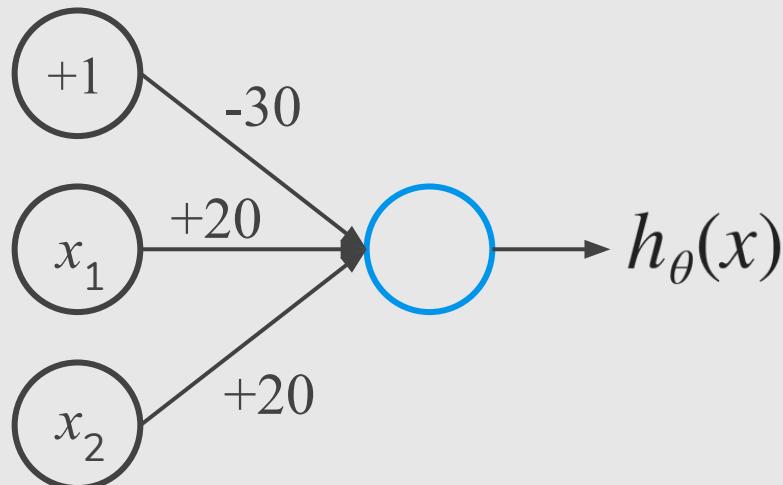
$$h_\theta(x) = g(-30 + 20x_1 + 20x_2)$$



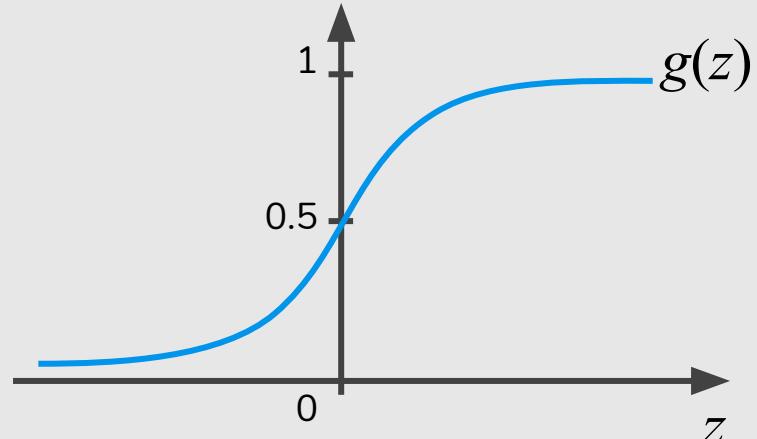
x_1	x_2	$h_\theta(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	
1	1	

Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



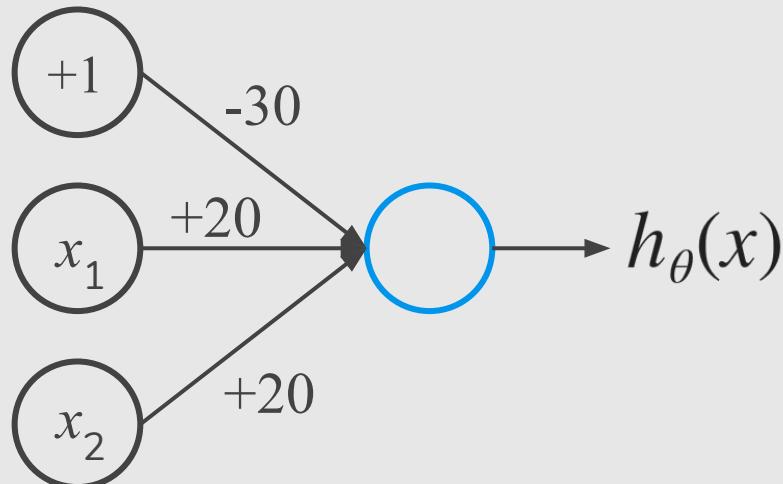
$$h_\theta(x) = g(-30 + 20x_1 + 20x_2)$$



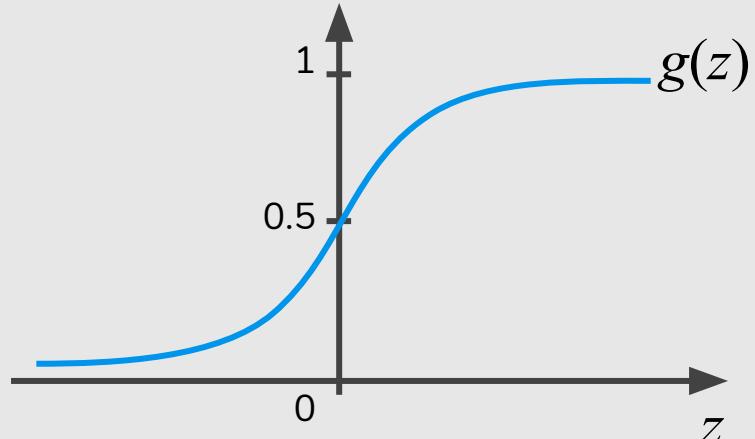
x_1	x_2	$h_\theta(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	1

Simple Example: AND

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ AND } x_2$$



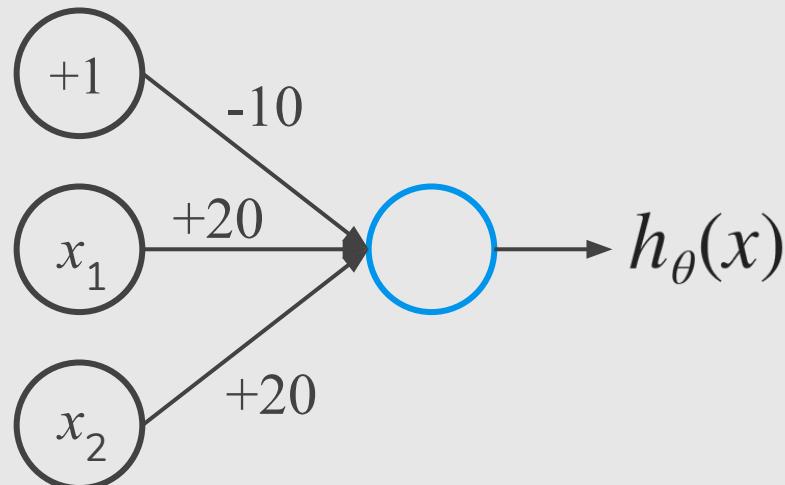
$$h_\theta(x) = g(-30 + 20x_1 + 20x_2)$$



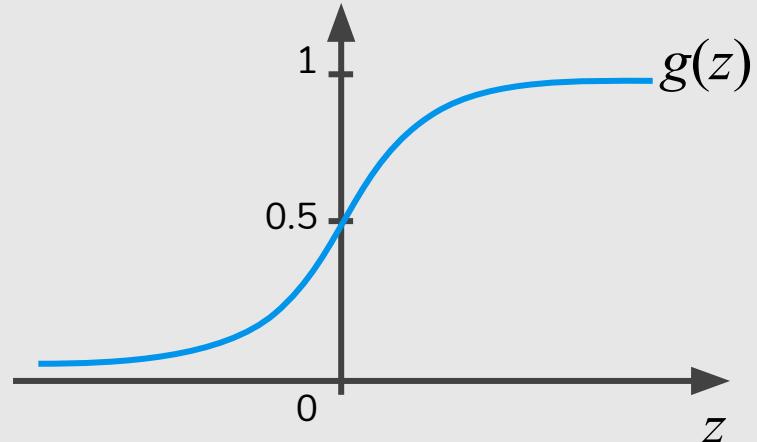
x_1	x_2	$h_\theta(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

Simple Example: OR

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ OR } x_2$$



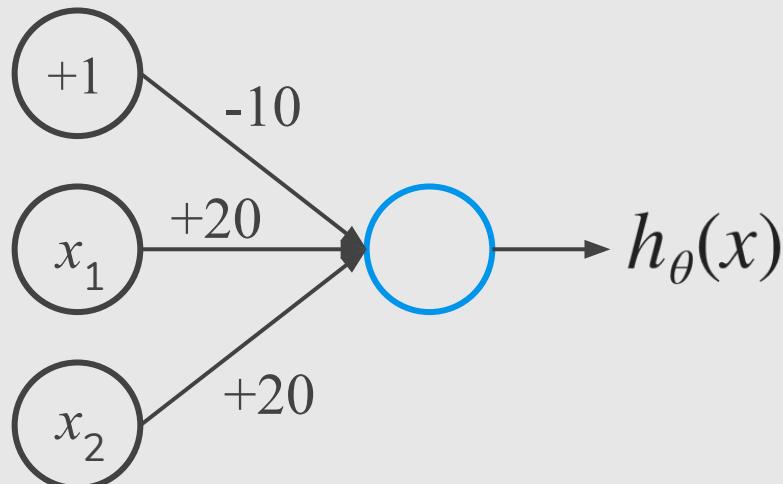
$$h_\theta(x) = g(-10 + 20x_1 + 20x_2)$$



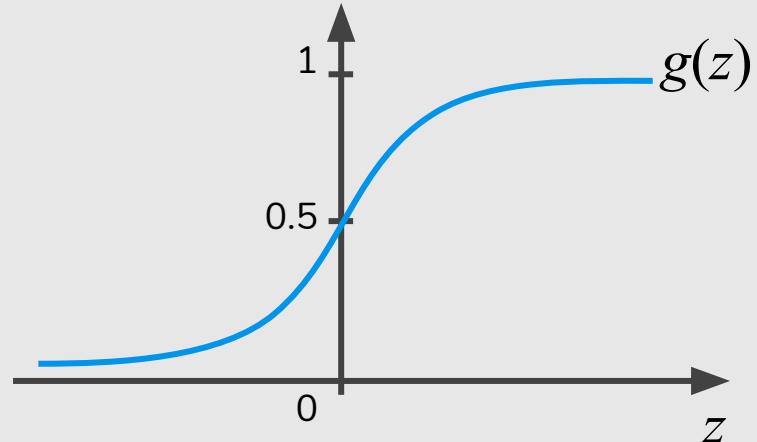
x_1	x_2	$h_\theta(x)$
0	0	0
0	1	0.5
1	0	0.5
1	1	1

Simple Example: OR

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ OR } x_2$$



$$h_\theta(x) = g(-10 + 20x_1 + 20x_2)$$



x_1	x_2	$h_\theta(x)$
0	0	$g(-10) \approx 0$
0	1	$g(10) \approx 1$
1	0	$g(10) \approx 1$
1	1	$g(30) \approx 1$

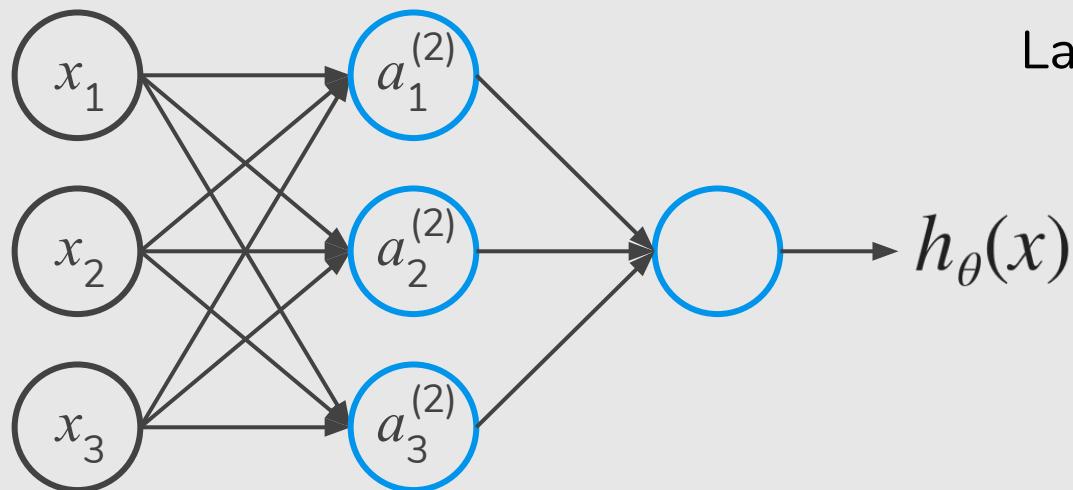
Neural Network

Neural Network

Layer 1 = Input layer

Layer 2 = Hidden layer

Layer 3 = Output layer

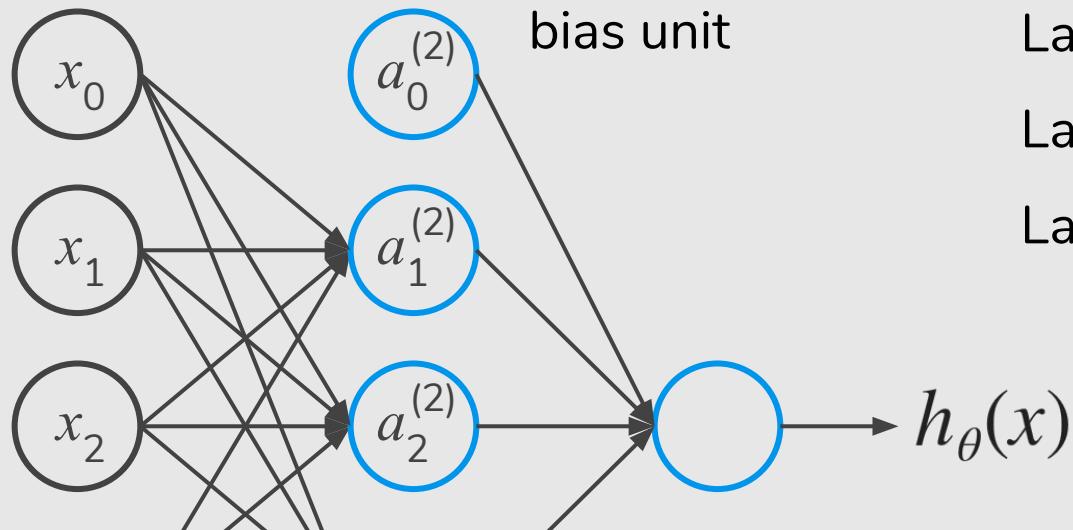


Layer 1

Layer 2

Layer 3

Neural Network



Layer 1

Layer 2

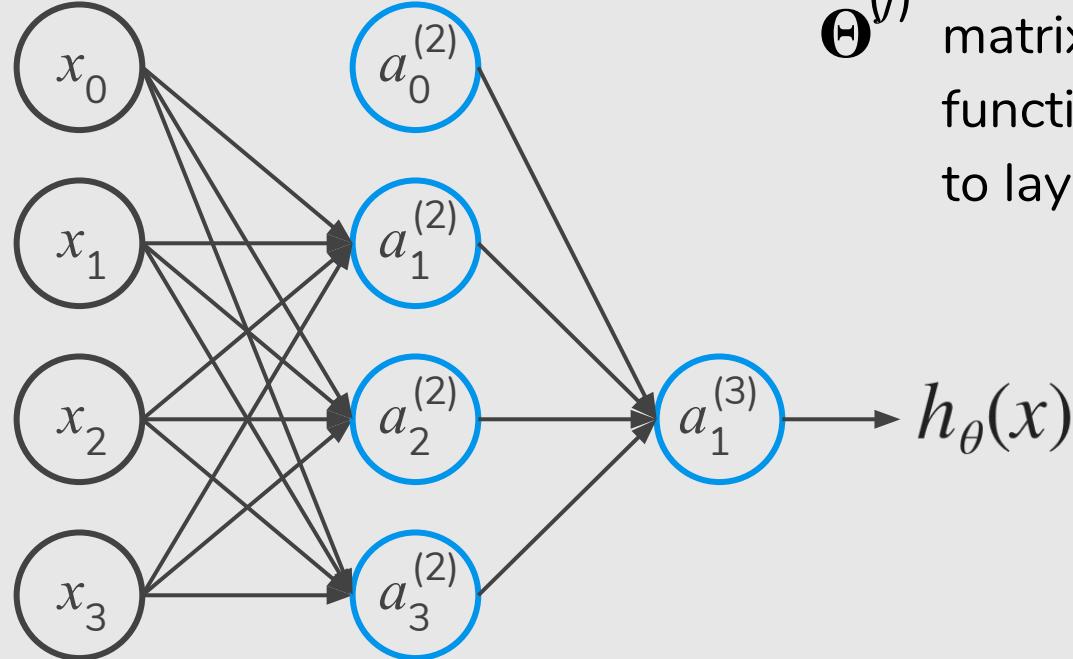
Layer 3

Layer 1 = Input layer

Layer 2 = Hidden layer

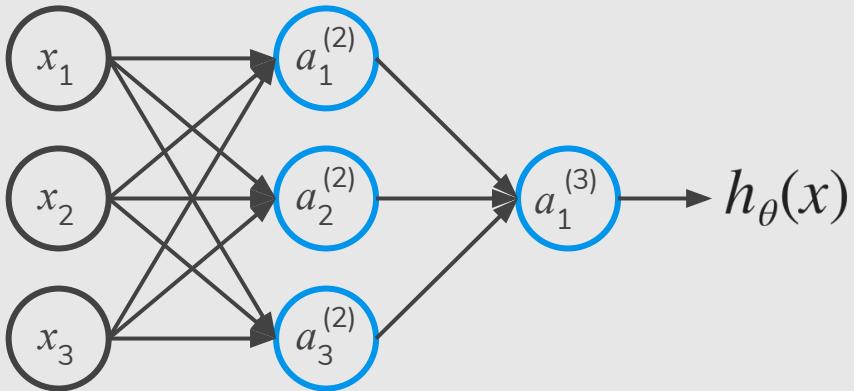
Layer 3 = Output layer

Neural Network



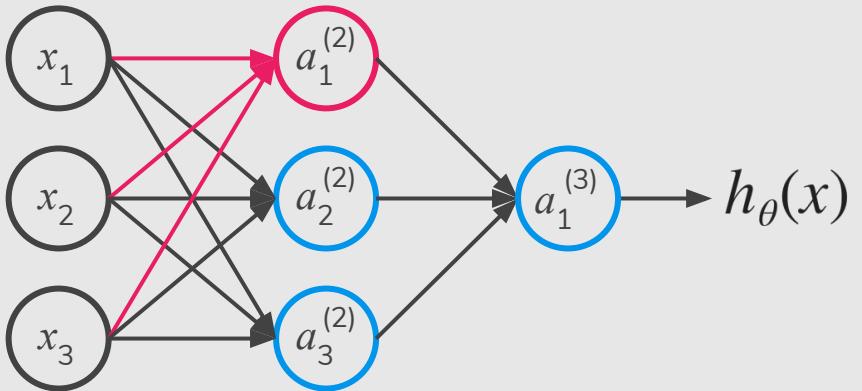
$a_i^{(j)}$ “activation” of unit i in layer j

$\Theta^{(j)}$ matrix of weights controlling function mapping from layer j to layer $j + 1$



$a_i^{(j)}$ “activation” of unit i in layer j

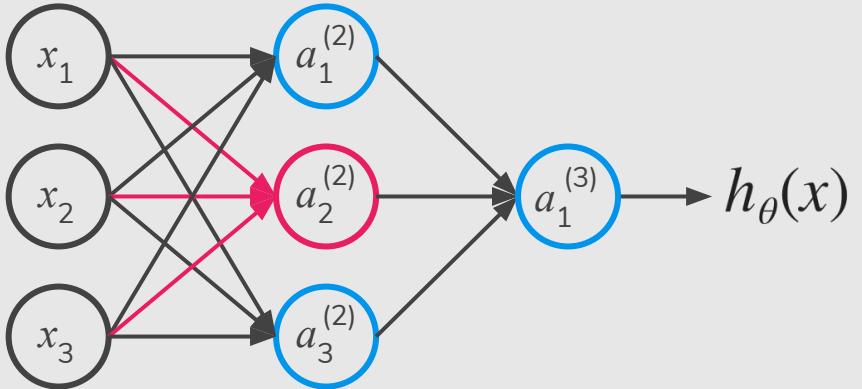
$\Theta^{(j)}$ matrix of weights controlling function mapping from layer j to layer $j + 1$



$a_i^{(j)}$ “activation” of unit i in layer j

$\Theta^{(j)}$ matrix of weights controlling function mapping from layer j to layer $j + 1$

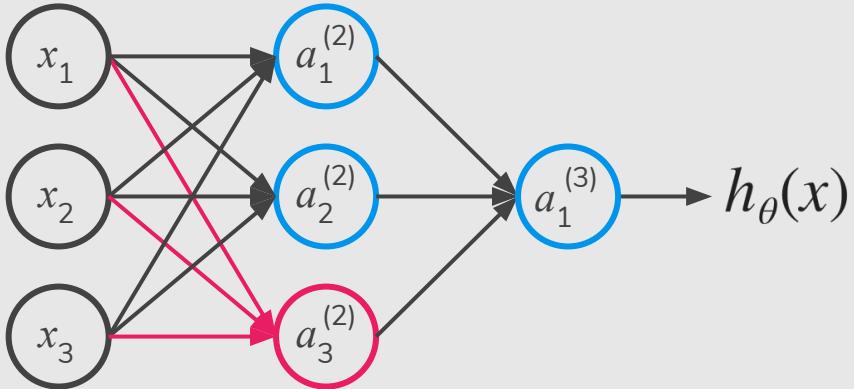
$$a_1^{(2)} = g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3)$$



$a_i^{(j)}$ “activation” of unit i in layer j
 $\Theta^{(j)}$ matrix of weights controlling
 function mapping from layer j
 to layer $j + 1$

$$a_1^{(2)} = g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3)$$



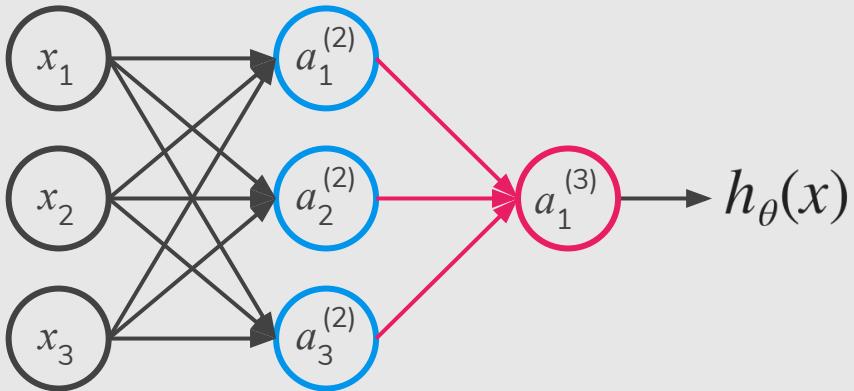
$a_i^{(j)}$ “activation” of unit i in layer j

$\Theta^{(j)}$ matrix of weights controlling function mapping from layer j to layer $j + 1$

$$a_1^{(2)} = g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)}x_0 + \Theta_{31}^{(1)}x_1 + \Theta_{32}^{(1)}x_2 + \Theta_{33}^{(1)}x_3)$$



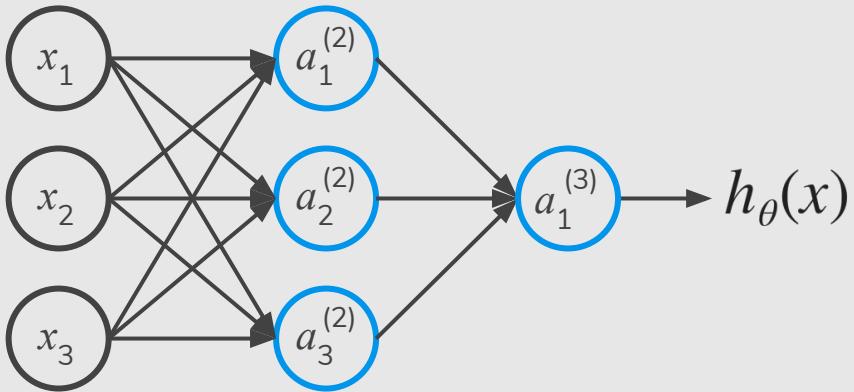
$a_i^{(j)}$ “activation” of unit i in layer j
 $\Theta^{(j)}$ matrix of weights controlling
 function mapping from layer j
 to layer $j + 1$

$$a_1^{(2)} = g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)}x_0 + \Theta_{31}^{(1)}x_1 + \Theta_{32}^{(1)}x_2 + \Theta_{33}^{(1)}x_3)$$

$$h_\Theta(x) = a_1^{(3)} = g(\Theta_{10}^{(2)}a_0^{(2)} + \Theta_{11}^{(2)}a_1^{(2)} + \Theta_{12}^{(2)}a_2^{(2)} + \Theta_{13}^{(2)}a_3^{(2)})$$

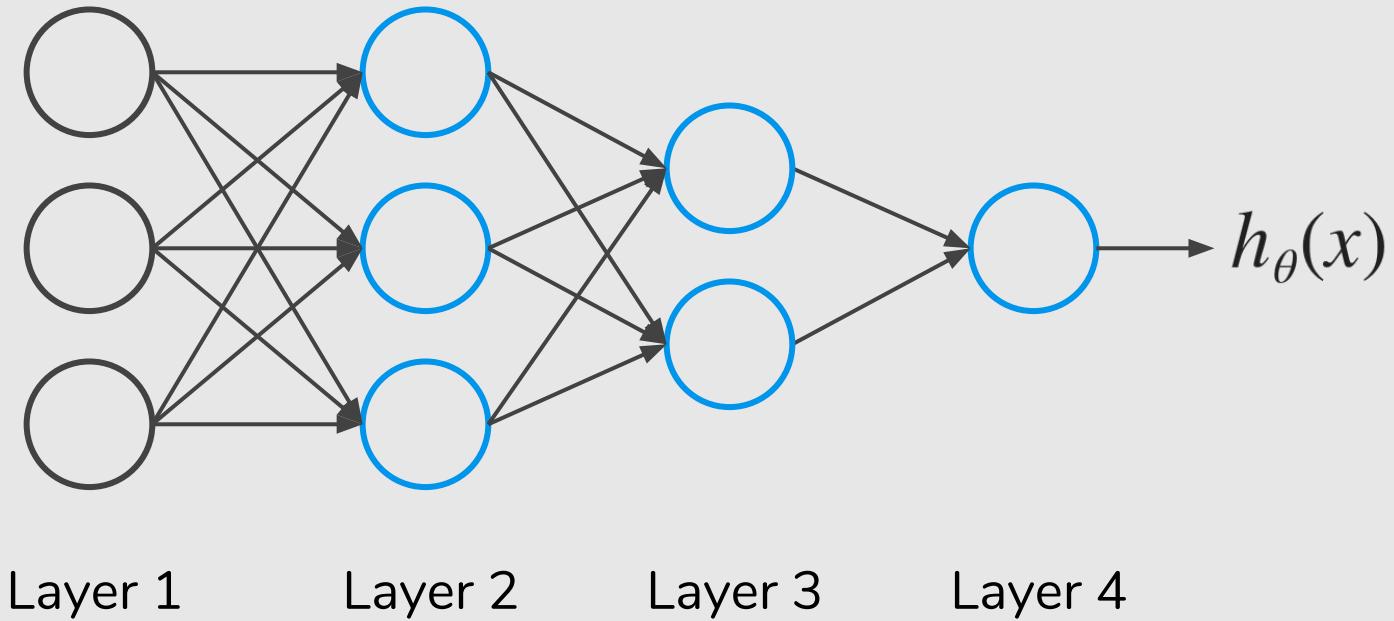


$a_i^{(j)}$ “activation” of unit i in layer j
 $\Theta^{(j)}$ matrix of weights controlling
 function mapping from layer j
 to layer $j + 1$

Feedforward Neural Network (forward propagating)

$$h_\Theta(x) = a_1^{(3)} = g(\Theta_{10}^{(2)}a_0^{(2)} + \Theta_{11}^{(2)}a_1^{(2)} + \Theta_{12}^{(2)}a_2^{(2)} + \Theta_{13}^{(2)}a_3^{(2)})$$

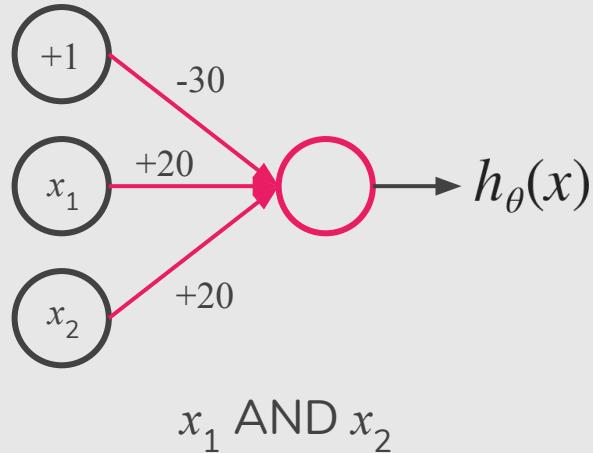
Other Network Architectures



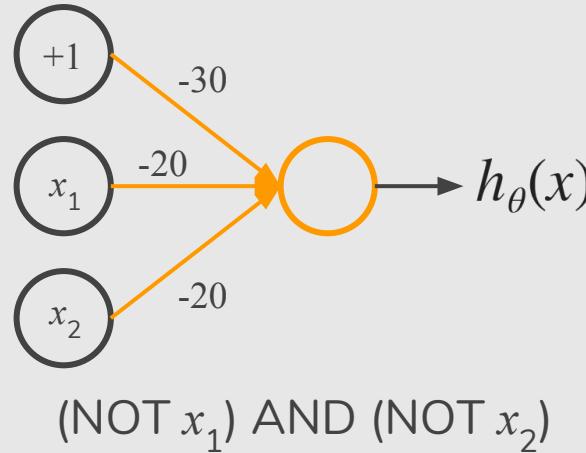
Example: XNOR $x_1, x_2 \in \{0,1\}$ $y = x_1 \text{ XNOR } x_2$

Example: XNOR

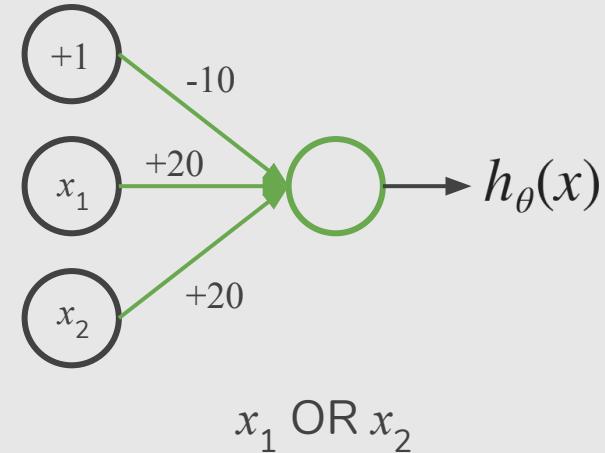
$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ XNOR } x_2$$



$x_1 \text{ AND } x_2$



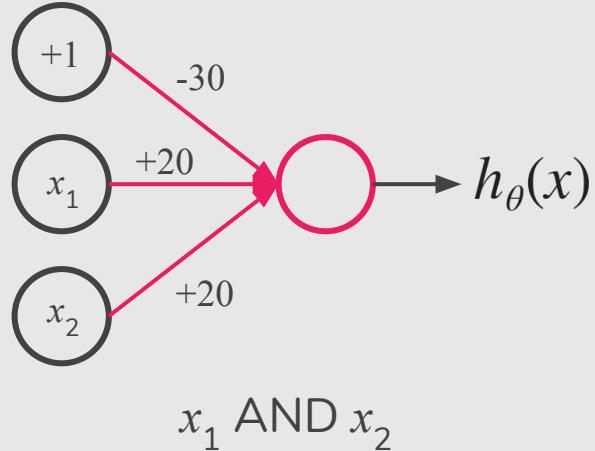
$(\text{NOT } x_1) \text{ AND } (\text{NOT } x_2)$



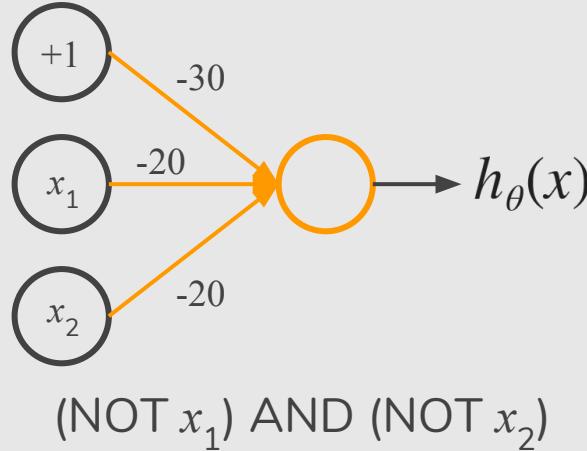
$x_1 \text{ OR } x_2$

Example: XNOR

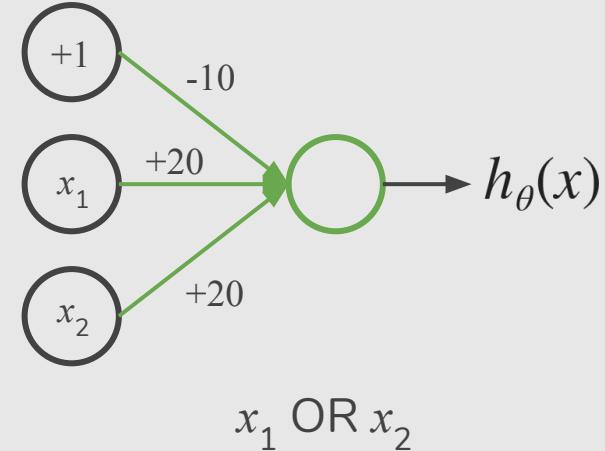
$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ XNOR } x_2$$



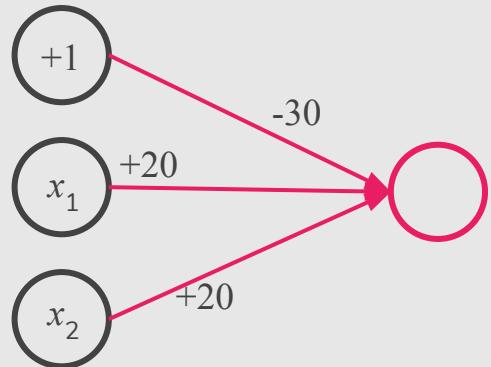
$x_1 \text{ AND } x_2$



$(\text{NOT } x_1) \text{ AND } (\text{NOT } x_2)$

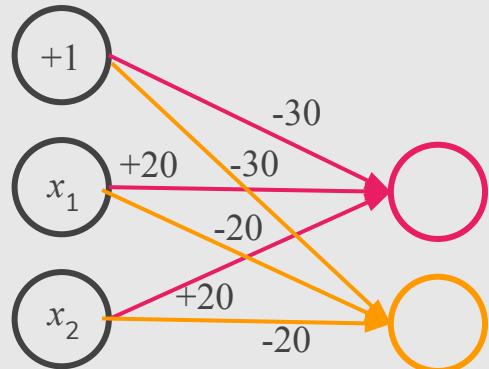
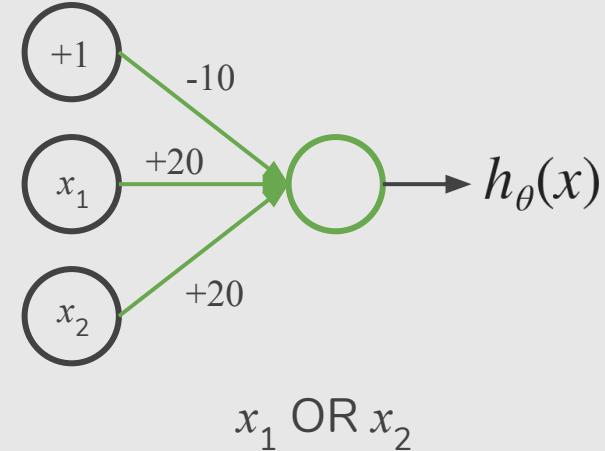
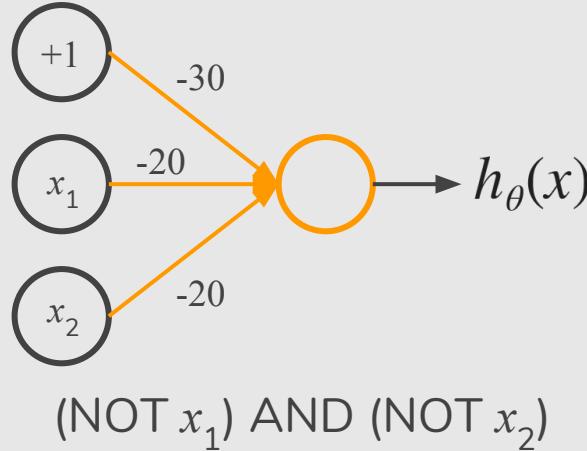
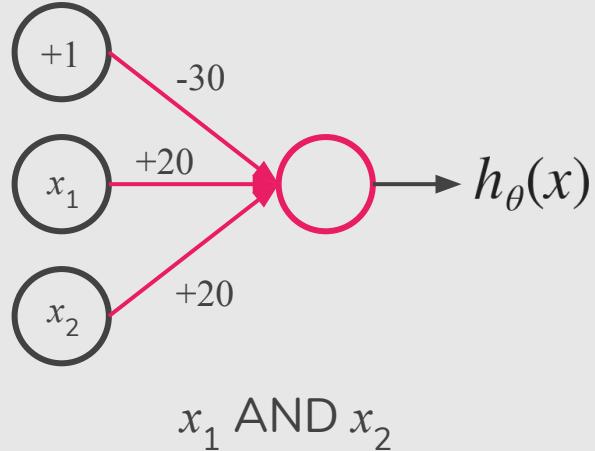


$x_1 \text{ OR } x_2$



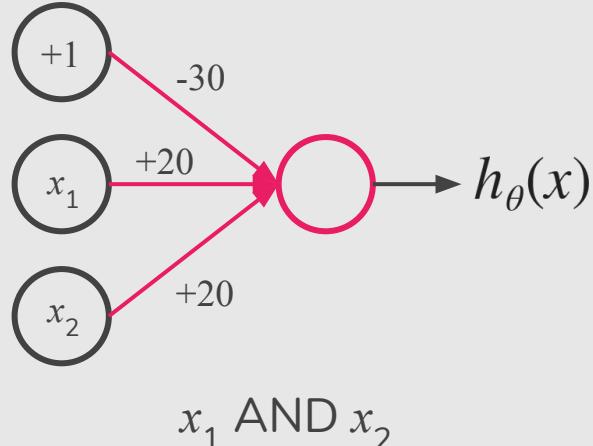
Example: XNOR

$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ XNOR } x_2$$

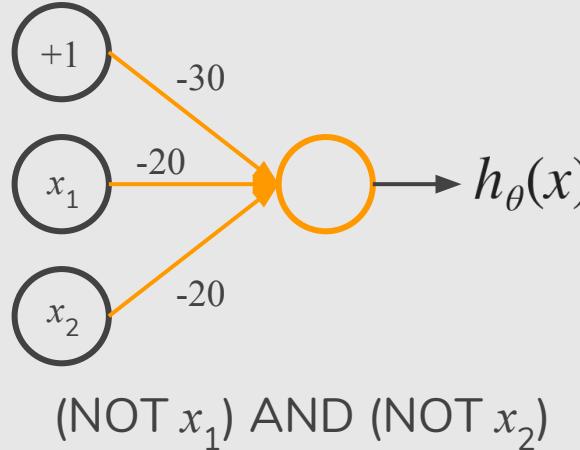


Example: XNOR

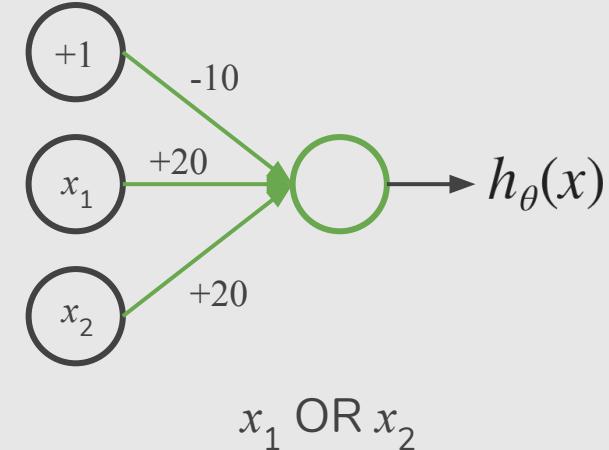
$$x_1, x_2 \in \{0,1\} \quad y = x_1 \text{ XNOR } x_2$$



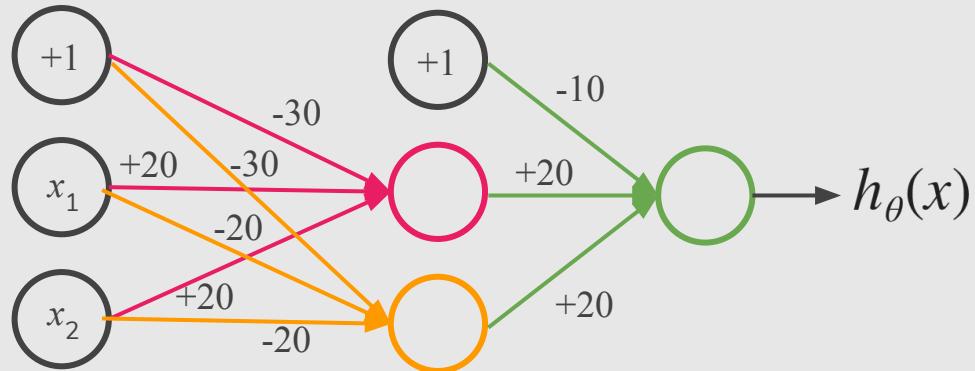
x_1 AND x_2



(NOT x_1) AND (NOT x_2)

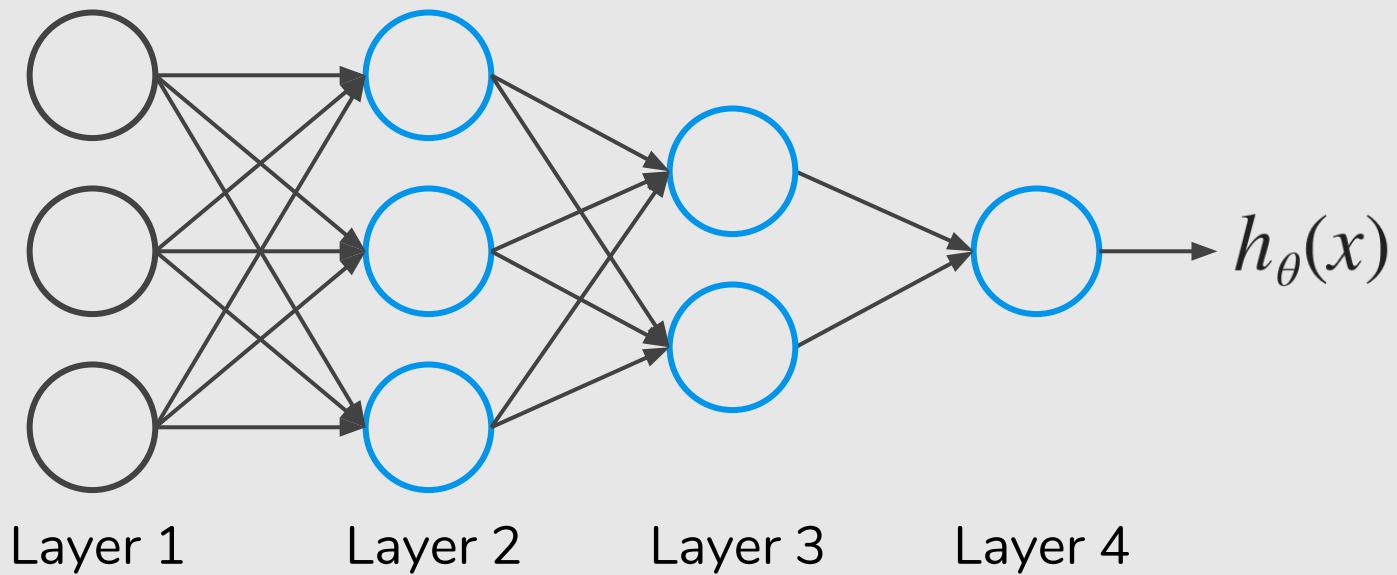


x_1 OR x_2



x_1	x_2	$h_\theta(x)$
0	0	1
0	1	0
1	0	0
1	1	1

Neural Network Intuition

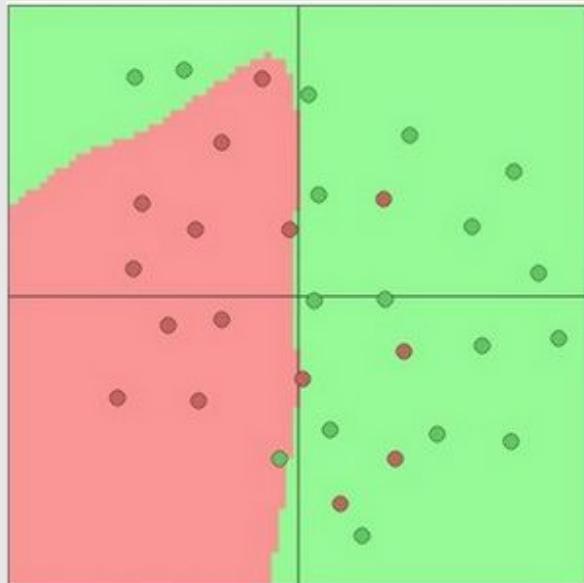


Complexity

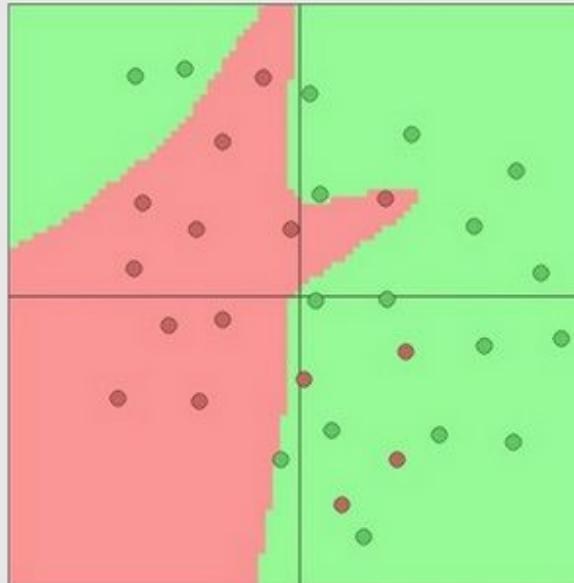
Neural Network Intuition

Toy 2d classification with 2-layer neural network

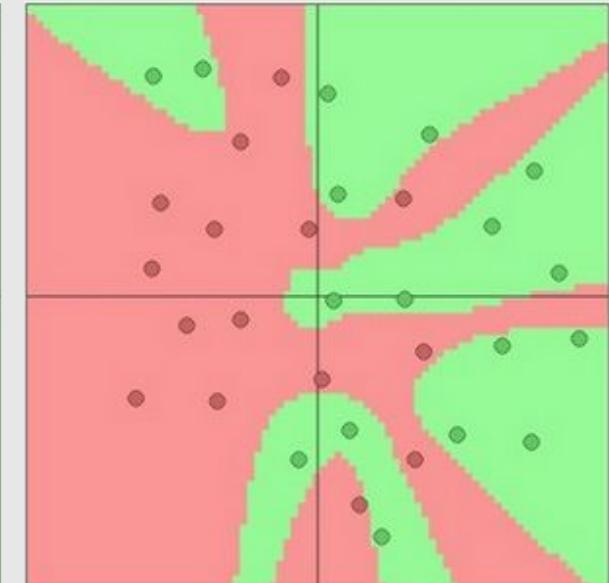
<http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>



3 hidden neurons

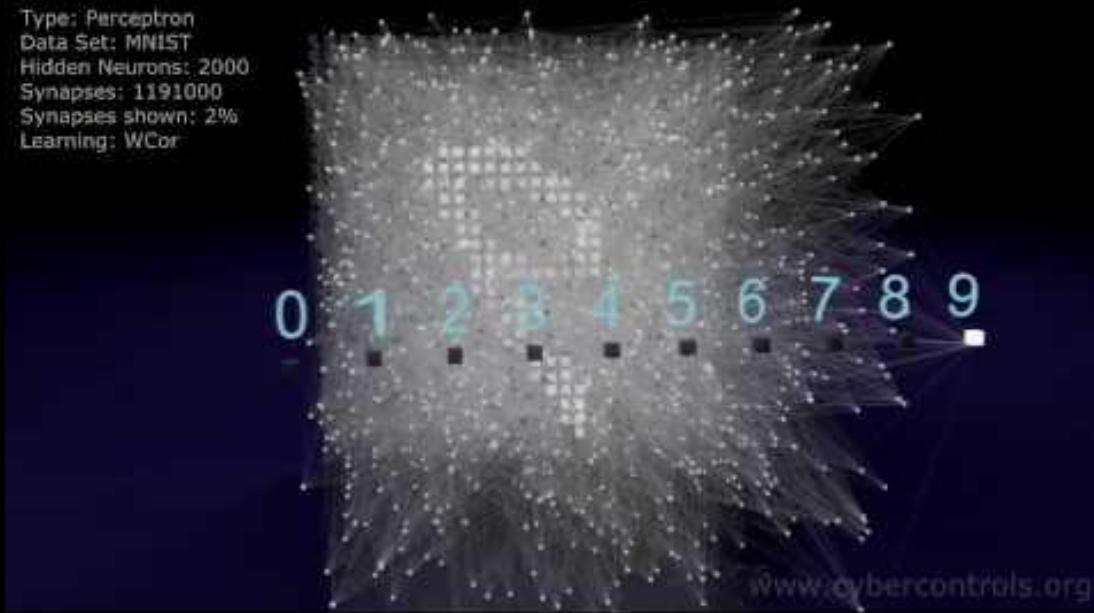


6 hidden neurons



20 hidden neurons

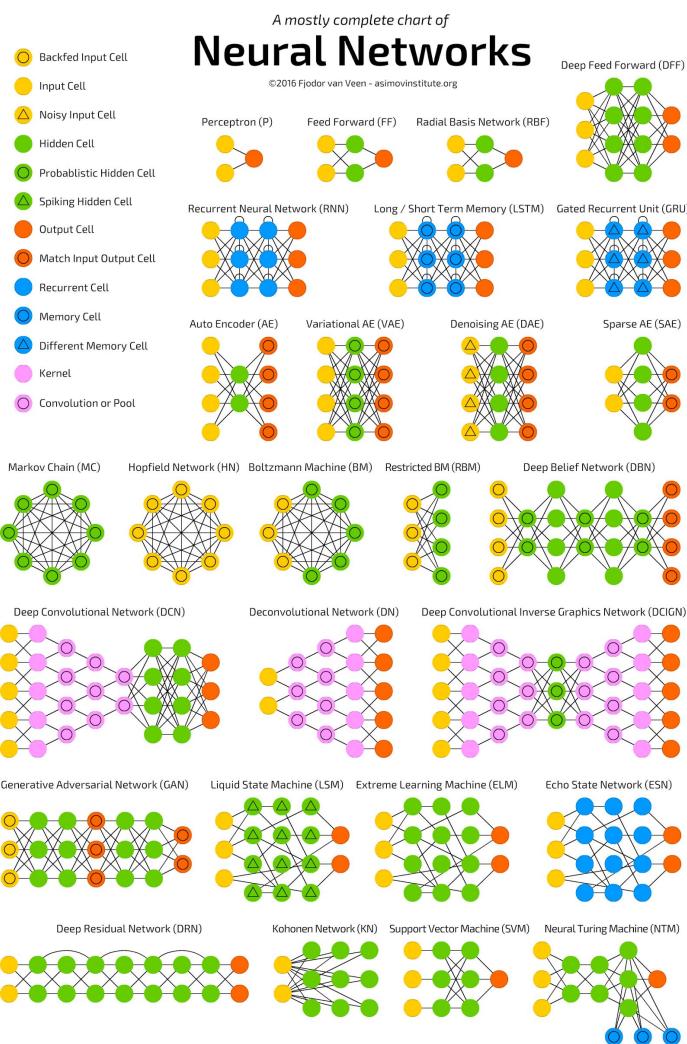
Type: Perceptron
Data Set: MNIST
Hidden Neurons: 2000
Synapses: 1191000
Synapses shown: 2%
Learning: WCor



<https://youtu.be/3JQ3hYko51Y>

Neural Network Zoo

<http://www.asimovinstitute.org/neural-network-zoo/>



To be continued ...

References

Machine Learning Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chap. 10
- Pattern Recognition and Machine Learning, Chap. 5
- Pattern Classification, Chap. 6
- Free online book: <http://neuralnetworksanddeeplearning.com>

Machine Learning Courses

- <https://www.coursera.org/learn/machine-learning>, Week 4 & 5
- <https://www.coursera.org/learn/neural-networks>